Transition Plan

OneView UDB Tuning – Assignee: *Name omitted*

June reorg activities have been moved to July. Targeting the July infrastructure weekend, which is Sunday July 20th. Timings are to be determined. PVS testing of the reorgs are currently scheduled for Thursday, June 26th – 27th. In order to run these reorgs HADR must be taken down, reorgs/index changes run, and then backup/restore to the standby. Whoever runs the reorgs MUST monitor the db2logs/archlogs directories VERY carefully. The ECMCASE table is large, being reorged for the first time, and will generate many logs. They will most likely have to drop all of the indexes on the table and rebuild them after the table is reorg-ed in order to do the run stats. There are four other small tables being targeted as well (less than 100,000 records for each). Changes are summarized below:

Table #1: CCSOWNER.ECMCASE, *Row Count*: 201,193,785

-Drop Primary Key CCSOWNER.ECMCASE_PK and rebuild it with a new column.
-Drop Index CCSOWNER.XIF1ECMCASE **permanently** from table (not included in the index rebuild script).
-Reorg the table (first time I ran the reorg it took about 8 hours before it bombed out due to the active log filling in the index rebuild step. I recommend dropping all indexes).
-Rebuild all dropped indexes (including primary key).
-Runstats on table and all indexes.

Table #2: CCA.ALERT, *Row Count*: 48,716

-Drop Indexes CCA.ALERT_I5 and CCA.ALERT_I9 **permanently** from table
-Reorg the table
-Runstats on table and all remaining indexes

Table #3: INTXN.TASK_OPEN_SUM, *Row Count*: 106,510

-Drop Indexes INTXN.TASK_OPEN_SUM_IX3, INTXN.TASK_OPEN_SUM_IX4, and INTXN.TASK_OPEN_SUM_IX5 **permanently** from table.
-Reorg the table
-Runstats on table and all remaining indexes

Table #4: PARTY.EMP, *Row Count*: 3,131

-Drop Indexes PARTY.EMP_I1, PARTY.EMP_I3, and PARTY.EMP_I4 **permanently** from table.
-Reorg the table
-Runstats on table and all remaining indexes

Table #5: CCA.REF_DATA, *Row Count*: 29,598

-No unused indexes to drop
-Reorg the table
-Runstats on the table and all indexes

All the scripts are located on OneView PVS db server cilonevs0014 in this folder:

```
  cilonevs0014 OVIHPVS *PVS* /ovihipvs/tools/releases/dm/REORGS_062614
-> ls -l
total 20
drwxrwxr-x 2 ovihipvs db2iadm1 4096 Jun  5 11:17 CCA_ALERT
drwxrwxr-x 2 ovihipvs db2iadm1 4096 Jun  5 11:12 CCSOWNER_ECMCASE
drwxrwxr-x 2 ovihipvs db2iadm1 4096 Jun  5 11:36 EMP_REFDATA
drwxrwxr-x 2 ovihipvs db2iadm1 4096 Jun  5 11:25 INTXN_TASK_OPEN_SUM
drwxrwxr-x 2 ovihipvs db2iadm1 4096 Jun  5 11:32 PARTY_EMP
```

BPM/BAM UDB – Assignee: *Name omitted*

Reorgs have been scheduled since the middle of April and running via ESP scheduler without issues at this point. They run once a week on Sunday mornings at 7am.

Installing a new monitor model has caused issues in the past. The issues usually arise from having both monitor models active at the same time; a view is created that joins the EVENTS table from the old model to the new one, which causes a lot of contention on the database since the application has to go to both tables when pulling data. Also when a new model is created a new set of tables/views (~10) are created on the database. Since the reorg list is not generated dynamically every week those tables will not get reorg-ed unless DBA generates a list for the new tables.

As a result the app team will have to engage the DBAs when installing a new model so they can extract the DDL and run it for them to create the new tables/views. That way the DBAs will know what the new tables that have been created are and can add them to the reorg list. Certain tables will never have more than a few thousand records so these tables SHOULD NOT be included in the runstats that run weekly (by marking them as "NRS" in the reorg list).

Name omitted has been working on creating some indexes on some problem tables to ease the load on the SQL calls that happen on a daily basis. We have been implementing them in PVS with the app team. The process as of now is to leave HADR up and running (tables have nowhere near the volume that is in OneView). The app team brings down the JVMs, I run the index change based on *Name omitted*'s recommendation, and then reorg/runstat the table. *Name omitted* will then run an EXPLAIN to make sure DB2 took the correct access path, and then the app team will bring back the JVMs to do their checkout. For PRD implementations, we use the same window for the reorgs, and run through a similar process. The app team brings the JVMs down, I build the new index, reorg ALL the tables for both BPMSIPRD and BPMMIPRD as part of the weekly maintenance, then *Name omitted* runs his EXPLAINs, and the app team brings up the JVMs and does their checkout. As of now, runs smoothly and doesn't cause issues with HADR. So this process is ironed out.

WPS/Advantage UDB – Assignee: *Name omitted*

Nothing new in this area. They are continuing to do nightly reorgs to help performance, these have been in place since mid-January of this year. All of the PMRs that were opened for the performance issues back in mid-December have been closed for the meantime, as IBM has basically stated that this is how the application behaves and there is not much more we can do about it (we've implemented reorgs, made changes to their JVMs, etc). So far the app has been stable and should continue to be going forward. Will need to monitor it closely during the end of year open enrollment period as the volume jumps up dramatically.

RXDM Oracle – Assignee: *Name omitted*

Pharmacy Data Mart that was transitioned to me from Rakesh. He knows the database well. As of now I have been supporting their releases, which they have once or twice a month. They follow the process and open a CM and assign me a task/as an approver. Most releases are scheduled for Thursday nights. I run through their scripts in their PVS environment (sometimes it takes multiple runs because they don't get everything correct the first time). I send the outputs to their development team so they can make changes to the SQL if required. Once everything is completely tested/vetted, I give my approval for the change.

All of their scripts are located in this library: \\wiwvsg199\rbs\8 Implementation\Release Information

Most of their SQL is pretty straightforward, usually just table changes/creation and granting access to those tables for various users.

CA Monitoring Tool – Assignee: UDB Prod Support

All of the production UDB databases besides the HADR databases have had the changes put into place to install the CA Monitoring tool. The changes are below:

db2 update dbm cfg using SYSCTRL_GROUP MONITOR – **Need to recycle the database instance**
db2 grant connect on database to user saadmin – **Need to grant this user to ALL databases on instance**

The reason the HADR databases have not had the changes put into place is that one of the changes requires a recycle of the instance. In an HADR environment if you recycle either the primary or standby while HADR is up and running then TSAMP (cluster manager) may cause the server hosting that instance to reboot because it cannot connect to the instance that was recycled. We have seen this happen many times when doing server maintenance for the HADR dbs. In order to update the HADR dbs, we most likely will need an outage so that we can take HADR down and bounce both instances at the same time.

OneView wants to be one of the last applications to be migrated so they aren't scheduled until mid-July. The OneView team is working on getting a window so that we can make these changes in PVS to test them out (they had issues when installing CA when OneView was brought over back in July/August of last year, hence their reluctance). For BPM/BAM and WPS/Advantage, I have updated the PVS environments and Ram is working with the CA team to test out CA on those dbs. Everything else is good.

InfoCenter – Assignee: *Name omitted*

Mohan has been working on collecting the UDB data needed to display in the InfoCenter. I've started the development of a few pages but it is still a work in progress. We did an extract of what was currently in APEX in order to maintain that information in the new InfoCenter going forward. A lot of the backup/size information that was not being collected previously (especially in the case of the Linux server) is now being pulled. Matt is working with the Linux team to set up the "db2dba" ID so that it can SFTP to all of the Linux servers to get the data required in one place, and then transfer that to the Oracle back end for the InfoCenter.

OEM for DB2 Plugin – Assignee: *Name omitted*

I had been working with Nik to apply an OEM DB2 plugin on a few test UDB servers, to see if we can pull similar information that we use OEM for Oracle for our UDB databases. Right now the following test servers have been set up for the OEM plugin:

Standard Changes – Assignee: Oracle Prod Support Team

Worked with Rakesh and Mike L. to develop the new set of standard changes for our group. We submitted nine updated standard changes to the Change Management group but they have not yet approved the new set of changes. There is a total of nineteen that were created to handle various DBA functions both 24/7 and from 6pm – 6am. Need to follow up with Change Management team for their approvals.

They are located in this directory: \Midrange_DBA\ORACLE\Standard_Change_Templates

UDB Playbook/Troubleshooting Guide – Assignee: UDB Prod Support

Developed a UDB Troubleshooting guide to list out what the process is when a technician is engaged in a P1 situation, which also includes some custom scripts and useful commands to run when diagnosing a problem. This is going to be a living and breathing document, so that it can be maintained by the technicians going forward. They will be able to update it according to new issues that may pop up (especially with the 10.5 upgrade that is being planned for later this year).

It is located here:
link removed

Ops Console UDB – Assignee: TBD

These were a set of UDB databases I created for *Name omitted* and the ETL group back in January – March. They were built using SQL generated by IBM and the production database has been in place since the end of March. They were very small, and have reorgs running once a week. As of yet have not had any issues for these databases.

Outstanding Tickets: None
Outstanding CMs: None
Work Requests: None