

```
In [28]: ## Provide a wider display for easier viewing
from IPython.core.display import display, HTML
display(HTML("<style>.container { width:100% !important; }</style>"))
```

```
In [2]: ## Remove future warnings for Pandas
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

## Import the necessary Libraries
%matplotlib inline
import pandas as pd
import numpy as np
from sklearn import *
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import norm
import tensorflow as tf
from sklearn.multioutput import MultiOutputRegressor
from sklearn.model_selection import cross_val_score, train_test_split, GridSearchCV, StratifiedKFold, KFold
from sklearn.metrics import *
import category_encoders as ce
from sklearn.preprocessing import *
## from mlxtend.feature_selection import SequentialFeatureSelector as sfs
from sklearn.compose import TransformedTargetRegressor
```

**Quation 1. (50 points)** Use numeric prediction techniques to build a predictive model for the HW3.xlsx dataset. This dataset is provided on the course website and contains data about whether or not different consumers made a purchase in response to a test mailing of a certain catalog and, in case of a purchase, how much money each consumer spent. The data file has a brief description of all the attributes in a separate worksheet. Note that this dataset has two possible outcome variables: Purchase (0/1 value: whether or not the purchase was made) and Spending (numeric value: amount spent).

## Dictionary

Codelist				
Var. #	Variable Name	Description	Variable Type	Code Description
1.	US	Is it a US address?	binary	1: yes 0: no
2 - 16	Source_*	Source catalog for the record (15 possible sources)	binary	1: yes 0: no
17.	Freq.	Number of transactions in last year at source catalog	numeric	
18.	last_update_days_ago	How many days ago was last update to cust. record	numeric	
19.	1st_update_days_ago	How many days ago was 1st update to cust. record	numeric	
20.	Web_order	Customer placed at least 1 order via web	binary	1: yes 0: no
21.	Gender=mal	Customer is male	binary	1: yes 0: no
22.	Address_is_res	Address is a residence	binary	1: yes 0: no
23.	Purchase	Person made purchase in test mailing	binary	1: yes 0: no
24.	Spending	Amount spent by customer in test mailing (\$)	numeric	

## A. Data pre-processing and pre-analysis

1. Read in the data
2. Explore the features and target variables to assess what parameters will need to be changed
3. Prepare & transform data for data mining process

```
In [3]: ## Read in the data

purchases_df = pd.read_excel("HW3.xlsx")
```

In [4]: *## Observe the first five values to make sure data was read in correctly*

```
purchases_df.head()
```

Out[4]:

	sequence_number	US	source_a	source_c	source_b	source_d	source_e	source_m	source_o	source_h	...	source_x	source_w	Freq	last_
0	1	1	0	0	1	0	0	0	0	0	...	0	0	2	
1	2	1	0	0	0	0	1	0	0	0	...	0	0	0	
2	3	1	0	0	0	0	0	0	0	0	...	0	0	2	
3	4	1	0	1	0	0	0	0	0	0	...	0	0	1	
4	5	1	0	1	0	0	0	0	0	0	...	0	0	1	

5 rows × 25 columns



In [5]: *## Run summary statistics on every column except for the "Source", "Last Update", "First Update" columns*  
*## Just to check for missing values*

```
## Create a list of the columns we want to subset  
summary_columns = ['US', 'Freq', 'Web order', 'Gender=male',  
                  'Address_is_res', 'Purchase', 'Spending']  
  
purchases_df[summary_columns].describe()
```

Out[5]:

	US	Freq	Web order	Gender=male	Address_is_res	Purchase	Spending
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	0.824500	1.417000	0.426000	0.524500	0.221000	0.500000	102.560745
std	0.380489	1.405738	0.494617	0.499524	0.415024	0.500125	186.749816
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	1.000000	0.000000	1.000000	0.000000	0.500000	1.855000
75%	1.000000	2.000000	1.000000	1.000000	0.000000	1.000000	152.532500
max	1.000000	15.000000	1.000000	1.000000	1.000000	1.000000	1500.060000

```
In [6]: ## We check to see what our "target" variable is - Purchase  
## It looks like it is a binary classification problem.  
## Since we're using numeric prediction, we can leave this the same  
  
purchases_df.Purchase.unique()
```

```
Out[6]: array([1, 0], dtype=int64)
```

```
In [7]: ## View splits for the targets  
  
total_obs = purchases_df.groupby("Purchase")["US"].count().sum()  
total_purchases = purchases_df.groupby("Purchase")["US"].count()[0]  
total_nopurchase = purchases_df.groupby("Purchase")["US"].count()[1]  
  
## Evenly distributed right in the middle  
print("{} total customers purchased, \  
      {} % of overall observations".format(total_purchases, round(total_purchases/total_obs*100, 2)))  
print("{} total customers did not purchase, \  
      {} % of overall observations".format(total_nopurchase, round(total_nopurchase/total_obs*100, 2)))  
  
1000 total customers purchased,      50.0 % of overall observations  
1000 total customers did not purchase,      50.0 % of overall observations
```

**We see an even split of the data. I don't see a need to stratify in this instance because of the split of the information. If they were skewed one way or the other then it might be needed.**

**Let's start building our models!**

```
In [10]: ## Remove any columns that will not be used as features  
## I am making a conscious decision to remove sequence_number - I believe it is extraneous and not needed  
## It won't help the predictive power of the models  
  
feature_cols = purchases_df.columns[~purchases_df.columns.isin(["sequence_number", "Purchase", "Spending"])]  
feature_cols
```

```
Out[10]: Index(['US', 'source_a', 'source_c', 'source_b', 'source_d', 'source_e',  
               'source_m', 'source_o', 'source_h', 'source_r', 'source_s', 'source_t',  
               'source_u', 'source_p', 'source_x', 'source_w', 'Freq',  
               'last_update_days_ago', '1st_update_days_ago', 'Web order',  
               'Gender=male', 'Address_is_res'],  
            dtype='object')
```

```
In [11]: ## Create our "X" variable that contains all the features we are curious about plotting  
X = np.array(purchases_df[feature_cols])  
  
## Create our "y" variable which is our target variable  
y = np.array(purchases_df["Spending"])  
  
## Split data training 70 % and testing 30%  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

**First thing I will do is assess which features are actually needed. To do this, I will use both a Principle Component Analysis, along with a Feature Importance Plot.**

```
In [13]: ## We normalize our training and testing data PCA to work correctly
## We don't want to skew the results of the plot because some features are not on the same scale
## We perform this normalization AFTER splitting the data - again, so that we don't skew the training
## data with the testing data.

## Normalization is the process of scaling individual samples to have unit norm. This process can be useful
## if you plan to use a quadratic form such as the dot-product or any other kernel to quantify the similarity of any pair of samples.
X_train_norm = Normalizer(norm = "l1").fit_transform(X_train)
X_train_robust = RobustScaler().fit_transform(X_train)
X_train_power = PowerTransformer(method='yeo-johnson', standardize=False).fit_transform(X_train)
X_train_standard = StandardScaler().fit_transform(X_train)

## Import PCA from sklearn
from sklearn.decomposition import PCA

## Initialize two new PCA instances, we'll use this to plot the training data using two different transformations
## Normalizer will likely be skewed by the outliers in each of the features being used
## Robust is going to transform feature values to be larger than the previous scalers and more importantly are approximately similar to original data
## PowerTransformer is a family of parametric, monotonic transformations that aim to map data from any distribution to as close to a Gaussian distribution
## as possible in order to stabilize variance and minimize skewness.

pca = PCA()
pca_r = PCA()
pca_p = PCA()
pca_s = PCA()

## The goal of this plot is to determine what features need to be included in our models
## In the first few homeworks, we threw the kitchen sink at the models. Here, we are going to be
## more refined in our analysis.

## Train the PCA instance using the normalized training data
pca.fit(X_train_norm)
pca_r.fit(X_train_robust)
pca_p.fit(X_train_power)
pca_s.fit(X_train_standard)

plt.figure(1, figsize=(15, 10))
plt.clf()
plt.axes([.2, .2, .7, .7])


plt.plot(pca.explained_variance_ratio_, 'r--', linewidth = 2)
plt.plot(pca_r.explained_variance_ratio_, 'b--', linewidth = 2)
plt.plot(pca_p.explained_variance_ratio_, 'y--', linewidth = 2)
plt.plot(pca_s.explained_variance_ratio_, 'g--', linewidth = 2)
```

```
## Set plot labels
plt.xlabel('Number of features required to explain variance')
plt.ylabel('Amount of Variance Explained by Features (0.0 to 1.0)')

## Explicitly set the x-axis data so we can see where the drop-off is
plt.xticks(np.arange(0, 21, step=1))

## Show the graph!
plt.show()
```

```
C:\Python\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.  
warnings.warn(msg, DataConversionWarning)  
C:\Python\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.  
warnings.warn(msg, DataConversionWarning)
```



**We only need three features at most, based on this principal component analysis.**

**Now let's see which features are the ones that actually explain the most variance.**



In [14]: *## SET A VARIABLE TO SWAP IN AND OUT BETWEEN THE DIFFERENT TRAINING INSTANCES*

```
#Z = X_train_robust
#Z = X_train_norm
#Z = X_train_power
Z = X_train_standard
```

*## Start with identifying the best features using a Random Forest classifier*

*## Create a new classifier*

```
clf_rf_5 = ensemble.RandomForestRegressor()
clr_rf_5 = clf_rf_5.fit(Z, y_train)
```

*## Save our importances to a variable*

```
importances = clr_rf_5.feature_importances_
```

*## Get the standard deviation for each feature*

```
std = np.std([tree.feature_importances_ for tree in clf_rf_5.estimators_],
             axis=0)
indices = np.argsort(importances)[::-1]
```

*## Print the feature ranking*

```
print("Feature ranking:")
```

*## Print the top five features, and their importance based on the Random Forest Classifier*

```
for f in range(0, 5):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))
```

*## Plot the feature importances of the Random Forest Regressor - to see this visually*

*## Set the plot size*

```
plt.figure(1, figsize=(15, 10))
```

*## Set the title*

```
plt.title("Feature importances")
```

*## Plot a graph using all of the normalized features*

```
plt.bar(range(Z.shape[1]), importances[indices],
        color="b", yerr=std[indices], align="center")
```

```
plt.xticks(range(Z.shape[1]), feature_cols[indices], rotation=90)
```


```
plt.xlim([-1, Z.shape[1]])
```

*## Show the graph!*

```
plt.show()
```

Feature ranking:

1. feature 16 (0.549566)
2. feature 17 (0.169993)
3. feature 18 (0.114418)
4. feature 21 (0.036042)
5. feature 9 (0.021158)



**Confirming my earlier observation, only three of the features explain 82% of the variance between all the features.**

```
In [15]: ### Build a custom function to transform feature column or input into a scaled and then labeled item
def robustscale_labelencoder(array):

    ## Initialize a RobustScaler object
    rs = RobustScaler()
    ## Initialize a LabelEncoder object
    le = LabelEncoder()

    ## Reshape the array to conform to the right dimensions needed for RobustScaler
    array = array.reshape(-1, 1)

    ## Transform array using Robust Scaling
    robust_array = rs.fit_transform(array)

    ## Reshape the array back to its original shape for Label Encoding
    robust_array = robust_array.flatten()

    ## Now apply the LabelEncoding on a better scaled version of the data
    labeled_array = le.fit_transform(robust_array)

    ## Return the output
    return labeled_array
```

```
In [18]: ## Use that cool function to transform the training target for use in the LR Classifier to observe what
## features are important
# y_train_robust_labeled = robustscale_labelencoder(y_train)

## Use StandardScaling on the training target. Not likely to help with training the model, from what I have seen
## in previous attempts to train the models

y_train_standard = StandardScaler().fit_transform(y_train.reshape(-1, 1))

y_train_labeled = LabelEncoder().fit_transform(y_train)
```

```
In [20]: ## Build RF classifier to use in feature selection
clf = linear_model.LogisticRegression(C=1e5)

# Build step forward feature selection
sfs1 = sfs(clf,
            k_features=6,
            forward=True, # Otherwise, this will be the backward selection
            floating=False,
            n_jobs=10, # The number of CPUs to use for evaluating
            verbose=2,
            scoring='accuracy',
            cv=5)

# Perform SFFS
sfs1 = sfs1.fit(X_train_standard, y_train_labeled)
```

[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 15 out of 22 | elapsed: 54.2s remaining: 25.2s

[Parallel(n\_jobs=10)]: Done 22 out of 22 | elapsed: 58.7s finished

[2019-10-30 14:17:06] Features: 1/6 -- score: 0.44553618151441493[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 13 out of 21 | elapsed: 1.2min remaining: 42.4s

[Parallel(n\_jobs=10)]: Done 21 out of 21 | elapsed: 1.2min finished

[2019-10-30 14:18:22] Features: 2/6 -- score: 0.47563816778106094[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 12 out of 20 | elapsed: 1.4min remaining: 56.2s

[Parallel(n\_jobs=10)]: Done 20 out of 20 | elapsed: 1.5min finished

[2019-10-30 14:19:50] Features: 3/6 -- score: 0.5165274559056995[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 10 out of 19 | elapsed: 59.4s remaining: 53.4s

[Parallel(n\_jobs=10)]: Done 19 out of 19 | elapsed: 1.8min finished

[2019-10-30 14:21:39] Features: 4/6 -- score: 0.5331336778749486[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 9 out of 18 | elapsed: 1.1min remaining: 1.1min

[Parallel(n\_jobs=10)]: Done 18 out of 18 | elapsed: 2.0min finished

[2019-10-30 14:23:43] Features: 5/6 -- score: 0.5403199214808022[Parallel(n\_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n\_jobs=10)]: Done 7 out of 17 | elapsed: 1.3min remaining: 1.9min

[Parallel(n\_jobs=10)]: Done 17 out of 17 | elapsed: 2.3min finished

[2019-10-30 14:26:03] Features: 6/6 -- score: 0.5471182950056331


```
In [21]: ## Now we print out what feature columns of Linear Regressor picked out - we can use these to subset later  
## We'll remember to grab these columns further along in the analysis  
feat_cols = list(sfs1.k_feature_idx_)  
print(feat_cols)  
  
[3, 7, 8, 16, 17, 21]
```

**We do a little visual exploration of the features to see what the model is using for data points; this will help us with encoding the target correctly, as well as selecting the proper distance metric for kNN and the various forms of regression.**

```
In [22]: ### Reset seaborn to the default background - for better viewing
sns.set()

## Build a quick plot comparing two of the features column with our spending using a simple regression line
sns.pairplot(purchases_df, x_vars=['last_update_days_ago', 'Freq'], y_vars='Spending', height = 8, aspect = 0.7, kind =
'reg')

## Show the plot
plt.show()
```



**Interesting observation to be made here - the "last update" feature column has a tremendous amount of spread between how much a customer spent, and shows almost no linear relationship at all.**

**Conversely, the "freq" feature column is a little bit more linear, but not quite. So this is telling me that my linear model is probably going to suffer because it won't be able to distinguish between targets very well.**

**I will have to scale the features and targets correctly, or many of the data mining models will perform badly.**



In [23]: *## Plotting out the distribution of spending - to see how much a spread my target is*

*## Import a normal distribution from scipy*

**from** **scipy.stats** **import** norm

*## Build a new plot*


plt.figure(1, figsize=(15, 6))

*## Use a special function in seaborn to build a distribution plot and also include a*

sns.distplot(purchases\_df.Spending, fit = norm, kde = **True**, bins = 20)

*## Show the plot*

plt.show()



**Most of the spending is less than \$20! We will probably want to standardize/normalize the target first, and then use that as the prediction.**

```
In [25]: ## Create two new encoders for the target - we'll use this to see  
rs = RobustScaler()  
pt = PowerTransformer(method='yeo-johnson', standardize=False)  
le = LabelEncoder()  
nm = Normalizer(norm = "l1")  
  
## Reshape the data because these scaling methods are only supposed to be used by 2D arrays  
y_train_reshape = y_train.reshape(-1, 1)  
  
## Create four versions of the training data, to see the distributions for each one  
robust_target = rs.fit_transform(y_train_reshape)  
power_target = pt.fit_transform(y_train_reshape)  
label_target = le.fit_transform(y_train)  
normal_target = nm.fit_transform(y_train_reshape)
```

```
In [26]: ## Plotting out the distribution of spending using the new target variables with new encoding

target_list = [robust_target, power_target, label_target, normal_target]
titles = ["Spending Distribution Using Robust Scaling", "Spending Distribution Using Power Transformation",
          "Spend Distribution Using Label Encoding", "Spending Distribution Using Normalization"]

for i,t in zip(target_list, titles):

    # Build a new plot
    plt.figure(1, figsize=(15, 6))

    ## Use a special function in seaborn to build a distribution plot and also include a
    sns.distplot(i, fit = norm, kde = True, bins = 20)


    plt.title(t)

    ## Show the plot
    plt.show()
```


Spending Distribution Using Robust Scaling




Spending Distribution Using Power Transformation



Spend Distribution Using Label Encoding



Spending Distribution Using Normalization



**Here's a analysis of the behavior of the different encoding strategies for the target variable. Based on the above distribution patterns, and knowing how the original "Spending" column is distributed, I will use normalize the target variable for training the model.**

## **What this pre-processing work has told me is that**

1. The target variable is heavily skewed to the left side, most purchases are going to be for low amounts of money
2. We will want to normalize the target variable, to get them on the same scale, otherwise the regression will suffer
3. We're dealing with a lot of outliers, especially the cases where the spend is quite large - we will want to use L1 normalization
4. Many of the features are binary variables and don't really account for the variance or spread
5. In fact, many of them are probably not needed. So we can build models using all the features or a subset

## **B. Model Creation and Evaluation**

1. Create parameter grids for each model
2. Used nested cross validation to determine the best model
3. Tune the hyper parameters for the best model
4. Evaluate the models on the testing data

```
In [40]: ### RE INITIALIZE OUR TARGETS AND THEN TRANSFORM THE LABELS AND TARGET FOR TRAINING

## Create our "X" variable that contains all the features we are curious about plotting
X = np.array(purchases_df[feature_cols])

## Create our "y" variable which is our target variable
y = np.array(purchases_df["Spending"])

## Split data training 70 % and testing 30%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)

## Normalize the features
X_train_normal = nm.fit_transform(X_train)

## Reshape y_train (target variable) for normalization technique
y_train_reshaped = y_train.reshape(-1, 1)
y_train_labeled = le.fit_transform(y_train)

## Print the results of the transformation
print(X_train_normal)
print()
print(y_train_labeled)
```

```
[[0.00024969 0.          0.00024969 ... 0.          0.00024969 0.          ]
 [0.00020665 0.          0.          ... 0.00020665 0.00020665 0.          ]
 [0.00030516 0.          0.          ... 0.00030516 0.          0.00030516]
 ...
 [0.00039573 0.          0.00039573 ... 0.00039573 0.          0.00039573]
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 0.00035625 0.          0.          ]]

[349 190   2 ...   0   0   0]
```

**Here are the different models I will be using for my analysis of this regression problem**

1. Linear Regression (requires all numeric)
2. k-Nearest Neighbors
3. Decision Tree (requires pre-processing)
4. SVM Regression (requires pre-processing of data)
5. Ensemble/Gradient Boost
6. Neural Network (Keras & KerasTuner)

In [36]: ##### SETTING UP PARAMETER GRIDS FOR THE DIFFERENT DATA MINING MODELS WE ARE GOING TO USE #####

```
## Set up a grid for the Logit Regressor
## Using L1 only, this is a sparse data set. And this data mining technique is going to suffer anyways
lr_p_grid = {"penalty": ["l1"],
             "C": [1, 5, 10, 50, 1000],
             "solver": ["liblinear"]}

## Set up a grid for kNN Regressor
## Going to use 1-30 neighbors, and two different distance calculations
knnr_p_grid = {"n_neighbors": list(range(1, 31)),
               "weights": ["uniform", "distance"]}

## Set up a grid for the DecisionTree Regressor
dtr_p_grid = {"criterion": ["mse", "mae"],
              "splitter": ["best", "random"],
              "max_features": [4, 6, 10, 15],
              "max_depth": [5, 10, 15]}

## Set up a grid for the Support Vector Regressor
svr_p_grid = {"C": [1, 5, 10, 50, 1000],
              "gamma": [0.0001, 0.0005, 0.001, 0.005],
              "kernel": ["poly", "rbf"]}

## Set up a grid for GBoost
#### Use least squares for the regression
gbr_p_grid = {'loss': ['ls'],
              'n_estimators': [100, 200, 300, 400, 500],
              'max_depth': [3, 4, 5],
              'min_samples_split': [2, 4, 6],
              'learning_rate': [0.01]}
```

In [37]: ## Set a number of trials to run for the models

```
num_trials = 20

## Empty arrays to store scores for classifier
nested_scores_lr = np.zeros(num_trials)
nested_scores_knnr = np.zeros(num_trials)
nested_scores_dtr = np.zeros(num_trials)
nested_scores_svr = np.zeros(num_trials)
nested_scores_gbr = np.zeros(num_trials)
```



```
In [38]: ## Create new regressors for each data mining technique
```

```
## Linear Regression  
lr = linear_model.LogisticRegression()  
  
## k-Nearest Neighbors  
knnr = neighbors.KNeighborsRegressor()  
  
## Decision Tree  
dtr = tree.DecisionTreeRegressor()  
  
## Support Vector Machine  
svr = svm.SVR()  
  
## Gradient Boost  
gbr = ensemble.GradientBoostingRegressor()
```

```
In [39]: ## All of the new regressors initialized correctly
```

```
print(lr, knnr, dtr, svr, gbr)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='warn',  
    n_jobs=None, penalty='l2', random_state=None, solver='warn',  
    tol=0.0001, verbose=0, warm_start=False) KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkows  
ki',  
    metric_params=None, n_jobs=None, n_neighbors=5, p=2,  
    weights='uniform') DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,  
    max_leaf_nodes=None, min_impurity_decrease=0.0,  
    min_impurity_split=None, min_samples_leaf=1,  
    min_samples_split=2, min_weight_fraction_leaf=0.0,  
    presort=False, random_state=None, splitter='best') SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=  
0.1,  
    gamma='auto_deprecated', kernel='rbf', max_iter=-1, shrinking=True,  
    tol=0.001, verbose=False) GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,  
    learning_rate=0.1, loss='ls', max_depth=3, max_features=None,  
    max_leaf_nodes=None, min_impurity_decrease=0.0,  
    min_impurity_split=None, min_samples_leaf=1,  
    min_samples_split=2, min_weight_fraction_leaf=0.0,  
    n_estimators=100, n_iter_no_change=None, presort='auto',  
    random_state=None, subsample=1.0, tol=0.0001,  
    validation_fraction=0.1, verbose=0, warm_start=False)
```

Now we'll use Nested Cross Validation, combined with GridSearch to find out what models perform well.

I am going to build one loop for each regression method to reduce the load on my machine

```
In [52]: ##### Training ground - just to make sure the data is being used well
```

```
dtr.fit(X_train_normal, y_train_labeled)
```

```
Out[52]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,  
                                max_leaf_nodes=None, min_impurity_decrease=0.0,  
                                min_impurity_split=None, min_samples_leaf=1,  
                                min_samples_split=2, min_weight_fraction_leaf=0.0,  
                                presort=False, random_state=None, splitter='best')
```

```
In [59]: ## We get a lot of convergence warnings, because the data mining procedures get lost  
## I just include this to remove the warnings.
```

```
import warnings  
warnings.filterwarnings(action = "ignore", module = "sklearn")
```

```
In [60]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Logit Regression

    lreg= GridSearchCV(estimator=lr, param_grid=lr_p_grid, cv=inner_cv)
    lreg.fit(X_train_normal, y_train_labeled)

    nested_score = cross_val_score(lreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    nested_scores_lr[i] = nested_score.mean()
```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-60-b4a3015e8336> in <module>
    12     lreg.fit(X_train_normal, y_train_labeled)
    13
--> 14     nested_score = cross_val_score(lreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    15     nested_scores_lr[i] = nested_score.mean()

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_val_score(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, error_score)
    400         fit_params=fit_params,
    401         pre_dispatch=pre_dispatch,
--> 402         error_score=error_score)
    403     return cv_results['test_score']
    404

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_validate(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, return_train_score, return_estimator, error_score)
    238         return_times=True, return_estimator=return_estimator,
    239         error_score=error_score)
--> 240     for train, test in cv.split(X, y, groups))
    241
    242     zipped_scores = list(zip(*scores))

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
    918         self._iterating = self._original_iterator is not None
    919
--> 920         while self.dispatch_one_batch(iterator):
    921             pass
    922

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
    757         return False
    758     else:
--> 759         self._dispatch(tasks)
    760         return True
    761

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
    714         with self._lock:
    715             job_idx = len(self._jobs)
--> 716             job = self._backend.apply_async(batch, callback=cb)
    717             # A job can complete so quickly than its callback is
    718             # called before we get here, causing self._jobs to

C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
    180     def apply_async(self, func, callback=None):

```

```

181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel_backends.py in __init__(self, batch)
547         # Don't delay the application, to avoid keeping the input
548         # arguments in memory
--> 549         self.results = batch()
550
551     def get(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, return_estimator, error_score)
526         estimator.fit(X_train, **fit_params)
527     else:
--> 528         estimator.fit(X_train, y_train, **fit_params)
529
530     except Exception as e:

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in fit(self, X, y, groups, **fit_params)
720         return results_container[0]
721
--> 722         self._run_search(evaluate_candidates)
723
724         results = results_container[0]

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in _run_search(self, evaluate_candidates)
1189     def _run_search(self, evaluate_candidates):
1190         """Search all candidates in param_grid"""
-> 1191         evaluate_candidates(ParameterGrid(self.param_grid))
1192
1193

```

```
C:\Python\lib\site-packages\sklearn\model_selection\_search.py in evaluate_candidates(candidate_params)
709         for parameters, (train, test)
710         in product(candidate_params,
--> 711                   cv.split(X, y, groups)))
712
713     all_candidate_params.extend(candidate_params)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
918     self._iterating = self._original_iterator is not None
919
--> 920     while self.dispatch_one_batch(iterator):
921         pass
922
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
757     return False
758     else:
--> 759         self._dispatch(tasks)
760         return True
761
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
714     with self._lock:
715         job_idx = len(self._jobs)
--> 716         job = self._backend.apply_async(batch, callback=cb)
717         # A job can complete so quickly that its callback is
718         # called before we get here, causing self._jobs to
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
180     def apply_async(self, func, callback=None):
181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in __init__(self, batch)
547     # Don't delay the application, to avoid keeping the input
548     # arguments in memory
--> 549     self.results = batch()
550
551     def get(self):
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223     with parallel_backend(self._backend, n_jobs=self._n_jobs):
224         return [func(*args, **kwargs)
--> 225                 for func, args, kwargs in self.items]
226
```

```

227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, t
est, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, retu
rn_estimator, error_score)
526         estimator.fit(X_train, **fit_params)
527     else:
--> 528         estimator.fit(X_train, y_train, **fit_params)
529
530     except Exception as e:

C:\Python\lib\site-packages\sklearn\linear_model\logistic.py in fit(self, X, y, sample_weight)
1303         self.class_weight, self.penalty, self.dual, self.verbose,
1304         self.max_iter, self.tol, self.random_state,
-> 1305         sample_weight=sample_weight)
1306         self.n_iter_ = np.array([n_iter_])
1307         return self

C:\Python\lib\site-packages\sklearn\svm\base.py in _fit_liblinear(X, y, C, fit_intercept, intercept_scaling, class_wi
ght, penalty, dual, verbose, max_iter, tol, random_state, multi_class, loss, epsilon, sample_weight)
921         X, y_ind, sp.isspmatrix(X), solver_type, tol, bias, C,
922         class_weight_, max_iter, rnd.randint(np.iinfo('i').max),
--> 923         epsilon, sample_weight)
924         # Regarding rnd.randint(..) in the above signature:
925         # seed for srand in range [0..INT_MAX); due to limitations in Numpy

```

KeyboardInterrupt:

```
In [62]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Logit Regression

    knnreg= GridSearchCV(estimator=knnr, param_grid=knnr_p_grid, cv=inner_cv)
    knnreg.fit(X_train_normal, y_train_labeled)

    nested_score = cross_val_score(knnreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    nested_scores_knnr[i] = nested_score.mean()
```

```
In [64]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Logit Regression

    dtreg= GridSearchCV(estimator=dtr, param_grid=dtr_p_grid, cv=inner_cv)
    dtreg.fit(X_train_normal, y_train_labeled)

    nested_score = cross_val_score(dtreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    nested_scores_dtr[i] = nested_score.mean()
```



```
In [67]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

    ## Nested CV for Logistic Regression

    svreg= GridSearchCV(estimator=svr, param_grid=svr_p_grid, cv=inner_cv)
    svreg.fit(X_train_normal, y_train_labeled)

    nested_score = cross_val_score(svreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    nested_scores_svr[i] = nested_score.mean()
```

```
In [70]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

    ## Nested CV for Logit Regression

    gbreg= GridSearchCV(estimator=gbr, param_grid=gbr_p_grid, cv=inner_cv)
    gbreg.fit(X_train_normal, y_train_labeled)

    nested_score = cross_val_score(gbreg, X = X_train_normal, y = y_train_labeled, cv = outer_cv)
    nested_scores_gbr[i] = nested_score.mean()
```

```
In [263]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(nested_scores_svr, 'r--', label = "SVR")
plt.plot(nested_scores_knnr, 'y--', label = "kNN Regression")
plt.plot(nested_scores_lr, 'g--', label = "Linear Regression")
plt.plot(nested_scores_dtr, 'b--', label = "Decision Tree")
plt.plot(nested_scores_gbr, 'v--', label = "Gradient Boost")

## Give some labels
plt.xlabel("No. of Trials")
plt.ylabel("Score")

## Title and Legend
plt.title("Nested Cross Validation Scores")
plt.legend(loc = 'center right')

## Show the graph
plt.show()
```



**What this plot tells us is our models really suffer and don't do well with multivariate regression, which is what this problem is. We're going to try using a different method, namely Keras & Tensorflow to build a Neural Network.**

Here we start building a neural network; it allows us to build from the ground up, continually learning, and will hopefully lead to much better performance.

```
In [18]: ## Import our libraries

from tensorflow.keras import backend as K
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, Dropout
from keras.utils import np_utils
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import SGD
```

Using TensorFlow backend.

```
In [9]: ## Remove any columns that will not be used as features
## I am making a conscious decision to remove sequence_number - I believe it is extraneous and not needed
## It won't help the predictive power of the models

feature_cols = purchases_df.columns[~purchases_df.columns.isin(["sequence_number", "Purchase", "Spending"])]
feature_cols
```

```
Out[9]: Index(['US', 'source_a', 'source_c', 'source_b', 'source_d', 'source_e',
              'source_m', 'source_o', 'source_h', 'source_r', 'source_s', 'source_t',
              'source_u', 'source_p', 'source_x', 'source_w', 'Freq',
              'last_update_days_ago', '1st_update_days_ago', 'Web order',
              'Gender=male', 'Address_is_res'],
              dtype='object')
```

```
In [10]: ## Create a new "X" variable that contains all the features we are curious about plotting
X = np.array(purchases_df[feature_cols])

## Create our "y" variable which is our target variable and remove the floating point decimals
y = np.array(purchases_df["Spending"])

## Split data training 70 % and testing 30%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

```
In [11]: ## Standardize dataset  
## Create new scaling object  
  
sc = StandardScaler()  
  
## Standardize the training data  
X_train_standard = sc.fit_transform(X_train)  
y_train_standard = sc.fit_transform(y_train.reshape(len(y_train),1))[:,0]
```

```
In [21]: ## Build a new model using Keras, with 22 feature columns  
  
## Rectified Linear Unit -  $f(x) = \max\_value$  for  $x \geq \max\_value$ ,  $f(x) = x$  for  $\text{threshold} \leq x < \max\_value$ ,  $f(x) = \alpha * (x - \text{threshold})$   
## Initializer - he_uniform - It draws samples from a uniform distribution within  $[-\text{limit}, \text{limit}]$  where  $\text{limit}$  is  $\sqrt{6 / \text{fan\_in}}$  where  $\text{fan\_in}$  is the number of input units in the weight tensor.  
model = Sequential()  
model.add(Dense(25, input_dim=22, activation='relu', kernel_initializer='he_uniform'))  
  
## Activation - Linear to perform linear regression  
model.add(Dense(1, activation='linear'))
```

```
In [23]: ## Standardize the testing data separately from the training data  
  
X_test_standard = sc.transform(X_test)  
y_test_standard = sc.transform(y_test.reshape(len(y_test),1))[:,0]
```

```
In [24]: ## Build the optimizer with stochastic gradient descent with a learning rate of 0.01 and a momentum of 0.9  
  
opt = SGD(lr=0.01, momentum=0.9)  
model.compile(loss="mean_squared_logarithmic_error", optimizer=opt, metrics = ["mse"])  
  
## fit model  
history = model.fit(X_train_standard, y_train_standard,  
                    validation_data=(X_test_standard, y_test_standard), epochs=100, verbose=0)
```

```
In [25]: ## Evaluate the model once it has been fit to determine performance  
  
train_mse = model.evaluate(X_train_standard, y_train_standard, verbose=0)  
test_mse = model.evaluate(X_test_standard, y_test_standard, verbose=0)
```

The Mean Squared Error, or MSE, loss is the default loss function to use for most regression problems, and this would normally be the logical choice. It is the preferred loss function under the inference framework of maximum likelihood if the distribution of the target is a Gaussian (normal) distribution.

In this case, it is a regression problem where the target values have a large spread of values, from \$0 to 1500, and when predicting a larger value, we do not want to punish the model as heavily as with MSE.

Instead, we can calculate the natural logarithm of each of the predicted values, then calculate the mean squared error. This is called the Mean Squared Logarithmic Error loss, or MSLE for short.

```
In [26]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV

## Set the figure size
plt.figure(figsize= (15, 10))


## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(history.history['loss'], label = "Training Data Loss")
plt.plot(history.history['val_loss'], label = "Testing Data Loss")

## Give some labels and title
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Title and Legend
plt.title("Model Loss & Mean Squared Logarithmic Error")
plt.legend()

## Show the graph
plt.show()
```

Model Loss & Mean Squared Logarithmic Error





**I created a line plot to show the Mean Squared Logarithmic Error (MSLE) loss over the training cycles on both the training data (blue line) and testing data (orange line).**

**We can see that the model learned reasonably quickly (after about 15 epochs) and test performance remained equivalent/steady from the moment the model was created because of the loss function being used.**

**The performance and convergence behavior of the model suggests that MSLE is a good match for a neural network model attempting to learn from this problem.**

**I will also want to see the Mean Squared Error (MSE) performance of the model, because the logarithmic loss function puts all of our "Purchases" on an exponential scale (the power of 10). As discovered earlier, most of the purchases are for less than \$20.00, so we have to make sure the model can handle identifying smaller amounts.**

In [28]: `from math import sqrt`

```
### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV

## Set the figure size
plt.figure(figsize= (15, 10))

## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(history.history['mse'], label = "Training Data Loss")
plt.plot(history.history['val_mse'], label = "Testing Data Loss")


## Add more line for plotting the RMSE of testing data
plt.plot([sqrt(i) for i in history.history['val_mse']], label = "Testing Data Loss (Using RMSE)")

## Give some labels and title
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Title and Legend
plt.title("Model Loss & Mean Squared Error")
plt.legend()

## Show the graph
plt.show()
```

Model Loss & Mean Squared Error



This plot shows the Mean Squared Error loss over the training cycles for training data (blue line) and testing data (orange line); I also created a line to show the Root Mean Squared Error (RMSE) for testing data to really show the model performance.

It appears that MSE may be showing signs of over-fitting the problem. When I look at the RMSE, it is much more stable and doesn't fluctuate quite as much when going through training cycles. While I am confident that this model is performing much better compared to my earlier attempts with the "out-of-box" regressors from sklearn, I will run a similar process on the data but remove any rows that did not have a purchase to evaluate the model performance using this method.

(b) (20 points) As a variation on this exercise, create a separate “restricted” dataset (i.e., a subset of the original dataset), which includes only purchase records (i.e., where Purchase = 1). Build numeric prediction models to predict Spending for this restricted dataset. All the same requirements as for task (a) apply.

```
In [29]: ## Create a subset of the data, so that we only try to predict the Spending - cool!

## Subset the dataframe created earlier to only include rows with Spending information
purchase_only_df = purchases_df[purchases_df.Purchase != 0]
```

```
In [30]: ## Look at the first five observations

purchase_only_df.head()
```

Out[30]:

	sequence_number	US	source_a	source_c	source_b	source_d	source_e	source_m	source_o	source_h	...	source_x	source_w	Freq	las
0	1	1	0	0	1	0	0	0	0	0	...	0	0	2	
2	3	1	0	0	0	0	0	0	0	0	...	0	0	2	
8	9	1	1	0	0	0	0	0	0	0	...	0	0	4	
9	10	1	1	0	0	0	0	0	0	0	...	0	0	1	
13	14	1	1	0	0	0	0	0	0	0	...	0	0	5	

5 rows × 25 columns

```
In [31]: ## Confirm that no records without a purchase made it in
```

```
purchase_only_df.groupby("Purchase")["US"].count()
```

```
Out[31]: Purchase
1      1000
Name: US, dtype: int64
```

```
In [33]: ## Create a new "X" variable that contains all the features
```

```
X = np.array(purchase_only_df[feature_cols])
```

```
## Create our "y" variable which is our target variable and remove the floating point decimals
```

```
y = np.array(purchase_only_df["Spending"])
```

```
## Split data training 70 % and testing 30%
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

```
In [34]: ## Standardize training features and target variable
```

```
sc = StandardScaler()
```

```
X_train_standard = sc.fit_transform(X_train)
```

```
y_train_standard = sc.fit_transform(y_train.reshape(len(y_train),1))[:,0]
```

```
In [35]: ## Standardize testing features and target variable separately so that the training data doesn't skew or affect the model performance
```

```
## This is known as DATA LEAKAGE.
```

```
X_test_standard = sc.transform(X_test)
```

```
y_test_standard = sc.transform(y_test.reshape(len(y_test),1))[:,0]
```

## C. Use helpful Keras & Keras Tuner libraries to help us perform model evaluation at the same time we develop new classifiers

1. Build the model function
2. Build the hyper tuner function
3. Tune the model and explore the data space to make predictions
4. Assess the model performance.

**We can use some specific custom functions which are explained below.**

```
In [ ]: ### Import some libraries to help us build and tune the model at the same time
```

```
from tensorflow import keras  
from kerastuner.tuners import RandomSearch  
from tensorflow.keras import layers
```

```
In [45]: ## As a first step, we build a function to actually put together our model
```

```
def build_model(hp):  
  
    ## Base Layer  
    model = Sequential()  
  
    ## Add first layer, test with 10 up to 22 features  
    model.add(layers.Dense(units = hp.Int ("units",  
                                         min_value = 10,  
                                         max_value = 22,  
                                         step = 2),  
                      ## Use same initializer as model above  
                      activation = "relu", kernel_initializer = "he_uniform"))  
  
    ## Add target layer using linear regressor  
    model.add(layers.Dense(1, activation = "linear"))  
  
    ## Use different Learning rates to test the model  
    model.compile(  
        optimizer = keras.optimizers.SGD(  
            hp.Choice("learning_rate", values = [1e-2, 1e-3, 1e-4])),  
  
    ## Use same loss function as before  
    loss = "mean_squared_logarithmic_error",  
  
    ## Test using MSE  
  
    metrics = ["mse"])  
  
    ## Return completed model  
    return model
```

In [46]: *## Second step, we build a "hypertuner", this is what will actually fine tune our model as it is being created*

*### Use RandomSearch*

```
tuner = RandomSearch(  
  
    ## Use model function from above  
    build_model,  
  
    ## What is objective function? Using Loss here  
    objective = "val_loss",  
  
    ## Set number of trials  
    max_trials = 5,  
  
    ## Number of executions  
    executions_per_trial = 3,  
  
    ## Set to short dir path  
    directory = "C:\\"  
  
)
```

In [47]: *## Run our search with Keras Tuner!*

```
tuner.search(X_train_standard, y_train_standard, epochs = 10,  
            validation_data = (X_test_standard, y_test_standard))
```



Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 9s - loss: 0.2278 - mse: 2.077 - ETA: 0s - loss: 0.1906 - mse: 1.850 -  
1s 1ms/sample - loss: 0.1897 - mse: 1.8145 - val\_loss: 0.2142 - val\_mse: 39.1103

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.3552 - mse: 3.114 - 0s 113us/sample - loss: 0.1753 - mse:  
1.8438 - val\_loss: 0.2142 - val\_mse: 44.3810

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.1800 - mse: 1.620 - 0s 125us/sample - loss: 0.1688 - mse:  
1.8697 - val\_loss: 0.2142 - val\_mse: 47.8374

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.0647 - mse: 1.080 - ETA: 0s - loss: 0.1623 - mse: 1.822 -  
0s 140us/sample - loss: 0.1655 - mse: 1.9067 - val\_loss: 0.2142 - val\_mse: 49.9367

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.1368 - mse: 2.110 - ETA: 0s - loss: 0.1617 - mse: 1.921 -  
0s 140us/sample - loss: 0.1632 - mse: 1.9334 - val\_loss: 0.2142 - val\_mse: 51.1677

Epoch 6/10

700/700 [=====] - ETA: 0s - loss: 0.1585 - mse: 1.573 - 0s 125us/sample - loss: 0.1612 - mse:  
1.8967 - val\_loss: 0.2142 - val\_mse: 52.5365

Epoch 7/10

700/700 [=====] - ETA: 0s - loss: 0.1937 - mse: 1.670 - ETA: 0s - loss: 0.1383 - mse: 1.852 -  
0s 145us/sample - loss: 0.1596 - mse: 1.9086 - val\_loss: 0.2142 - val\_mse: 53.2560

Epoch 8/10

700/700 [=====] - ETA: 0s - loss: 0.1789 - mse: 1.591 - ETA: 0s - loss: 0.1681 - mse: 1.969 -  
0s 136us/sample - loss: 0.1580 - mse: 1.8965 - val\_loss: 0.2142 - val\_mse: 53.7244

Epoch 9/10

700/700 [=====] - ETA: 0s - loss: 0.1993 - mse: 1.889 - ETA: 0s - loss: 0.1622 - mse: 1.918 -  
0s 139us/sample - loss: 0.1564 - mse: 1.8611 - val\_loss: 0.2142 - val\_mse: 53.9843

Epoch 10/10

700/700 [=====] - ETA: 0s - loss: 0.1509 - mse: 1.541 - ETA: 0s - loss: 0.1598 - mse: 1.886 -  
0s 134us/sample - loss: 0.1547 - mse: 1.8316 - val\_loss: 0.2142 - val\_mse: 54.2806

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 12s - loss: 0.3805 - mse: 2.29 - 1s 1ms/sample - loss: 0.4098 - mse:  
2.6610 - val\_loss: 0.2142 - val\_mse: 8.6957

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.2388 - mse: 1.723 - ETA: 0s - loss: 0.2865 - mse: 2.146 -  
0s 135us/sample - loss: 0.2780 - mse: 2.0689 - val\_loss: 0.2142 - val\_mse: 14.1440

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.2782 - mse: 2.654 - ETA: 0s - loss: 0.1944 - mse: 1.804 -  
0s 183us/sample - loss: 0.2025 - mse: 1.9278 - val\_loss: 0.2142 - val\_mse: 18.3506

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.1523 - mse: 1.465 - ETA: 0s - loss: 0.1792 - mse: 2.003 -  
0s 139us/sample - loss: 0.1713 - mse: 1.9377 - val\_loss: 0.2142 - val\_mse: 20.9990

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.0361 - mse: 0.891 - 0s 107us/sample - loss: 0.1598 - mse:  
1.9953 - val\_loss: 0.2142 - val\_mse: 22.6407

Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.2181 - mse: 2.926 - 0s 131us/sample - loss: 0.1556 - mse: 2.0410 - val\_loss: 0.2142 - val\_mse: 23.6571  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.2695 - mse: 4.305 - ETA: 0s - loss: 0.1439 - mse: 1.961 - 0s 152us/sample - loss: 0.1539 - mse: 2.0656 - val\_loss: 0.2142 - val\_mse: 24.5813  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.0666 - mse: 1.752 - ETA: 0s - loss: 0.1495 - mse: 2.102 - 0s 141us/sample - loss: 0.1524 - mse: 2.0858 - val\_loss: 0.2142 - val\_mse: 25.8277  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.0622 - mse: 1.494 - 0s 107us/sample - loss: 0.1515 - mse: 2.1006 - val\_loss: 0.2142 - val\_mse: 26.6425  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.2446 - mse: 3.798 - 0s 124us/sample - loss: 0.1509 - mse: 2.1524 - val\_loss: 0.2142 - val\_mse: 27.1363  
Train on 700 samples, validate on 300 samples  
Epoch 1/10  
700/700 [=====] - ETA: 14s - loss: 0.0780 - mse: 2.32 - 1s 1ms/sample - loss: 0.2018 - mse: 2.5607 - val\_loss: 1.1964 - val\_mse: 8.1160  
Epoch 2/10  
700/700 [=====] - ETA: 0s - loss: 0.1176 - mse: 2.172 - ETA: 0s - loss: 0.1897 - mse: 2.715 - 0s 162us/sample - loss: 0.1793 - mse: 2.6575 - val\_loss: 0.8969 - val\_mse: 5.4383  
Epoch 3/10  
700/700 [=====] - ETA: 0s - loss: 0.0769 - mse: 0.888 - ETA: 0s - loss: 0.1604 - mse: 2.538 - 0s 143us/sample - loss: 0.1674 - mse: 2.7262 - val\_loss: 0.6836 - val\_mse: 3.9247  
Epoch 4/10  
700/700 [=====] - ETA: 0s - loss: 0.2739 - mse: 3.457 - ETA: 0s - loss: 0.1662 - mse: 2.968 - 0s 158us/sample - loss: 0.1604 - mse: 2.8242 - val\_loss: 0.5364 - val\_mse: 3.0590  
Epoch 5/10  
700/700 [=====] - ETA: 0s - loss: 0.1883 - mse: 2.779 - ETA: 0s - loss: 0.1427 - mse: 2.591 - 0s 171us/sample - loss: 0.1560 - mse: 2.8656 - val\_loss: 0.4387 - val\_mse: 2.5670  
Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.0993 - mse: 2.312 - ETA: 0s - loss: 0.1529 - mse: 2.974 - 0s 325us/sample - loss: 0.1531 - mse: 2.9327 - val\_loss: 0.3717 - val\_mse: 2.2779  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.1290 - mse: 2.855 - ETA: 0s - loss: 0.1408 - mse: 2.685 - 0s 153us/sample - loss: 0.1509 - mse: 2.9367 - val\_loss: 0.3217 - val\_mse: 2.0908  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.0923 - mse: 1.974 - ETA: 0s - loss: 0.1468 - mse: 2.868 - 0s 134us/sample - loss: 0.1492 - mse: 2.9831 - val\_loss: 0.2911 - val\_mse: 1.9873  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1389 - mse: 3.140 - ETA: 0s - loss: 0.1449 - mse: 2.891 - 0s 182us/sample - loss: 0.1473 - mse: 2.9709 - val\_loss: 0.2718 - val\_mse: 1.9185  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.0388 - mse: 1.840 - ETA: 0s - loss: 0.1420 - mse: 2.870 - 0s 140us/sample - loss: 0.1456 - mse: 2.9571 - val\_loss: 0.2577 - val\_mse: 1.8692

**Trial complete**

**Trial summary**

**Hp values:**

|learning\_rate: 0.01

|units: 18

|Score: 0.22868611564238864

|Best step: 0

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 8s - loss: 0.4462 - mse: 3.820 - 1s 903us/sample - loss: 0.3916 - mse: 2.8532 - val\_loss: 0.2239 - val\_mse: 5.1755

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.2346 - mse: 2.658 - ETA: 0s - loss: 0.2996 - mse: 2.604 - 0s 281us/sample - loss: 0.2902 - mse: 2.4810 - val\_loss: 0.2142 - val\_mse: 7.8681

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.1844 - mse: 1.773 - ETA: 0s - loss: 0.2344 - mse: 2.350 - 0s 170us/sample - loss: 0.2260 - mse: 2.3335 - val\_loss: 0.2142 - val\_mse: 11.0840

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.2982 - mse: 3.190 - ETA: 0s - loss: 0.1869 - mse: 2.315 - 0s 159us/sample - loss: 0.1853 - mse: 2.2941 - val\_loss: 0.2142 - val\_mse: 14.1641

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.1275 - mse: 2.577 - ETA: 0s - loss: 0.1654 - mse: 2.285 - 0s 150us/sample - loss: 0.1606 - mse: 2.3215 - val\_loss: 0.2142 - val\_mse: 16.5521

Epoch 6/10

700/700 [=====] - ETA: 0s - loss: 0.2502 - mse: 2.687 - ETA: 0s - loss: 0.1418 - mse: 2.217 - 0s 134us/sample - loss: 0.1458 - mse: 2.3336 - val\_loss: 0.2142 - val\_mse: 18.3576

Epoch 7/10

700/700 [=====] - ETA: 0s - loss: 0.2216 - mse: 2.749 - ETA: 0s - loss: 0.1338 - mse: 2.342 - 0s 140us/sample - loss: 0.1364 - mse: 2.3676 - val\_loss: 0.2142 - val\_mse: 19.6279

Epoch 8/10

700/700 [=====] - ETA: 0s - loss: 0.0993 - mse: 1.318 - 0s 103us/sample - loss: 0.1302 - mse: 2.3859 - val\_loss: 0.2142 - val\_mse: 20.4884

Epoch 9/10

700/700 [=====] - ETA: 0s - loss: 0.1309 - mse: 3.163 - ETA: 0s - loss: 0.1287 - mse: 2.395 - 0s 125us/sample - loss: 0.1258 - mse: 2.3845 - val\_loss: 0.2142 - val\_mse: 21.1285

Epoch 10/10

700/700 [=====] - ETA: 0s - loss: 0.0708 - mse: 1.045 - 0s 141us/sample - loss: 0.1225 - mse: 2.3942 - val\_loss: 0.2142 - val\_mse: 21.5549

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 11s - loss: 0.1152 - mse: 1.56 - 1s 1ms/sample - loss: 0.1340 - mse: 1.4337 - val\_loss: 0.2387 - val\_mse: 15.1861

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.2146 - mse: 2.789 - 0s 105us/sample - loss: 0.1232 - mse: 1.4449 - val\_loss: 0.2343 - val\_mse: 16.7029

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.1791 - mse: 1.293 - 0s 97us/sample - loss: 0.1168 - mse: 1.4632 - val\_loss: 0.2313 - val\_mse: 17.6525

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.0625 - mse: 1.063 - 0s 98us/sample - loss: 0.1125 - mse: 1.4683 - val\_loss: 0.2293 - val\_mse: 18.5404

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.1017 - mse: 1.157 - 0s 104us/sample - loss: 0.1095 - mse: 1.4635 - val\_loss: 0.2278 - val\_mse: 19.1628

Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.1535 - mse: 1.116 - 0s 114us/sample - loss: 0.1073 - mse: 1.4673 - val\_loss: 0.2267 - val\_mse: 19.6844  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.1605 - mse: 1.796 - 0s 105us/sample - loss: 0.1054 - mse: 1.4721 - val\_loss: 0.2257 - val\_mse: 19.9802  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.1295 - mse: 2.297 - 0s 118us/sample - loss: 0.1039 - mse: 1.4636 - val\_loss: 0.2248 - val\_mse: 20.2435  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1547 - mse: 3.121 - 0s 109us/sample - loss: 0.1026 - mse: 1.4596 - val\_loss: 0.2241 - val\_mse: 20.3862  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.0720 - mse: 1.528 - ETA: 0s - loss: 0.0999 - mse: 1.446 - 0s 136us/sample - loss: 0.1015 - mse: 1.4560 - val\_loss: 0.2235 - val\_mse: 20.4760  
Train on 700 samples, validate on 300 samples  
Epoch 1/10  
700/700 [=====] - ETA: 15s - loss: 0.1417 - mse: 2.22 - 1s 1ms/sample - loss: 0.1631 - mse: 2.6476 - val\_loss: 0.2142 - val\_mse: 17.4176  
Epoch 2/10  
700/700 [=====] - ETA: 0s - loss: 0.3243 - mse: 4.395 - 0s 105us/sample - loss: 0.1613 - mse: 2.6647 - val\_loss: 0.2142 - val\_mse: 17.7705  
Epoch 3/10  
700/700 [=====] - ETA: 0s - loss: 0.1627 - mse: 2.931 - 0s 90us/sample - loss: 0.1600 - mse: 2.6685 - val\_loss: 0.2142 - val\_mse: 18.0416  
Epoch 4/10  
700/700 [=====] - ETA: 0s - loss: 0.2778 - mse: 3.680 - 0s 93us/sample - loss: 0.1589 - mse: 2.6669 - val\_loss: 0.2142 - val\_mse: 18.2795  
Epoch 5/10  
700/700 [=====] - ETA: 0s - loss: 0.0927 - mse: 2.204 - 0s 94us/sample - loss: 0.1580 - mse: 2.6695 - val\_loss: 0.2142 - val\_mse: 18.4466  
Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.3625 - mse: 5.350 - 0s 107us/sample - loss: 0.1571 - mse: 2.6534 - val\_loss: 0.2142 - val\_mse: 18.6982  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.0545 - mse: 1.385 - 0s 100us/sample - loss: 0.1562 - mse: 2.6450 - val\_loss: 0.2142 - val\_mse: 18.8721  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.0135 - mse: 1.371 - 0s 94us/sample - loss: 0.1553 - mse: 2.6324 - val\_loss: 0.2142 - val\_mse: 19.0736  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.2039 - mse: 2.646 - 0s 90us/sample - loss: 0.1543 - mse: 2.6133 - val\_loss: 0.2142 - val\_mse: 19.2449  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.2501 - mse: 3.994 - 0s 93us/sample - loss: 0.1535 - mse: 2.5933 - val\_loss: 0.2142 - val\_mse: 19.3571

**Trial complete**

**Trial summary**

**Hp values:**

|learning\_rate: 0.01

|units: 12

|Score: 0.21729836497041913

|Best step: 0

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 10s - loss: 0.2651 - mse: 2.17 - 1s 1ms/sample - loss: 0.3101 - mse: 2.6623 - val\_loss: 2.4770 - val\_mse: 24.6645

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.3102 - mse: 2.894 - 0s 140us/sample - loss: 0.3027 - mse: 2.6381 - val\_loss: 2.3992 - val\_mse: 23.1556

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.2140 - mse: 1.424 - 0s 150us/sample - loss: 0.2961 - mse: 2.6197 - val\_loss: 2.3233 - val\_mse: 21.7476

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.2880 - mse: 2.605 - ETA: 0s - loss: 0.2770 - mse: 2.468 - 0s 207us/sample - loss: 0.2898 - mse: 2.6023 - val\_loss: 2.2501 - val\_mse: 20.4467

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.2619 - mse: 2.675 - ETA: 0s - loss: 0.2812 - mse: 2.608 - 0s 201us/sample - loss: 0.2831 - mse: 2.5887 - val\_loss: 2.1737 - val\_mse: 19.1451

Epoch 6/10

700/700 [=====] - ETA: 0s - loss: 0.2500 - mse: 2.177 - ETA: 0s - loss: 0.2488 - mse: 2.231 - 0s 241us/sample - loss: 0.2768 - mse: 2.5788 - val\_loss: 2.1012 - val\_mse: 17.9638

Epoch 7/10

700/700 [=====] - ETA: 0s - loss: 0.3181 - mse: 2.400 - ETA: 0s - loss: 0.2690 - mse: 2.499 - 0s 198us/sample - loss: 0.2711 - mse: 2.5705 - val\_loss: 2.0307 - val\_mse: 16.8641

Epoch 8/10

700/700 [=====] - ETA: 0s - loss: 0.4406 - mse: 3.150 - 0s 161us/sample - loss: 0.2657 - mse: 2.5655 - val\_loss: 1.9604 - val\_mse: 15.8139

Epoch 9/10

700/700 [=====] - ETA: 0s - loss: 0.3042 - mse: 3.496 - 0s 143us/sample - loss: 0.2606 - mse: 2.5598 - val\_loss: 1.8931 - val\_mse: 14.8520

Epoch 10/10

700/700 [=====] - ETA: 0s - loss: 0.2542 - mse: 1.847 - 0s 145us/sample - loss: 0.2559 - mse: 2.5572 - val\_loss: 1.8284 - val\_mse: 13.9639

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 9s - loss: 0.2335 - mse: 1.488 - ETA: 0s - loss: 0.3022 - mse: 2.306 - 1s 1ms/sample - loss: 0.3024 - mse: 2.3003 - val\_loss: 0.8833 - val\_mse: 4.6209

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.2335 - mse: 1.491 - ETA: 0s - loss: 0.2867 - mse: 2.195 - 0s 184us/sample - loss: 0.2938 - mse: 2.2656 - val\_loss: 0.8117 - val\_mse: 4.1825

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.3157 - mse: 2.691 - 0s 152us/sample - loss: 0.2858 - mse: 2.2319 - val\_loss: 0.7439 - val\_mse: 3.7937

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.2200 - mse: 1.295 - 0s 146us/sample - loss: 0.2784 - mse: 2.2038 - val\_loss: 0.6813 - val\_mse: 3.4569

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.2627 - mse: 2.647 - ETA: 0s - loss: 0.2775 - mse: 2.209 - 0s 191us/sample - loss: 0.2717 - mse: 2.1754 - val\_loss: 0.6233 - val\_mse: 3.1624

Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.2704 - mse: 2.276 - ETA: 0s - loss: 0.2544 - mse: 2.095 -  
0s 195us/sample - loss: 0.2654 - mse: 2.1505 - val\_loss: 0.5678 - val\_mse: 2.8963  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.2269 - mse: 1.175 - 0s 141us/sample - loss: 0.2595 - mse:  
2.1287 - val\_loss: 0.5169 - val\_mse: 2.6648  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.1824 - mse: 1.251 - 0s 145us/sample - loss: 0.2540 - mse:  
2.1104 - val\_loss: 0.4699 - val\_mse: 2.4613  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1754 - mse: 1.726 - ETA: 0s - loss: 0.2444 - mse: 2.034 -  
0s 172us/sample - loss: 0.2489 - mse: 2.0906 - val\_loss: 0.4271 - val\_mse: 2.2836  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.2354 - mse: 2.224 - 0s 141us/sample - loss: 0.2442 - mse:  
2.0748 - val\_loss: 0.3890 - val\_mse: 2.1304  
Train on 700 samples, validate on 300 samples  
Epoch 1/10  
700/700 [=====] - ETA: 9s - loss: 0.3172 - mse: 1.763 - 1s 1ms/sample - loss: 0.2506 - mse:  
1.5827 - val\_loss: 0.2533 - val\_mse: 6.0947  
Epoch 2/10  
700/700 [=====] - ETA: 0s - loss: 0.2711 - mse: 1.673 - 0s 154us/sample - loss: 0.2371 - mse:  
1.5390 - val\_loss: 0.2421 - val\_mse: 6.5102  
Epoch 3/10  
700/700 [=====] - ETA: 0s - loss: 0.1921 - mse: 1.195 - 0s 187us/sample - loss: 0.2248 - mse:  
1.5018 - val\_loss: 0.2326 - val\_mse: 6.9277  
Epoch 4/10  
700/700 [=====] - ETA: 0s - loss: 0.2065 - mse: 1.367 - 0s 154us/sample - loss: 0.2139 - mse:  
1.4674 - val\_loss: 0.2246 - val\_mse: 7.3356  
Epoch 5/10  
700/700 [=====] - ETA: 0s - loss: 0.1593 - mse: 0.970 - 0s 147us/sample - loss: 0.2044 - mse:  
1.4408 - val\_loss: 0.2175 - val\_mse: 7.7353  
Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.2328 - mse: 1.282 - ETA: 0s - loss: 0.1924 - mse: 1.365 -  
0s 187us/sample - loss: 0.1959 - mse: 1.4196 - val\_loss: 0.2115 - val\_mse: 8.1215  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.1490 - mse: 0.973 - 0s 154us/sample - loss: 0.1883 - mse:  
1.4008 - val\_loss: 0.2061 - val\_mse: 8.5106  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.1552 - mse: 0.876 - 0s 149us/sample - loss: 0.1815 - mse:  
1.3857 - val\_loss: 0.2014 - val\_mse: 8.8968  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1705 - mse: 1.201 - ETA: 0s - loss: 0.1735 - mse: 1.371 -  
0s 219us/sample - loss: 0.1753 - mse: 1.3735 - val\_loss: 0.1974 - val\_mse: 9.2647  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.2586 - mse: 1.666 - ETA: 0s - loss: 0.1719 - mse: 1.296 -  
0s 189us/sample - loss: 0.1699 - mse: 1.3649 - val\_loss: 0.1940 - val\_mse: 9.6145



**Trial complete**

**Trial summary**

**Hp values:**

|learning\_rate: 0.001

|units: 20

|Score: 0.8037920699516933

|Best step: 0

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 10s - loss: 0.1684 - mse: 1.06 - ETA: 0s - loss: 0.2880 - mse: 1.8078  
- 1s 1ms/sample - loss: 0.2879 - mse: 1.8118 - val\_loss: 0.2142 - val\_mse: 5.8789

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.2582 - mse: 1.399 - ETA: 0s - loss: 0.2858 - mse: 1.818 -  
0s 166us/sample - loss: 0.2772 - mse: 1.7584 - val\_loss: 0.2142 - val\_mse: 6.6973

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.2301 - mse: 1.589 - ETA: 0s - loss: 0.2728 - mse: 1.727 -  
0s 158us/sample - loss: 0.2672 - mse: 1.7110 - val\_loss: 0.2142 - val\_mse: 7.5712

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.1561 - mse: 0.957 - ETA: 0s - loss: 0.2541 - mse: 1.594 -  
0s 160us/sample - loss: 0.2579 - mse: 1.6655 - val\_loss: 0.2142 - val\_mse: 8.5069

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.2213 - mse: 1.196 - 0s 107us/sample - loss: 0.2491 - mse:  
1.6255 - val\_loss: 0.2142 - val\_mse: 9.4815

Epoch 6/10

700/700 [=====] - ETA: 0s - loss: 0.3181 - mse: 2.045 - 0s 104us/sample - loss: 0.2409 - mse:  
1.5899 - val\_loss: 0.2142 - val\_mse: 10.4935

Epoch 7/10

700/700 [=====] - ETA: 0s - loss: 0.2379 - mse: 1.709 - 0s 125us/sample - loss: 0.2332 - mse:  
1.5573 - val\_loss: 0.2142 - val\_mse: 11.5392

Epoch 8/10

700/700 [=====] - ETA: 0s - loss: 0.2155 - mse: 1.319 - 0s 164us/sample - loss: 0.2259 - mse:  
1.5280 - val\_loss: 0.2142 - val\_mse: 12.6076

Epoch 9/10

700/700 [=====] - ETA: 0s - loss: 0.2162 - mse: 1.915 - ETA: 0s - loss: 0.2185 - mse: 1.522 -  
0s 144us/sample - loss: 0.2192 - mse: 1.5015 - val\_loss: 0.2142 - val\_mse: 13.7033

Epoch 10/10

700/700 [=====] - ETA: 0s - loss: 0.2722 - mse: 1.578 - ETA: 0s - loss: 0.2209 - mse: 1.533 -  
0s 153us/sample - loss: 0.2128 - mse: 1.4781 - val\_loss: 0.2142 - val\_mse: 14.8050

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 12s - loss: 0.5751 - mse: 4.01 - 1s 1ms/sample - loss: 0.5541 - mse:  
3.5412 - val\_loss: 4.1518 - val\_mse: 94.2987

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.1860 - mse: 0.901 - 0s 104us/sample - loss: 0.5354 - mse:  
3.4029 - val\_loss: 4.0642 - val\_mse: 89.6174

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.2453 - mse: 1.089 - 0s 123us/sample - loss: 0.5175 - mse:  
3.2755 - val\_loss: 3.9774 - val\_mse: 85.1492

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.6303 - mse: 4.502 - 0s 103us/sample - loss: 0.5005 - mse:  
3.1563 - val\_loss: 3.8916 - val\_mse: 80.8862

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.6183 - mse: 4.296 - 0s 110us/sample - loss: 0.4840 - mse:  
3.0440 - val\_loss: 3.8059 - val\_mse: 76.7925

Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.5311 - mse: 2.882 - ETA: 0s - loss: 0.4763 - mse: 3.000 -  
0s 133us/sample - loss: 0.4682 - mse: 2.9377 - val\_loss: 3.7216 - val\_mse: 72.9088  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.5509 - mse: 3.114 - 0s 116us/sample - loss: 0.4530 - mse:  
2.8379 - val\_loss: 3.6380 - val\_mse: 69.2117  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.4358 - mse: 2.661 - 0s 103us/sample - loss: 0.4384 - mse:  
2.7452 - val\_loss: 3.5556 - val\_mse: 65.7013  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.3149 - mse: 1.724 - 0s 121us/sample - loss: 0.4245 - mse:  
2.6576 - val\_loss: 3.4744 - val\_mse: 62.3672  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.4545 - mse: 2.781 - ETA: 0s - loss: 0.4116 - mse: 2.574 -  
0s 140us/sample - loss: 0.4111 - mse: 2.5775 - val\_loss: 3.3946 - val\_mse: 59.2071  
Train on 700 samples, validate on 300 samples  
Epoch 1/10  
700/700 [=====] - ETA: 15s - loss: 0.3293 - mse: 1.93 - ETA: 0s - loss: 0.2918 - mse: 1.6378  
- 1s 1ms/sample - loss: 0.2930 - mse: 1.6461 - val\_loss: 1.7003 - val\_mse: 14.0277  
Epoch 2/10  
700/700 [=====] - ETA: 0s - loss: 0.2150 - mse: 0.885 - ETA: 0s - loss: 0.2724 - mse: 1.552 -  
0s 156us/sample - loss: 0.2741 - mse: 1.5723 - val\_loss: 1.6080 - val\_mse: 12.7749  
Epoch 3/10  
700/700 [=====] - ETA: 0s - loss: 0.2257 - mse: 1.257 - ETA: 0s - loss: 0.2387 - mse: 1.373 -  
0s 152us/sample - loss: 0.2570 - mse: 1.5053 - val\_loss: 1.5194 - val\_mse: 11.6409  
Epoch 4/10  
700/700 [=====] - ETA: 0s - loss: 0.2415 - mse: 1.205 - ETA: 0s - loss: 0.2423 - mse: 1.469 -  
0s 144us/sample - loss: 0.2419 - mse: 1.4473 - val\_loss: 1.4355 - val\_mse: 10.6255  
Epoch 5/10  
700/700 [=====] - ETA: 0s - loss: 0.1693 - mse: 0.830 - 0s 125us/sample - loss: 0.2285 - mse:  
1.3974 - val\_loss: 1.3558 - val\_mse: 9.7135  
Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.2309 - mse: 1.129 - ETA: 0s - loss: 0.2160 - mse: 1.344 -  
0s 145us/sample - loss: 0.2167 - mse: 1.3524 - val\_loss: 1.2805 - val\_mse: 8.8974  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.1792 - mse: 1.061 - ETA: 0s - loss: 0.2046 - mse: 1.287 -  
0s 152us/sample - loss: 0.2064 - mse: 1.3129 - val\_loss: 1.2102 - val\_mse: 8.1732  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.1709 - mse: 1.278 - 0s 160us/sample - loss: 0.1976 - mse:  
1.2794 - val\_loss: 1.1447 - val\_mse: 7.5321  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1772 - mse: 1.384 - 0s 125us/sample - loss: 0.1899 - mse:  
1.2504 - val\_loss: 1.0833 - val\_mse: 6.9587  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.1624 - mse: 0.784 - 0s 119us/sample - loss: 0.1833 - mse:  
1.2242 - val\_loss: 1.0266 - val\_mse: 6.4530

**Trial complete**

**Trial summary**

**Hp values:**

|learning\_rate: 0.001

|units: 14

|Score: 1.545140876173973

|Best step: 0

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 9s - loss: 0.5462 - mse: 3.590 - 1s 913us/sample - loss: 0.3963 - mse: 2.7392 - val\_loss: 1.2289 - val\_mse: 7.7978

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.3283 - mse: 2.158 - ETA: 0s - loss: 0.3920 - mse: 2.714 - 0s 185us/sample - loss: 0.3948 - mse: 2.7304 - val\_loss: 1.2232 - val\_mse: 7.7457

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.2493 - mse: 1.629 - 0s 144us/sample - loss: 0.3933 - mse: 2.7214 - val\_loss: 1.2175 - val\_mse: 7.6942

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.3017 - mse: 1.837 - 0s 139us/sample - loss: 0.3918 - mse: 2.7125 - val\_loss: 1.2119 - val\_mse: 7.6433

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.1810 - mse: 1.235 - 0s 148us/sample - loss: 0.3903 - mse: 2.7038 - val\_loss: 1.2063 - val\_mse: 7.5928

Epoch 6/10

700/700 [=====] - ETA: 0s - loss: 0.3773 - mse: 2.363 - ETA: 0s - loss: 0.3683 - mse: 2.512 - 0s 188us/sample - loss: 0.3888 - mse: 2.6953 - val\_loss: 1.2006 - val\_mse: 7.5426

Epoch 7/10

700/700 [=====] - ETA: 0s - loss: 0.5643 - mse: 4.109 - 0s 160us/sample - loss: 0.3874 - mse: 2.6869 - val\_loss: 1.1951 - val\_mse: 7.4932

Epoch 8/10

700/700 [=====] - ETA: 0s - loss: 0.2708 - mse: 1.613 - 0s 157us/sample - loss: 0.3859 - mse: 2.6784 - val\_loss: 1.1896 - val\_mse: 7.4443

Epoch 9/10

700/700 [=====] - ETA: 0s - loss: 0.3981 - mse: 2.541 - ETA: 0s - loss: 0.3962 - mse: 2.725 - 0s 186us/sample - loss: 0.3845 - mse: 2.6700 - val\_loss: 1.1841 - val\_mse: 7.3958

Epoch 10/10

700/700 [=====] - ETA: 0s - loss: 0.6233 - mse: 4.472 - 0s 135us/sample - loss: 0.3830 - mse: 2.6618 - val\_loss: 1.1786 - val\_mse: 7.3478

Train on 700 samples, validate on 300 samples

Epoch 1/10

700/700 [=====] - ETA: 10s - loss: 0.1572 - mse: 2.94 - ETA: 0s - loss: 0.1636 - mse: 3.8171 - 1s 1ms/sample - loss: 0.1583 - mse: 3.6636 - val\_loss: 0.2071 - val\_mse: 2.2801

Epoch 2/10

700/700 [=====] - ETA: 0s - loss: 0.3037 - mse: 5.404 - 0s 114us/sample - loss: 0.1583 - mse: 3.6639 - val\_loss: 0.2071 - val\_mse: 2.2826

Epoch 3/10

700/700 [=====] - ETA: 0s - loss: 0.1428 - mse: 3.473 - ETA: 0s - loss: 0.1621 - mse: 3.836 - 0s 187us/sample - loss: 0.1582 - mse: 3.6641 - val\_loss: 0.2071 - val\_mse: 2.2850

Epoch 4/10

700/700 [=====] - ETA: 0s - loss: 0.0842 - mse: 2.139 - ETA: 0s - loss: 0.1621 - mse: 3.745 - 0s 128us/sample - loss: 0.1582 - mse: 3.6644 - val\_loss: 0.2072 - val\_mse: 2.2875

Epoch 5/10

700/700 [=====] - ETA: 0s - loss: 0.1412 - mse: 3.576 - 0s 111us/sample - loss: 0.1582 - mse: 3.6647 - val\_loss: 0.2072 - val\_mse: 2.2899

Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.0535 - mse: 1.353 - 0s 115us/sample - loss: 0.1581 - mse: 3.6650 - val\_loss: 0.2072 - val\_mse: 2.2923  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.0713 - mse: 2.274 - ETA: 0s - loss: 0.1558 - mse: 3.729 - 0s 128us/sample - loss: 0.1581 - mse: 3.6653 - val\_loss: 0.2072 - val\_mse: 2.2948  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.0890 - mse: 3.854 - 0s 118us/sample - loss: 0.1581 - mse: 3.6654 - val\_loss: 0.2072 - val\_mse: 2.2972  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.2913 - mse: 4.875 - 0s 115us/sample - loss: 0.1580 - mse: 3.6657 - val\_loss: 0.2073 - val\_mse: 2.2996  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.1627 - mse: 4.327 - 0s 120us/sample - loss: 0.1580 - mse: 3.6660 - val\_loss: 0.2073 - val\_mse: 2.3020  
Train on 700 samples, validate on 300 samples  
Epoch 1/10  
700/700 [=====] - ETA: 10s - loss: 0.0654 - mse: 4.54 - 1s 994us/sample - loss: 0.1929 - mse: 4.9452 - val\_loss: 0.2142 - val\_mse: 7.3170  
Epoch 2/10  
700/700 [=====] - ETA: 0s - loss: 0.2105 - mse: 5.969 - 0s 114us/sample - loss: 0.1928 - mse: 4.9463 - val\_loss: 0.2142 - val\_mse: 7.3343  
Epoch 3/10  
700/700 [=====] - ETA: 0s - loss: 0.1286 - mse: 2.833 - 0s 114us/sample - loss: 0.1927 - mse: 4.9477 - val\_loss: 0.2142 - val\_mse: 7.3515  
Epoch 4/10  
700/700 [=====] - ETA: 0s - loss: 0.1573 - mse: 5.448 - 0s 104us/sample - loss: 0.1926 - mse: 4.9488 - val\_loss: 0.2142 - val\_mse: 7.3688  
Epoch 5/10  
700/700 [=====] - ETA: 0s - loss: 0.2316 - mse: 3.835 - 0s 108us/sample - loss: 0.1925 - mse: 4.9502 - val\_loss: 0.2142 - val\_mse: 7.3861  
Epoch 6/10  
700/700 [=====] - ETA: 0s - loss: 0.1341 - mse: 3.091 - ETA: 0s - loss: 0.1975 - mse: 5.041 - 0s 127us/sample - loss: 0.1924 - mse: 4.9515 - val\_loss: 0.2142 - val\_mse: 7.4033  
Epoch 7/10  
700/700 [=====] - ETA: 0s - loss: 0.2731 - mse: 4.268 - ETA: 0s - loss: 0.1869 - mse: 4.679 - 0s 151us/sample - loss: 0.1923 - mse: 4.9528 - val\_loss: 0.2142 - val\_mse: 7.4205  
Epoch 8/10  
700/700 [=====] - ETA: 0s - loss: 0.1058 - mse: 4.611 - ETA: 0s - loss: 0.2017 - mse: 4.885 - 0s 134us/sample - loss: 0.1922 - mse: 4.9541 - val\_loss: 0.2142 - val\_mse: 7.4376  
Epoch 9/10  
700/700 [=====] - ETA: 0s - loss: 0.1310 - mse: 4.517 - 0s 108us/sample - loss: 0.1922 - mse: 4.9551 - val\_loss: 0.2142 - val\_mse: 7.4548  
Epoch 10/10  
700/700 [=====] - ETA: 0s - loss: 0.1541 - mse: 4.405 - ETA: 0s - loss: 0.1910 - mse: 4.962 - 0s 150us/sample - loss: 0.1921 - mse: 4.9565 - val\_loss: 0.2142 - val\_mse: 7.4718

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.0001

|units: 16

|Score: 0.5333168996042675

|Best step: 0

```
In [48]: ## Print out the summary results of our data mining process using Keras  
  
tuner.results_summary()
```

## Results summary

|Results in C:\untitled\_project

|Showing 10 best trials

|Objective: Objective(name='val\_loss', direction='min') Score: 0.21729836497041913

|Objective: Objective(name='val\_loss', direction='min') Score: 0.22868611564238864

|Objective: Objective(name='val\_loss', direction='min') Score: 0.5333168996042675

|Objective: Objective(name='val\_loss', direction='min') Score: 0.8037920699516933

|Objective: Objective(name='val\_loss', direction='min') Score: 1.545140876173973

```
In [50]: ## Save the best models to evaluate their performance below  
  
models = tuner.get_best_models(num_models = 2)
```

```
In [51]: ## Save the single best model to a new variable  
  
best_model = models[0]
```

In [53]: *### Use our best model and train it again*

```
non_spend_history = best_model.fit(X_train_standard, y_train_standard,  
                                   validation_data = (X_test_standard, y_test_standard), epochs = 50, verbose = 0)
```

In [54]: *## Evaluate the "best" model*

```
train_mse = model.evaluate(X_train_standard, y_train_standard, verbose=0)  
test_mse = model.evaluate(X_test_standard, y_test_standard, verbose=0)
```



```
In [56]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot the MSLE for this model

## Set the figure size
plt.figure(figsize= (15, 10))


## Plot the training data "Loss" compared to the testing data "Loss"
plt.plot(non_spend_history.history['loss'], label = "Training Data Loss")
plt.plot(non_spend_history.history['val_loss'], label = "Testing Data Loss")

## Provide some labels
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Provide a title and Legend
plt.title("Model Loss & Mean Squared Logarithmic Error")
plt.legend()

## Show the graph
plt.show()
```

Model Loss & Mean Squared Logarithmic Error



```
In [58]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV

## Set the figure size
plt.figure(figsize= (15, 10))


## Plot the training data "Loss" compared to the testing data "Loss"
plt.plot(non_spend_history.history['mse'], label = "Training Data Loss")
plt.plot(non_spend_history.history['val_mse'], label = "Testing Data Loss")

## Add more line for plotting the RMSE of testing data
plt.plot([sqrt(i) for i in non_spend_history.history['val_mse']], label = "Testing Data Loss (Using RMSE)")

## Provide some labels
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Provide a title and Legend
plt.title("Model Loss & Mean Squared Error")
plt.legend(loc = "right")

## Show the graph
plt.show()
```



## D. Conclusions / Model Evaluation

From the various models tested, it seems the best predictor of "Spending" is the model that was developed using Keras & Keras Tuner. The reasoning behind this is that the model has less training data to use to learn, since we are removing a lot of the variance in the data by subsetting it to only include Purchases.

In our original explorations, we were using training data that had a lot of 0.00 values in the Spending column (no purchase). This meant that our modeling techniques were struggling with identifying the actual purchasing behavior because there were a lot of false positives. The mining technique would see two sets of features that were similar but one would end with spend and one would not, and do a poor job assigning or estimating the spend of a customer.

The basic idea is that there are two possible (and almost opposite) reasons for a data mining technique to not perform well.

In the first case, we might have a model that is too complicated for the amount of data we have. This situation, known as high variance, leads to model over-fitting. We know that we are facing a high variance problem with this data set, and it is confirmed when we see that the training error is much lower than testing error.

High variance problems can be addressed by reducing the number of features, and... yes, by increasing the number of data points, which we can't do here.

These models have many features, as compared to our training examples, which is why we see the over-fitting in the first attempt of using Keras and Keras Tuner.

However, we might have a model that is too simple to explain the data we have. In this case, which is known as high bias, adding more data will not help.

Here is the specific model selection and model testing steps. The "best\_model" would be declared the winner!

```
: ▼ ## Run our search with Keras Tuner!
  ▼ tuner.search(X_train_standard, y_train_standard, epochs = 10,
    validation_data = (X_test_standard, y_test_standard))
...

: ▼ ## Print out the summary results of our data mining process using Keras
  tuner.results_summary()
...

: ▼ ## Save the best models to evaluate their performance below
  models = tuner.get_best_models(num_models = 2)


: ▼ ## Save the single best model to a new variable
  best_model = models[0]

: ▼ ### Use our best model and train it again
  ▼ non_spend_history = best_model.fit(X_train_standard, y_train_standard,
    validation_data = (X_test_standard, y_test_standard), epochs = 50, verbose = 0)

: ▼ ## Evaluate the "best" model
  train_mse = model.evaluate(X_train_standard, y_train_standard, verbose=0)
  test_mse = model.evaluate(X_test_standard, y_test_standard, verbose=0)
```

I am highlighting two of the plots that were generated, that show a nice smoothing of the training data loss, which meant that the model was learning at a steady rate. This is unlike the first attempt of building a model with Keras, where we saw a lot of uneven learning.


Model Loss & Mean Squared Logarithmic Error





When I look at the RMSE for the preferred regression model, it is much more stable and doesn't fluctuate quite as much when going through training cycles. I am confident that this model is performing much better compared to my earlier attempts with the "out-of-box" regressors from sklearn.

Model Loss & Mean Squared Error



**Quation 2. (50 points)** Download the dataset on spam vs. non-spam emails from the following URL: <http://archive.ics.uci.edu/ml/datasets/Spambase>. Specifically, (i) file “spambase.data” contains the actual data, and (ii) files “spambase.names” and “spambase.DOCUMENTATION” contain the description of the data. This dataset has 4601 records, each record representing a different email message. Each record is described with 58 attributes (indicated in the aforementioned .names file): attributes 1-57 represent various content-based characteristics already extracted from each email message (related to the frequency of certain words or certain punctuation symbols in a message as well as to the usage of capital letters in a message), and the last attribute represents the class label for each message (spam or non-spam).

---

SPAM E-MAIL DATABASE ATTRIBUTES (in .names format)

48 continuous real [0,100] attributes of type word\_freq\_WORD  
= percentage of words in the e-mail that match WORD,  
i.e.  $100 * (\text{number of times the WORD appears in the e-mail}) /$   
total number of words in e-mail. A “word” in this case is any  
string of alphanumeric characters bounded by non-alphanumeric  
characters or end-of-string.

6 continuous real [0,100] attributes of type char\_freq\_CHAR  
= percentage of characters in the e-mail that match CHAR,  
i.e.  $100 * (\text{number of CHAR occurrences}) / \text{total characters in e-mail}$

1 continuous real [1,...] attribute of type capital\_run\_length\_average  
= average length of uninterrupted sequences of capital letters

1 continuous integer [1,...] attribute of type capital\_run\_length\_longest  
= length of longest uninterrupted sequence of capital letters

1 continuous integer [1,...] attribute of type capital\_run\_length\_total  
= sum of length of uninterrupted sequences of capital letters  
= total number of capital letters in the e-mail

1 nominal {0,1} class attribute of type spam  
= denotes whether the e-mail was considered spam (1) or not (0),  
i.e. unsolicited commercial e-mail.

For more information, see file ‘spambase.DOCUMENTATION’ at the  
UCI Machine Learning Repository: <http://www.ics.uci.edu/~mlearn/MLRepository.html>

# A. Data pre-processing and pre-analysis

1. Read in the data
2. Explore the features and target variables to assess what parameters will need to be changed
3. Prepare & transform data for data mining process

```
In [3]: ## Pull in all the column names based on the spam base documentation

colnames = ["word_freq_make", "word_freq_address", "word_freq_all", "word_freq_3d", "word_freq_our",
            "word_freq_over", "word_freq_remove", "word_freq_internet", "word_freq_order", "word_freq_mail",
            "word_freq_receive", "word_freq_will", "word_freq_people", "word_freq_report", "word_freq_addresses",
            "word_freq_free", "word_freq_business", "word_freq_email", "word_freq_you", "word_freq_credit",
            "word_freq_your", "word_freq_font", "word_freq_000", "word_freq_money", "word_freq_hp",
            "word_freq_hpl", "word_freq_george", "word_freq_650", "word_freq_lab", "word_freq_labs",
            "word_freq_telnet", "word_freq_857", "word_freq_data", "word_freq_415", "word_freq_85",
            "word_freq_technology", "word_freq_1999", "word_freq_parts", "word_freq_pm", "word_freq_direct",
            "word_freq_cs", "word_freq_meeting", "word_freq_original", "word_freq_project", "word_freq_re",
            "word_freq_edu", "word_freq_table", "word_freq_conference", "char_freq;", "char_freq(",
            "char_freq_", "char_freq!", "char_freq$", "char_freq#", "capital_run_length_average",
            "capital_run_length_longest", "capital_run_length_total", "spam"]
```

```
In [4]: ## Read in the file and save it as a dataframe

spambase_df = pd.read_csv("spambase.data", header = None, names = colnames)
```


```
In [5]: ## Look at the df to make sure it loaded correctly

spambase_df.tail()
```

Out[5]:

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	word_freq_order
4596	0.31	0.0	0.62	0.0	0.00	0.31	0.0	0.0	0.00
4597	0.00	0.0	0.00	0.0	0.00	0.00	0.0	0.0	0.00
4598	0.30	0.0	0.30	0.0	0.00	0.00	0.0	0.0	0.00
4599	0.96	0.0	0.00	0.0	0.32	0.00	0.0	0.0	0.00
4600	0.00	0.0	0.65	0.0	0.00	0.00	0.0	0.0	0.00

5 rows × 58 columns




In [6]: *## Generate some summary statistics of the dataframe, just to see the distribution of the various attributes*

```
spambase_df.describe()
```

Out[6]:

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	wo
count	4601.000000	4601.000000	4601.000000	4601.000000	4601.000000	4601.000000	4601.000000	4601.000000	
mean	0.104553	0.213015	0.280656	0.065425	0.312223	0.095901	0.114208	0.105295	
std	0.305358	1.290575	0.504143	1.395151	0.672513	0.273824	0.391441	0.401071	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.420000	0.000000	0.380000	0.000000	0.000000	0.000000	
max	4.540000	14.280000	5.100000	42.810000	10.000000	5.880000	7.270000	11.110000	

8 rows × 58 columns



In [7]: *## Grab all columns as features except our target which is spam!*

```
spam_feature_cols = spambase_df.columns[spambase_df.columns != "spam"]
```

In [8]: `spambase_df.spam.unique()`

Out[8]: `array([1, 0], dtype=int64)`

In [9]: `spambase_df.groupby("spam")["word_freq_make"].count()`

Out[9]:

spam	
0	2788
1	1813

Name: word\_freq\_make, dtype: int64

In [10]: *## View splits for the targets*

```
total_obvs = spambase_df.groupby("spam")["word_freq_make"].count().sum()
total_non_spam = spambase_df.groupby("spam")["word_freq_make"].count()[0]
total_spam = spambase_df.groupby("spam")["word_freq_make"].count()[1]

## Uneven balanced classes
print("{} total non spam messages, \
      {} % of overall observations".format(total_non_spam, round(total_non_spam/total_obvs*100, 2)))
print("{} total spam messages, \
      {} % of overall observations".format(total_spam, round(total_spam/total_obvs*100, 2)))
```

```
2788 total non spam messages,      60.6 % of overall observations
1813 total spam messages,         39.4 % of overall observations
```

**We have unbalanced classes, so we will stratify our splits before developing the model. For a classification task, this is chosen to ensure that the train and test sets have approximately the same percentage of samples of each target class as the complete data set.**

In [12]: *## Grabbing all the features available*

```
X = np.array(spambase_df[spam_feature_cols])
```

*## Target variable*

```
y = np.array(spambase_df["spam"])
```

In [13]: *## Split data training 70 % and testing 30%*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42, stratify = y)
```

```
In [260]: ## We normalize our training and testing data PCA to work correctly
## We don't want to skew the results of the plot because some features are not on the same scale
## We perform this normalization AFTER splitting the data - again, so that we don't skew the training
## data with the testing data.

## Normalization is the process of scaling individual samples to have unit norm. This process can be useful
## if you plan to use a quadratic form such as the dot-product or any other kernel to quantify the similarity of any pair of samples.
X_train_norm = Normalizer(norm = "l1").fit_transform(X_train)
X_train_robust = RobustScaler().fit_transform(X_train)
X_train_power = PowerTransformer(method='yeo-johnson', standardize=False).fit_transform(X_train)
X_train_standard = StandardScaler().fit_transform(X_train)

## Import PCA from sklearn
from sklearn.decomposition import PCA

## Initialize two new PCA instances, we'll use this to plot the training data using two different transformations
## Normalizer will likely be skewed by the outliers in each of the features being used
## Robust is going to transform feature values to be larger than the previous scalers and more importantly are approximately similar to original data
## PowerTransformer is a family of parametric, monotonic transformations that aim to map data from any distribution to as close to a Gaussian distribution
## as possible in order to stabilize variance and minimize skewness.

pca = PCA()
pca_r = PCA()
pca_p = PCA()
pca_s = PCA()

## The goal of this plot is to determine what features need to be included in our models
## In the first few homeworks, we threw the kitchen sink at the models. Here, we are going to be
## more refined in our analysis.

## Train the PCA instance using the normalized training data
pca.fit(X_train_norm)
pca_r.fit(X_train_robust)
pca_p.fit(X_train_power)
pca_s.fit(X_train_standard)

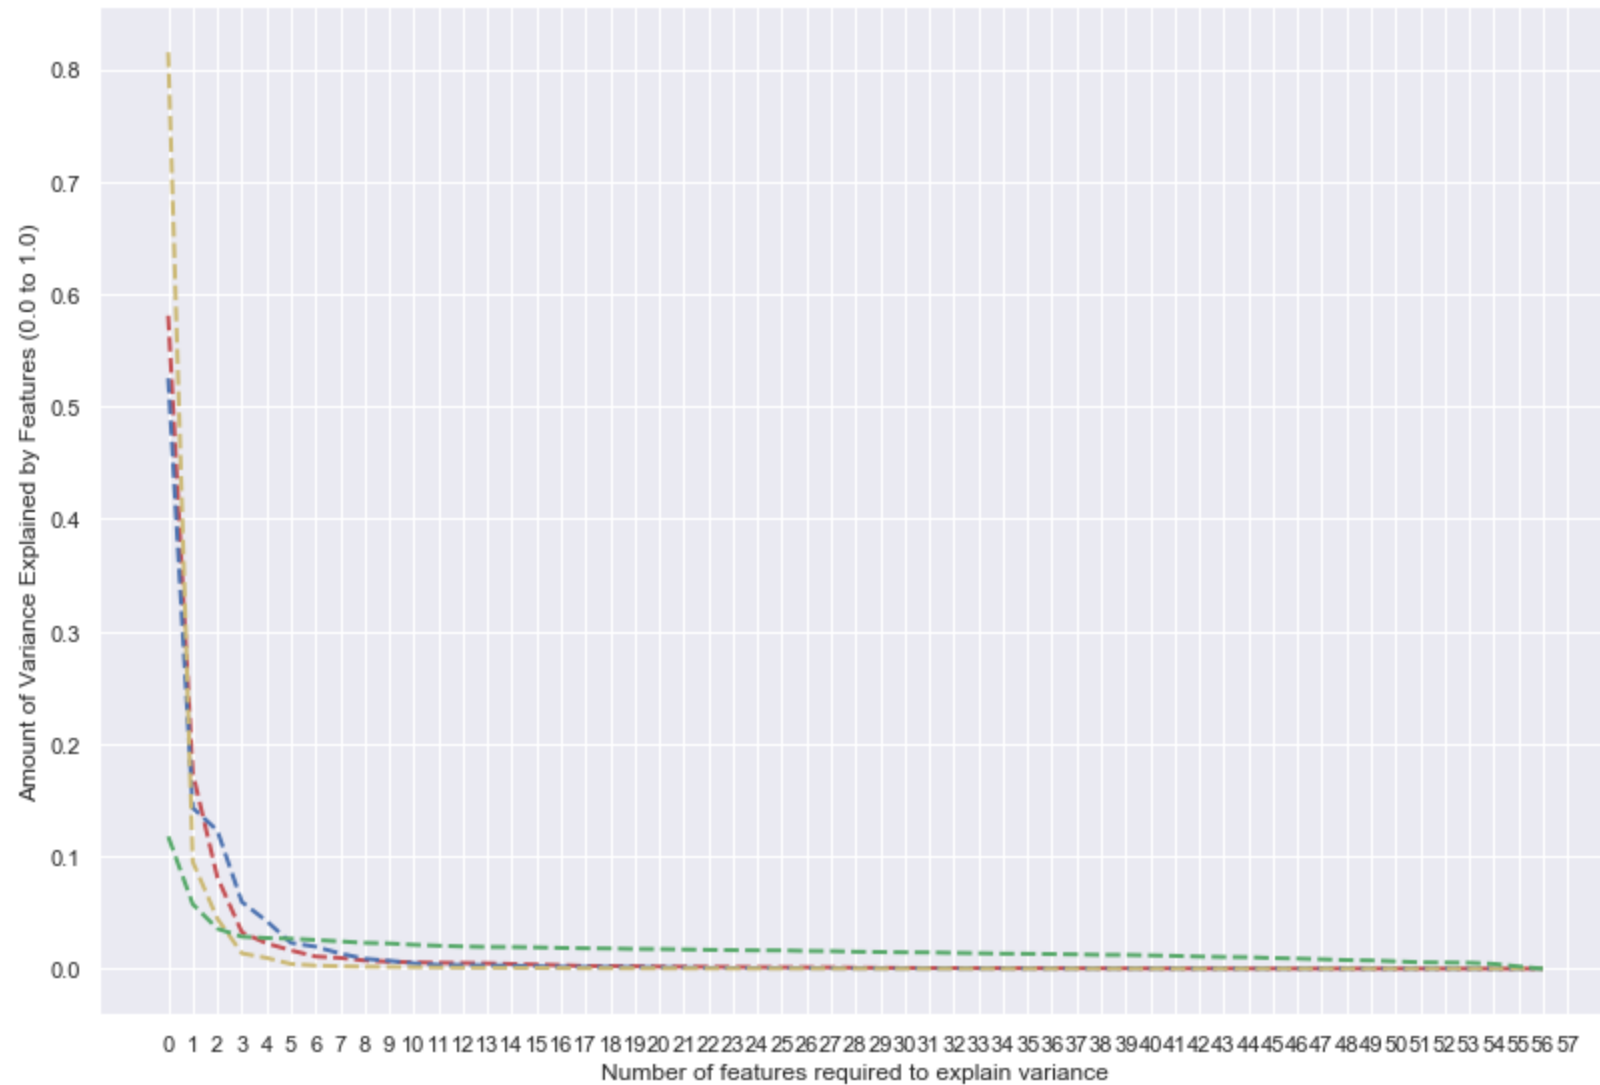
plt.figure(1, figsize=(15, 10))
plt.clf()
plt.axes([.2, .2, .7, .7])

plt.plot(pca.explained_variance_ratio_, 'r--', linewidth = 2)
plt.plot(pca_r.explained_variance_ratio_, 'b--', linewidth = 2)
plt.plot(pca_p.explained_variance_ratio_, 'y--', linewidth = 2)
plt.plot(pca_s.explained_variance_ratio_, 'g--', linewidth = 2)
```

```
## Set plot labels
plt.xlabel('Number of features required to explain variance')
plt.ylabel('Amount of Variance Explained by Features (0.0 to 1.0)')

## Explicitly set the x-axis data so we can see where the drop-off is
plt.xticks(np.arange(0, 58, step=1))

## Show the graph!
plt.show()
```





**Only a very limited amount of features explain the variance in the data set. For some of the classifiers, more information might actually lead to worse performance, either through overfitting or not being able to converge successfully.**

In [267]: *## SET A VARIABLE TO SWAP IN AND OUT BETWEEN THE DIFFERENT TRAINING INSTANCES*

```
#Z = X_train_robust
#Z = X_train_norm
#Z = X_train_power
Z = X_train_standard

## Start with identifying the best features using a Random Forest classifier

## Create a new classifier
clf_rf_5 = ensemble.RandomForestClassifier()
clr_rf_5 = clf_rf_5.fit(Z, y_train)

## Save our importances to a variable
importances = clr_rf_5.feature_importances_

## Get the standard deviation for each feature
std = np.std([tree.feature_importances_ for tree in clf_rf_5.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

## Print the feature ranking
print("Feature ranking:")

## Print the top ten features, and their importance based on the Random Forest Classifier
for f in range(0, 10):
    print("%.2d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

## Plot the feature importances of the Random Forest Regressor - to see this visually

## Set the plot size
plt.figure(1, figsize=(15, 10))

## Set the title
plt.title("Feature importances")

## Plot a graph using all of the normalized features
plt.bar(range(Z.shape[1]), importances[indices],
        color="b", yerr=std[indices], align="center")


plt.xticks(range(Z.shape[1]), spam_feature_cols[indices], rotation=90)
plt.xlim([-1, Z.shape[1]])

## Show the graph!
plt.show()
```

Feature ranking:

1. feature 51 (0.141203)
2. feature 23 (0.076135)
3. feature 6 (0.073537)
4. feature 52 (0.073031)
5. feature 54 (0.061571)
6. feature 55 (0.061143)
7. feature 56 (0.058942)
8. feature 15 (0.049227)
9. feature 20 (0.047997)
10. feature 22 (0.032252)

Feature importances



**Similar to the analysis above, we don't see much need for many of these features; maybe only the first thirty or so will be important for performing this classification task.**

**One final feature analysis, using a technique known as step-forward feature selection.**

**Step forward feature selection starts with the evaluation of each individual feature, and selects that which results in the best performing selected algorithm model. What's the "best?" That depends entirely on the defined evaluation criteria (AUC, prediction accuracy, RMSE, etc.). I am using accuracy as this is a classification task.**

**Next, all possible combinations of the that selected feature and a subsequent feature are evaluated, and a second feature is selected, and so on, until the required predefined number of features is selected.**

```
In [268]: ## Build RF classifier to use in feature selection
clf = ensemble.RandomForestClassifier()

# Build step forward feature selection
sfs1 = sfs(clf,
            k_features=6,
            forward=True, # Otherwise, this will be the backward selection
            floating=False,
            n_jobs=10, # The number of CPUs to use for evaluating
            verbose=2,
            scoring='accuracy',
            cv=5)
```

```
# Perform SFFS
```

```
sfs1 = sfs1.fit(X_train_standard, y_train)
```

```
[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 4.9s
```

```
[Parallel(n_jobs=10)]: Done 57 out of 57 | elapsed: 5.3s finished
```

```
[2019-10-31 16:08:16] Features: 1/6 -- score: 0.7779549202525267[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 2.5s
```

```
[Parallel(n_jobs=10)]: Done 56 out of 56 | elapsed: 3.0s finished
```

```
[2019-10-31 16:08:20] Features: 2/6 -- score: 0.8388206859921853[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 2.4s
```

```
[Parallel(n_jobs=10)]: Done 55 out of 55 | elapsed: 2.9s finished
```

```
[2019-10-31 16:08:23] Features: 3/6 -- score: 0.8680132610309345[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 2.5s
```

```
[Parallel(n_jobs=10)]: Done 54 out of 54 | elapsed: 3.1s finished
```

```
[2019-10-31 16:08:26] Features: 4/6 -- score: 0.8903754533969794[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 2.4s
```

```
[Parallel(n_jobs=10)]: Done 53 out of 53 | elapsed: 2.9s finished
```

```
[2019-10-31 16:08:30] Features: 5/6 -- score: 0.9040371742279195[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.
```

```
[Parallel(n_jobs=10)]: Done 21 tasks | elapsed: 2.4s
```

```
[Parallel(n_jobs=10)]: Done 52 out of 52 | elapsed: 2.9s finished
```

```
[2019-10-31 16:08:33] Features: 6/6 -- score: 0.9155259375009127
```

```
In [287]: ## Now we print out what feature columns of Linear Regressor picked out - we can use these to subset later  
## We'll remember to grab these columns further along in the analysis  
feat_cols = list(sfs1.k_feature_idx_)  
print(feat_cols)
```

```
[6, 24, 26, 45, 52, 55]
```

```
In [204]: ## Display the column names related to our "premier" features  
for i in spambase_df.columns[feat_cols]:  
    print(i)
```

```
word_freq_remove  
word_freq_hp  
word_freq_george  
word_freq_edu  
char_freq_$  
capital_run_length_longest
```

**Based on this analysis, the top six features are presented above.**

```
In [14]: ## Create new scaling object  
  
sc = StandardScaler()  
  
## Standardize dataset  
  
X_train_standard = sc.fit_transform(X_train)  
X_test_standard = sc.transform(X_test)
```

## B. Model Creation and Evaluation

1. Create parameter grids for each model
2. Used nested cross validation to determine the best model
3. Tune the hyper parameters for the best model
4. Evaluate the models on the testing data

## Here are the different models I will be using for my analysis for this classification problem

1. Logistic(linear) Classifier
2. k-Nearest Neighbors Classifier
3. Decision Tree Classifier
4. SVM Classifier
5. Ensemble/Gradient Boost Classifier
6. Neural Network (Keras & KerasTuner)

In [135]: ##### SETTING UP PARAMETER GRIDS FOR THE DIFFERENT DATA MINING MODELS WE ARE GOING TO USE #####

```
## Set up a grid for the Logit Classifier
lc_p_grid = {"penalty": ["l2"],
             "C": [1, 5, 10, 50, 1000],
             "solver": ["liblinear"]}

## Set up a grid for kNN Classifier
## Going to use 1-30 neighbors, and two different distance calculations
knnc_p_grid = {"n_neighbors": list(range(1, 31)),
               "weights": ["uniform", "distance"]}

## Set up a grid for the DecisionTree Classifier
dtc_p_grid = {"criterion": ["gini"],
              "splitter": ["best"],
              "max_features": [15, 20, 25],
              "max_depth": [5, 10, 15]}

## Set up a grid for the Support Vector Classifier
svc_p_grid = {"C": [1, 10, 50, 1000],
              "gamma": [0.0001, 0.0005, 0.001, 0.005],
              "kernel": ["poly", "rbf"]}

## Set up a grid for GBoost Classifier
gbc_p_grid = {'loss': ["deviance"],
              'n_estimators': [100, 200, 300, 400, 500],
              'max_depth': [3, 4, 5],
              'min_samples_split': [2, 4, 6],
              'max_features': [5, 10],
              'criterion': ["mse"],
              'learning_rate': [0.01]}
```



In [136]: *## Set a number of trials to run for the models*

```
num_trials = 20

## Empty arrays to store scores for classifier
nested_scores_lc = np.zeros(num_trials)
nested_scores_knnc = np.zeros(num_trials)
nested_scores_dtc = np.zeros(num_trials)
nested_scores_svc = np.zeros(num_trials)
nested_scores_gbc = np.zeros(num_trials)
```

In [137]: *## Create new regressors for each data mining technique*

```
## Linear Regression
lc = linear_model.LogisticRegression()

## k-Nearest Neighbors
knnc = neighbors.KNeighborsClassifier()

## Decision Tree
dtc = tree.DecisionTreeClassifier()

## Support Vector Machine
svc = svm.SVC()

## Gradient Boost
gbc = ensemble.GradientBoostingClassifier()
```

In [303]: *## Loop for each trial*

```
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

    ## Nested CV for Logit Regression

    lclass= GridSearchCV(estimator=lc, param_grid=lc_p_grid, cv=inner_cv)
    lclass.fit(X_train_standard, y_train)

    nested_score = cross_val_score(lclass, X = X_train_standard, y = y_train, cv = outer_cv)
    nested_scores_lc[i] = nested_score.mean()
```

```
In [305]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Logit Regression

    knnclass= GridSearchCV(estimator=knnc, param_grid=knnc_p_grid, cv=inner_cv)
    knnclass.fit(X_train_standard, y_train)

    nested_score = cross_val_score(knnclass, X = X_train_standard, y = y_train, cv = outer_cv)
    nested_scores_knnr[i] = nested_score.mean()
```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-305-16a7d48e722e> in <module>
    12     knnclass.fit(X_train_standard, y_train)
    13
--> 14     nested_score = cross_val_score(knnclass, X = X_train_standard, y = y_train, cv = outer_cv)
    15     nested_scores_knnr[i] = nested_score.mean()

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_val_score(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, error_score)
    400         fit_params=fit_params,
    401         pre_dispatch=pre_dispatch,
--> 402         error_score=error_score)
    403     return cv_results['test_score']
    404

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_validate(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, return_train_score, return_estimator, error_score)
    238         return_times=True, return_estimator=return_estimator,
    239         error_score=error_score)
--> 240     for train, test in cv.split(X, y, groups))
    241
    242     zipped_scores = list(zip(*scores))

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
    915         # remaining jobs.
    916         self._iterating = False
--> 917         if self.dispatch_one_batch(iterator):
    918             self._iterating = self._original_iterator is not None
    919

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
    757         return False
    758     else:
--> 759         self._dispatch(tasks)
    760         return True
    761

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
    714         with self._lock:
    715             job_idx = len(self._jobs)
--> 716             job = self._backend.apply_async(batch, callback=cb)
    717             # A job can complete so quickly than its callback is
    718             # called before we get here, causing self._jobs to

C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
    180     def apply_async(self, func, callback=None):

```

```

181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel_backends.py in __init__(self, batch)
547         # Don't delay the application, to avoid keeping the input
548         # arguments in memory
--> 549         self.results = batch()
550
551     def get(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, return_estimator, error_score)
526         estimator.fit(X_train, **fit_params)
527     else:
--> 528         estimator.fit(X_train, y_train, **fit_params)
529
530     except Exception as e:

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in fit(self, X, y, groups, **fit_params)
720         return results_container[0]
721
--> 722         self._run_search(evaluate_candidates)
723
724         results = results_container[0]

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in _run_search(self, evaluate_candidates)
1189     def _run_search(self, evaluate_candidates):
1190         """Search all candidates in param_grid"""
-> 1191         evaluate_candidates(ParameterGrid(self.param_grid))
1192
1193

```

```
C:\Python\lib\site-packages\sklearn\model_selection\_search.py in evaluate_candidates(candidate_params)
709         for parameters, (train, test)
710         in product(candidate_params,
--> 711                   cv.split(X, y, groups)))
712
713     all_candidate_params.extend(candidate_params)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
918     self._iterating = self._original_iterator is not None
919
--> 920     while self.dispatch_one_batch(iterator):
921         pass
922
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
757     return False
758     else:
--> 759         self._dispatch(tasks)
760         return True
761
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
714     with self._lock:
715         job_idx = len(self._jobs)
--> 716         job = self._backend.apply_async(batch, callback=cb)
717         # A job can complete so quickly that its callback is
718         # called before we get here, causing self._jobs to
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
180     def apply_async(self, func, callback=None):
181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in __init__(self, batch)
547     # Don't delay the application, to avoid keeping the input
548     # arguments in memory
--> 549     self.results = batch()
550
551     def get(self):
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223     with parallel_backend(self._backend, n_jobs=self._n_jobs):
224         return [func(*args, **kwargs)
--> 225                 for func, args, kwargs in self.items]
226
```

```

227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, t
est, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, retu
rn_estimator, error_score)
570         if return_train_score:
571             train_scores = _score(estimator, X_train, y_train, scorer,
--> 572                                 is_multimetric)
573
574         if verbose > 2:

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _score(estimator, X_test, y_test, scorer, is_mul
timetric)
603         """
604         if is_multimetric:
--> 605             return _multimetric_score(estimator, X_test, y_test, scorer)
606         else:
607             if y_test is None:

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _multimetric_score(estimator, X_test, y_test, sc
orers)
633             score = scorer(estimator, X_test)
634         else:
--> 635             score = scorer(estimator, X_test, y_test)
636
637             if hasattr(score, 'item'):

C:\Python\lib\site-packages\sklearn\metrics\scorer.py in _passthrough_scorer(estimator, *args, **kwargs)
239 def _passthrough_scorer(estimator, *args, **kwargs):
240     """Function that wraps estimator.score"""
--> 241     return estimator.score(*args, **kwargs)
242
243

C:\Python\lib\site-packages\sklearn\base.py in score(self, X, y, sample_weight)
288         """
289         from .metrics import accuracy_score
--> 290         return accuracy_score(y, self.predict(X), sample_weight=sample_weight)
291
292

```

C:\Python\lib\site-packages\sklearn\neighbors\classification.py in predict(self, X)

```
147         X = check_array(X, accept_sparse='csr')
148
--> 149         neigh_dist, neigh_ind = self.kneighbors(X)
150         classes_ = self.classes_
151         _y = self._y
```

C:\Python\lib\site-packages\sklearn\neighbors\base.py in kneighbors(self, X, n\_neighbors, return\_distance)

```
453         delayed_query(
454             self._tree, X[s], n_neighbors, return_distance)
--> 455         for s in gen_even_slices(X.shape[0], n_jobs)
456     )
457     else:
```

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in \_\_call\_\_(self, iterable)

```
915         # remaining jobs.
916         self._iterating = False
--> 917         if self.dispatch_one_batch(iterator):
918             self._iterating = self._original_iterator is not None
919
```

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch\_one\_batch(self, iterator)

```
757         return False
758     else:
--> 759         self._dispatch(tasks)
760         return True
761
```

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in \_dispatch(self, batch)

```
714         with self._lock:
715             job_idx = len(self._jobs)
--> 716             job = self._backend.apply_async(batch, callback=cb)
717             # A job can complete so quickly than its callback is
718             # called before we get here, causing self._jobs to
```

C:\Python\lib\site-packages\sklearn\externals\joblib\\_parallel\_backends.py in apply\_async(self, func, callback)

```
180     def apply_async(self, func, callback=None):
181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)
```

C:\Python\lib\site-packages\sklearn\externals\joblib\\_parallel\_backends.py in \_\_init\_\_(self, batch)

```
547         # Don't delay the application, to avoid keeping the input
548         # arguments in memory
--> 549         self.results = batch()
550
551     def get(self):
```

```

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
    223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
    224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
    226
    227     def __len__(self):

```

```

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
    223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
    224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
    226
    227     def __len__(self):

```

```

C:\Python\lib\site-packages\sklearn\neighbors\base.py in _tree_query_parallel_helper(tree, data, n_neighbors, return_d
istance)
    290     under PyPy.
    291     """
--> 292     return tree.query(data, n_neighbors, return_distance)
    293
    294

```

**KeyboardInterrupt:**

```

In [318]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

    ## Nested CV for Logit Regression

    dtclass= GridSearchCV(estimator=dtc, param_grid=dtc_p_grid, cv=inner_cv)
    dtclass.fit(X_train_standard, y_train)

    nested_score = cross_val_score(dtreg, X = X_train_standard, y = y_train, cv = outer_cv)
    nested_scores_dtc[i] = nested_score.mean()

```



```
In [138]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Logit Regression

    svclass= GridSearchCV(estimator=svc, param_grid=svc_p_grid, cv=inner_cv)
    svclass.fit(X_train_standard, y_train)

    nested_score = cross_val_score(svclass, X = X_train_standard, y = y_train, cv = outer_cv)
    nested_scores_svc[i] = nested_score.mean()
```

```
In [365]: ## Loop for each trial
for i in range(20):

    ## Choose cross-validation techniques for the inner and outer loops,
    ## independently of the dataset.
    inner_cv = KFold(n_splits = 4, shuffle = True, random_state = i)
    outer_cv = KFold(n_splits = 4, shuffle = True, random_state = i)

## Nested CV for Gradient Boost Classifier

    gbclass= GridSearchCV(estimator=gbc, param_grid=gbc_p_grid, cv=inner_cv)
    gbclass.fit(X_train_standard, y_train)

    nested_score = cross_val_score(gbclass, X = X_train_standard, y = y_train, cv = outer_cv)
    nested_scores_gbc[i] = nested_score.mean()
```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-365-60868a14b3d3> in <module>
    12     gbclass.fit(X_train_standard, y_train)
    13
--> 14     nested_score = cross_val_score(gbclass, X = X_train_standard, y = y_train, cv = outer_cv)
    15     nested_scores_gbc[i] = nested_score.mean()

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_val_score(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, error_score)
    400         fit_params=fit_params,
    401         pre_dispatch=pre_dispatch,
--> 402         error_score=error_score)
    403     return cv_results['test_score']
    404

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in cross_validate(estimator, X, y, groups, scoring, cv, n_jobs, verbose, fit_params, pre_dispatch, return_train_score, return_estimator, error_score)
    238         return_times=True, return_estimator=return_estimator,
    239         error_score=error_score)
--> 240     for train, test in cv.split(X, y, groups))
    241
    242     zipped_scores = list(zip(*scores))

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
    918         self._iterating = self._original_iterator is not None
    919
--> 920         while self.dispatch_one_batch(iterator):
    921             pass
    922

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
    757         return False
    758     else:
--> 759         self._dispatch(tasks)
    760         return True
    761

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
    714         with self._lock:
    715             job_idx = len(self._jobs)
--> 716             job = self._backend.apply_async(batch, callback=cb)
    717             # A job can complete so quickly than its callback is
    718             # called before we get here, causing self._jobs to

C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
    180     def apply_async(self, func, callback=None):

```

```

181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel_backends.py in __init__(self, batch)
547         # Don't delay the application, to avoid keeping the input
548         # arguments in memory
--> 549         self.results = batch()
550
551     def get(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, return_estimator, error_score)
526         estimator.fit(X_train, **fit_params)
527     else:
--> 528         estimator.fit(X_train, y_train, **fit_params)
529
530     except Exception as e:

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in fit(self, X, y, groups, **fit_params)
720         return results_container[0]
721
--> 722         self._run_search(evaluate_candidates)
723
724         results = results_container[0]

C:\Python\lib\site-packages\sklearn\model_selection\_search.py in _run_search(self, evaluate_candidates)
1189     def _run_search(self, evaluate_candidates):
1190         """Search all candidates in param_grid"""
-> 1191         evaluate_candidates(ParameterGrid(self.param_grid))
1192
1193

```

```
C:\Python\lib\site-packages\sklearn\model_selection\_search.py in evaluate_candidates(candidate_params)
709         for parameters, (train, test)
710         in product(candidate_params,
--> 711                   cv.split(X, y, groups)))
712
713     all_candidate_params.extend(candidate_params)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self, iterable)
918     self._iterating = self._original_iterator is not None
919
--> 920     while self.dispatch_one_batch(iterator):
921         pass
922
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in dispatch_one_batch(self, iterator)
757     return False
758     else:
--> 759         self._dispatch(tasks)
760         return True
761
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in _dispatch(self, batch)
714     with self._lock:
715         job_idx = len(self._jobs)
--> 716         job = self._backend.apply_async(batch, callback=cb)
717         # A job can complete so quickly that its callback is
718         # called before we get here, causing self._jobs to
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in apply_async(self, func, callback)
180     def apply_async(self, func, callback=None):
181         """Schedule a func to be run"""
--> 182         result = ImmediateResult(func)
183         if callback:
184             callback(result)
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\_parallel_backends.py in __init__(self, batch)
547     # Don't delay the application, to avoid keeping the input
548     # arguments in memory
--> 549     self.results = batch()
550
551     def get(self):
```

```
C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in __call__(self)
223     with parallel_backend(self._backend, n_jobs=self._n_jobs):
224         return [func(*args, **kwargs)
--> 225                 for func, args, kwargs in self.items]
226
```

```

227     def __len__(self):

C:\Python\lib\site-packages\sklearn\externals\joblib\parallel.py in <listcomp>(.0)
223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
224             return [func(*args, **kwargs)
--> 225                     for func, args, kwargs in self.items]
226
227     def __len__(self):

C:\Python\lib\site-packages\sklearn\model_selection\_validation.py in _fit_and_score(estimator, X, y, scorer, train, t
est, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, retu
rn_estimator, error_score)
526         estimator.fit(X_train, **fit_params)
527     else:
--> 528         estimator.fit(X_train, y_train, **fit_params)
529
530     except Exception as e:

C:\Python\lib\site-packages\sklearn\ensemble\gradient_boosting.py in fit(self, X, y, sample_weight, monitor)
1463         n_stages = self._fit_stages(X, y, y_pred, sample_weight, self._rng,
1464                                     X_val, y_val, sample_weight_val,
-> 1465                                     begin_at_stage, monitor, X_idx_sorted)
1466
1467         # change shape of arrays after fit (early-stopping or additional ests)

C:\Python\lib\site-packages\sklearn\ensemble\gradient_boosting.py in _fit_stages(self, X, y, y_pred, sample_weight, ra
ndom_state, X_val, y_val, sample_weight_val, begin_at_stage, monitor, X_idx_sorted)
1527         y_pred = self._fit_stage(i, X, y, y_pred, sample_weight,
1528                                 sample_mask, random_state, X_idx_sorted,
-> 1529                                 X_csc, X_csr)
1530
1531         # track deviance (= loss)

C:\Python\lib\site-packages\sklearn\ensemble\gradient_boosting.py in _fit_stage(self, i, X, y, y_pred, sample_weight,
sample_mask, random_state, X_idx_sorted, X_csc, X_csr)
1192         X = X_csr if X_csr is not None else X
1193         tree.fit(X, residual, sample_weight=sample_weight,
-> 1194                check_input=False, X_idx_sorted=X_idx_sorted)
1195
1196         # update tree leaves

C:\Python\lib\site-packages\sklearn\tree\tree.py in fit(self, X, y, sample_weight, check_input, X_idx_sorted)
1140         sample_weight=sample_weight,
1141         check_input=check_input,
-> 1142         X_idx_sorted=X_idx_sorted)
1143         return self
1144

```

```
C:\Python\lib\site-packages\sklearn\tree\tree.py in fit(self, X, y, sample_weight, check_input, X_idx_sorted)
    364         min_impurity_split)
    365
--> 366         builder.build(self.tree_, X, y, sample_weight, X_idx_sorted)
    367
    368         if self.n_outputs_ == 1:
```

**KeyboardInterrupt:**

```
In [356]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV

## Set the figure size
plt.figure(figsize= (15, 10))


## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(nested_scores_svc, 'r--', label = "SVR")
plt.plot(nested_scores_knnc, 'y--', label = "kNN Classification")
plt.plot(nested_scores_lc, 'g--', label = "Linear Regression")
plt.plot(nested_scores_dtc, 'b--', label = "Decision Tree Classification")
plt.plot(nested_scores_gbc, 'v--', label = "Gradient Boost")

## Give some labels
plt.xlabel("No. of Trials")
plt.ylabel("Score")

## Title and Legend
plt.title("Nested Cross Validation Scores")
plt.legend(loc = 'lower left')

## Show the graph
plt.show()
```





**From our Cross Validation plot above, we see that a Support Vector Machine classifier performs the best at classifying spam email out of the different data mining procedures tested.**

**Since it is pretty clear that this was an effective model, I will do further exploration using the best parameters selected by the nested cross validation process.**

```
In [140]: ## Get our best params and their scores
print(svcclass.best_params_)
print()

## Print out how well it performed using the best params
print(svcclass.best_score_)

## save our best params so we can use them in our actual SVC model!
best_svc_params = svcclass.best_params_

{'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}

0.9347826086956522
```

```
In [147]: ## Initialize our DecisionTree classifier with the best params based on our GridSearch.
svc_clf = svm.SVC(*best_svc_params, probability = True)

## Train the model (fit the data)
# 'fit' builds a decision tree from the training set (X, y).
svc_clf = svc_clf.fit(X_train_standard, y_train)
```

```
In [148]: ## Evaluate performance by cross-validation
scores = cross_val_score(svc_clf, X_train_standard, y_train, cv = 10)
print(scores)

print()

# The mean score and the 95% confidence interval of our scores:
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

[0.92569659 0.91925466 0.91925466 0.94099379 0.93167702 0.92546584
 0.93478261 0.94099379 0.95341615 0.93457944]

Accuracy: 0.93 (+/- 0.02)
```

```
In [149]: ## Create three empty lists to store my precision, recall, and average precision scores
precision, recall, average_precision = [], [], []

## Get probabilities for our labels
probas_ = svc_clf.predict_proba(X_test_standard)
```

```
In [156]: ## Create my false positive rate, true positive rate, and threshold using my test data
fpr, tpr, thresholds = roc_curve(y_test, probas[:, 1], pos_label = 1)

## Create precision and recall scores to plot with
precision_score, recall_score, _ = precision_recall_curve(y_test, probas[:, 1])

## Calculate the overall AUC for the model
auc = np.trapz(tpr, fpr)

## Save the average precision for our model
avg_precision = average_precision_score(y_test, probas[:, 1])
```

```
In [206]: ## Build a confusion matrix from our Support Vector Machine model - another assessment of performance

## Try to predict the outcomes on our test data
predicted = svc_clf.predict(X_test_standard)

## Compare that with our ACTUAL values from the test data set
matrix = confusion_matrix(y_test, predicted)
print(matrix)
print()

## Create a report to show our precision(accuracy), recall, and f1 for predictions
report = classification_report(y_test, predicted)
print(report)
```

```
[[798  39]
 [ 51 493]]
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	837
1	0.93	0.91	0.92	544
accuracy			0.93	1381
macro avg	0.93	0.93	0.93	1381
weighted avg	0.93	0.93	0.93	1381

**The AUC Curve is very, very close to being an almost perfect line. This tells me this model is very effective at discovering and identifying spam emails.**

```
In [198]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Create a new figure to plot
plt.figure(figsize= (15, 10))

lw = 2

## Plot the model's Precision Recall Curve
plt.plot(precision_score, recall_score, linestyle = '-', label='Precision-Recall curve (area = {})'.format(round(avg_precision, 4)))

## Put in a line to demonstrate blind luck
plt.plot([0, 1], [0, 1], color = 'black', linestyle = '--', label = 'Luck')


## Plot the perfect Precision Recall Line to see how well our SVM Classifier performed
plt.plot([0, 1], [1, 1], color = 'orange')
plt.plot([1, 1], [1, 0.5], color = 'orange', label = "Perfect Precision Recall")

## Set the limits so they start at zero
plt.ylim([-0.01, 1.01])
plt.xlim([-0.01, 1.01])

## Set labels for x and y
plt.xlabel('Recall')
plt.ylabel('Precision')

## Set a title and legend
plt.title('Precision Recall Curve')
plt.legend(loc = 'center left')

## Show the curve!
plt.show()
```



Due to the slight class imbalance, I also generate a Precision-Recall curve; accuracy may not be an effective measure of model performance so it is good to just confirm the model is working as expected.

Here we confirm our earlier observation that the model is performing extremely well, with an area of 0.9568.

To reinforce and showcase how the model is performing at its classification task, I include two lines: one showing random luck (black) and a perfect Precision-Recall (orange).

I also want to investigate the cost for misclassification; it's not enough to make sure the model is working as expected; we want to penalize the model for mistakes (not too harshly) to make sure that we don't allow e-mails through that should not be allowed in.

I will punish false positives more harshly than false negative. The reasoning behind this is if an email that *should* have been identified as spam get through, then it should be penalized vs. if the classifier misclassified a good email as spam.

If a spam email gets through the filter, then any harmful software associated with that spam message may infect the user's computer down the line.

```
In [ ]: ## Initialize different thresholds and Costs that will be tied to those thresholds
thresholds = np.linspace(0, 1.0, num=21)

## Generate three generic Cost List for each matrix
Cost_List = np.linspace(0, 1.0, num=21)

## Build cost matrix for misidentifying spam
cost_matrix = np.array([[0, 1], [10, 0]])
```

```
In [199]: index = 0

for t in thresholds:

    predict_thre = np.where(probas[:, 1] > t, 1, 0) ## prediction based on the preset threshold
    clf_matrix = confusion_matrix(y_test, predict_thre)
    Cost_List[index] = clf_matrix[0][0]*cost_matrix[0][0] +clf_matrix[0][1]*cost_matrix[0][1] +clf_matrix[1][0]*cost_ma
    trix[1][0] +clf_matrix[1][1]*cost_matrix[1][1]
    index+=1


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot each Cost Line individually
plt.plot(thresholds, Cost_List, 'r--', label = "Average Cost for Misclassification")

## Give some labels
plt.xlabel("thresholds")
plt.ylabel("Cost")

## Title and Legend
plt.title("Average Misclassification Cost Function")
plt.legend(loc = 'upper right')

## Show the Cost Matrix Analysis
plt.show()
```



### C. Use Keras and Keras Tuner to develop some classifiers

1. Build the model function
2. Build the hyper tuner function
3. Tune the model and explore the data space to make predictions
4. Assess the model performance.



```
In [15]: ## Just confirming the shape of the training data to know how many features we have  
X_train_standard.shape
```

```
Out[15]: (3220, 57)
```

**Now that I can use Keras and hyper tuning, I will create a classifier using a similar process.**

```
In [16]: ## First, we build a function to actually put together our model  
  
def build_class_model(hp):  
  
    ## Base Layer  
    cmodel = Sequential()  
  
    ## Add first layer, test with 10 up to 22 features  
    cmodel.add(layers.Dense(units = hp.Int("units",  
                                         min_value = 20,  
                                         max_value = 57,  
                                         step = 2),  
                        ## Use same initializer as model above  
                        activation = "relu", kernel_initializer = "he_uniform"))  
  
    ## Add target layer using linear regressor  
    cmodel.add(layers.Dense(1, activation = "sigmoid"))  
  
    ## Use different learning rates to test the model  
    cmodel.compile(  
        optimizer = keras.optimizers.SGD(  
            hp.Choice("learning_rate", values = [1e-2, 1e-3, 1e-4])),  
  
        ## Use same loss function as before  
        loss = "binary_crossentropy",  
  
        ## Test using Accuracy  
        metrics = ["accuracy", "binary_accuracy"])  
  
    ## Return completed model  
    return cmodel
```

In [99]: *## Build hyper tuner, will perform cross validation for us*

*### Use RandomSearch*

```
classtuner = RandomSearch(  
  
    ## Use model function from above  
    build_class_model,  
  
    ## What is objective function? Using Loss here  
    objective = "val_accuracy",  
  
    ## Set number of trials  
    max_trials = 5,  
  
    ## Number of executions  
    executions_per_trial = 3,  
  
    ## Set to short dir path  
    directory = "C:\\"  
  
)
```

```
In [22]: classtuner.search(X_train_standard, y_train, epochs = 10,  
                           validation_data = (X_test_standard, y_test))
```

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 1:15 - loss: 0.6932 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 2s - loss: 0.8565 - accuracy: 0.4414 - binary\_accuracy: 0.4414 - ETA: 1s - loss: 0.8521 - accuracy: 0.4428 - binary\_accuracy: 0.44 - ETA: 0s - loss: 0.8269 - accuracy: 0.4636 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8334 - accuracy: 0.4611 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8260 - accuracy: 0.4605 - binary\_accuracy: 0.46 - 1s 400us/sample - loss: 0.8271 - accuracy: 0.4584 - binary\_accuracy: 0.4584 - val\_loss: 0.8325 - val\_accuracy: 0.4881 - val\_binary\_accuracy: 0.4881

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.8710 - accuracy: 0.4375 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.8556 - accuracy: 0.4315 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.8306 - accuracy: 0.4617 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8270 - accuracy: 0.4587 - binary\_accuracy: 0.45 - ETA: 0s - loss: 0.8212 - accuracy: 0.4604 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8184 - accuracy: 0.4633 - binary\_accuracy: 0.46 - 0s 123us/sample - loss: 0.8215 - accuracy: 0.4637 - binary\_accuracy: 0.4637 - val\_loss: 0.8270 - val\_accuracy: 0.4902 - val\_binary\_accuracy: 0.4902

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 1.0493 - accuracy: 0.4062 - binary\_accuracy: 0.40 - ETA: 0s - loss: 0.8164 - accuracy: 0.4492 - binary\_accuracy: 0.44 - ETA: 0s - loss: 0.8136 - accuracy: 0.4614 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8119 - accuracy: 0.4701 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.8165 - accuracy: 0.4719 - binary\_accuracy: 0.47 - 0s 108us/sample - loss: 0.8159 - accuracy: 0.4699 - binary\_accuracy: 0.4699 - val\_loss: 0.8217 - val\_accuracy: 0.4953 - val\_binary\_accuracy: 0.4953

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.9919 - accuracy: 0.3438 - binary\_accuracy: 0.34 - ETA: 0s - loss: 0.8072 - accuracy: 0.4688 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.8172 - accuracy: 0.4744 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.8118 - accuracy: 0.4750 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.8132 - accuracy: 0.4770 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.8116 - accuracy: 0.4782 - binary\_accuracy: 0.47 - 0s 122us/sample - loss: 0.8104 - accuracy: 0.4783 - binary\_accuracy: 0.4783 - val\_loss: 0.8164 - val\_accuracy: 0.5018 - val\_binary\_accuracy: 0.5018

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 1.0000 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.8319 - accuracy: 0.4851 - binary\_accuracy: 0.48 - ETA: 0s - loss: 0.8214 - accuracy: 0.4805 - binary\_accuracy: 0.48 - ETA: 0s - loss: 0.8073 - accuracy: 0.4916 - binary\_accuracy: 0.49 - ETA: 0s - loss: 0.8012 - accuracy: 0.4823 - binary\_accuracy: 0.48 - ETA: 0s - loss: 0.8040 - accuracy: 0.4807 - binary\_accuracy: 0.48 - 0s 155us/sample - loss: 0.8050 - accuracy: 0.4832 - binary\_accuracy: 0.4832 - val\_loss: 0.8112 - val\_accuracy: 0.5040 - val\_binary\_accuracy: 0.5040

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.7132 - accuracy: 0.5312 - binary\_accuracy: 0.53 - ETA: 0s - loss: 0.8309 - accuracy: 0.4458 - binary\_accuracy: 0.44 - ETA: 0s - loss: 0.8025 - accuracy: 0.4706 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.7988 - accuracy: 0.4778 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.7930 - accuracy: 0.4860 - binary\_accuracy: 0.48 - ETA: 0s - loss: 0.7976 - accuracy: 0.4840 - binary\_accuracy: 0.48 - ETA: 0s - loss: 0.7997 - accuracy: 0.4842 - binary\_accuracy: 0.48 - 1s 160us/sample - loss: 0.7997 - accuracy: 0.4870 - binary\_accuracy: 0.4870 - val\_loss: 0.8061 - val\_accuracy: 0.5083 - val\_binary\_accuracy: 0.5083

Epoch 7/10

1312/3220 [=====>.....] - ETA: 0s - loss: 0.8565 - accuracy: 0.4375 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.8094 - accuracy: 0.4700 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.8189 - accuracy: 0.4710 - binary\_accuracy: 0.47

**limit\_output extension: Maximum message size of 10000 exceeded with 10070 characters**

```
In [23]: best_models = classtuner.get_best_models(num_models = 2)
```

```
In [24]: top_model = best_models[0]
```

```
In [25]: ### Use our best model and reverify it against our test data  
  
spam_history = top_model.fit(X_train_standard, y_train,  
                             validation_data = (X_test_standard, y_test), epochs = 50, verbose = 0)
```

```
In [26]: train_acc = top_model.evaluate(X_train_standard, y_train, verbose=0)  
test_acc = top_model.evaluate(X_test_standard, y_test, verbose=0)
```

```
In [29]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(spam_history.history['accuracy'], label = "Training Data Accuracy")
plt.plot(spam_history.history['val_accuracy'], label = "Testing Data Accuracy")

## Give some labels and title
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Title and Legend
plt.title("Model Accuracy Comparison - Training & Test Data")
plt.legend()

## Show the graph
plt.show()
```



**Using sklearn metrics library, I can determine how well the model is performing by generating ROC, Precision, and Recall curves**

```
In [30]: ## Create three empty lists to store my precision, recall, and average precision scores  
precision, recall, average_precision = [], [], []  
  
## Get probabilities for our labels  
probas_ = top_model.predict_proba(X_test)
```

```
In [69]: ## Create my false positive rate, true positive rate, and threshold using my test data  
fpr, tpr, thresholds = roc_curve(y_test, probas_, pos_label = 1)  
  
## Create precision and recall scores to plot with  
precision_score, recall_score, _ = precision_recall_curve(y_test, probas_)  
  
## Calculate the overall AUC for the model  
auc = np.trapz(tpr, fpr)  
  
## Save the average precision for our model  
avg_precision = average_precision_score(y_test, probas_)
```



```
In [41]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Create a new figure to plot
plt.figure(figsize= (15, 10))

lw = 2

## Draw the line for my fpr and tpr
plt.plot(fpr, tpr, color = 'darkorange',
         label = 'ROC Curve (area = %0.2f)' % auc)


## Put in a line to demonstrate blind luck
plt.plot([0, 1], [0, 1], color = 'navy', linestyle = '--', label = 'Luck')

## Set the limits of the plot for better visualization
plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.05])

## Set labels for x and y
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

## Set a title and legend
plt.title('ROC Curve')
plt.legend(loc = 'lower right')

## Show the curve!
plt.show()
```



```
In [77]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Create a new figure to plot
plt.figure(figsize= (15, 10))

lw = 2

## Plot the model's Precision Recall Curve
plt.plot(precision_score, recall_score, label='Precision-Recall curve (area = {})'.format(round(avg_precision, 4)))


## Put in a line to demonstrate blind luck
plt.plot([0, 1], [0, 1], color = 'black', linestyle = '--', label = 'Luck')

## Set the limits so they start at zero
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])

## Set labels for x and y
plt.xlabel('Recall')
plt.ylabel('Precision')

## Set a title and legend
plt.title('ROC Curve')
plt.legend(loc = 'lower right')

## Show the curve!
plt.show()
```



**I will now be building a cost matrix and visualization for this; again this is to visualize and assess how the model is "punishing" mistakes.**

```
In [61]: ## Initialize different thresholds and Costs that will be tied to those thresholds  
thresholds = np.linspace(0, 1.0, num=21)  
  
## Generate three generic Cost List for each matrix  
Cost_List = np.linspace(0, 1.0, num=21)  
  
## Build cost matrix for misidentifying spam  
cost_matrix = np.array([[0, 1], [10, 0]])
```

```
In [64]: ## Show the cost matrix  
cost_matrix
```

```
Out[64]: array([[ 0,  1],  
               [10,  0]])
```

```
In [62]: ## Set the index to zero

index = 0

### Loop through our different thresholds and calculate the misclassification cost at each interval
for t in thresholds:

    predict_thre = np.where(probas_ > t, 1, 0) ## prediction based on the preset threshold
    clf_matrix = confusion_matrix(y_test, predict_thre)
    Cost_List[index] = clf_matrix[0][0]*cost_matrix[0][0] +clf_matrix[0][1]*cost_matrix[0][1] +clf_matrix[1][0]*cost_ma
trix[1][0] +clf_matrix[1][1]*cost_matrix[1][1]
    index+=1


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot each Cost Line individually
plt.plot(thresholds, Cost_List, 'r--', label = "Average Cost for Misclassification")

## Give some labels
plt.xlabel("thresholds")
plt.ylabel("Cost")

## Title and Legend
plt.title("Cost Matrix")
plt.legend(loc = 'right')

## Show the Cost Matrix Analysis
plt.show()
```



**I am going to try one more model, using a Linear activator instead. We'll see if we can improve the accuracy of the model.**

In [120]: *## First, we build a function to actually put together our model - with change to use a linear activation instead*

```
def build_class_model(hp):  
  
    ## Base layer  
    cmodel = Sequential()  
  
    ## Add first layer, test with 30 up to 100 layers  
    cmodel.add(layers.Dense(units = hp.Int("units",  
                                         min_value = 30,  
                                         max_value = 100,  
                                         step = 2),  
                        ## Use same initializer as model above  
                        activation = "linear", kernel_initializer = "he_uniform"))  
  
    ## Add target layer using linear regressor  
    cmodel.add(layers.Dense(1, activation = "linear"))  
  
    ## Use different learning rates to test the model  
    cmodel.compile(  
        optimizer = keras.optimizers.SGD(  
            hp.Choice("learning_rate", values = [1e-2, 1e-3, 1e-4])),  
  
        ## Use same loss function as before  
        loss = "binary_crossentropy",  
  
        ## Test using Accuracy  
        metrics = ["accuracy", "binary_accuracy"])  
  
    ## Return completed model  
    return cmodel
```



In [121]: *## Build hyper tuner, will perform cross validation for us*

*### Use RandomSearch*

```
classtuner = RandomSearch(  
  
    ## Use model function from above  
    build_class_model,  
  
    ## What is objective function? Using Loss here  
    objective = "val_accuracy",  
  
    ## Set number of trials  
    max_trials = 5,  
  
    ## Number of executions  
    executions_per_trial = 3,  
  
    ## Set to short dir path  
    directory = "C:\\"  
  
)
```

```
In [122]: classtuner.search(X_train_standard, y_train, epochs = 10,  
                           validation_data = (X_test_standard, y_test))
```

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 41s - loss: 0.6667 - accuracy: 0.5938 - binary\_accuracy: 0.593 - ETA: 1s - loss: 0.6533 - accuracy: 0.5960 - binary\_accuracy: 0.596 - ETA: 0s - loss: 0.6518 - accuracy: 0.6094 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6526 - accuracy: 0.6068 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6480 - accuracy: 0.6157 - binary\_accuracy: 0.61 - 1s 266us/sample - loss: 0.6494 - accuracy: 0.6140 - binary\_accuracy: 0.6140 - val\_loss: 0.6507 - val\_accuracy: 0.6119 - val\_binary\_accuracy: 0.6119

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.7399 - accuracy: 0.4688 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.6469 - accuracy: 0.6175 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6442 - accuracy: 0.6263 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6481 - accuracy: 0.6174 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6484 - accuracy: 0.6159 - binary\_accuracy: 0.61 - 0s 100us/sample - loss: 0.6489 - accuracy: 0.6146 - binary\_accuracy: 0.6146 - val\_loss: 0.6502 - val\_accuracy: 0.6133 - val\_binary\_accuracy: 0.6133

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.6465 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6462 - accuracy: 0.6116 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6496 - accuracy: 0.6079 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6456 - accuracy: 0.6197 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6484 - accuracy: 0.6144 - binary\_accuracy: 0.61 - 0s 96us/sample - loss: 0.6485 - accuracy: 0.6146 - binary\_accuracy: 0.6146 - val\_loss: 0.6498 - val\_accuracy: 0.6140 - val\_binary\_accuracy: 0.6140

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.6453 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6439 - accuracy: 0.6262 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6454 - accuracy: 0.6256 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6476 - accuracy: 0.6192 - binary\_accuracy: 0.61 - 0s 91us/sample - loss: 0.6480 - accuracy: 0.6143 - binary\_accuracy: 0.6143 - val\_loss: 0.6494 - val\_accuracy: 0.6148 - val\_binary\_accuracy: 0.6148

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.5964 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6447 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6472 - accuracy: 0.6152 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6465 - accuracy: 0.6170 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6463 - accuracy: 0.6170 - binary\_accuracy: 0.61 - 0s 92us/sample - loss: 0.6476 - accuracy: 0.6137 - binary\_accuracy: 0.6137 - val\_loss: 0.6489 - val\_accuracy: 0.6140 - val\_binary\_accuracy: 0.6140

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.6227 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6404 - accuracy: 0.6377 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6496 - accuracy: 0.6138 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6491 - accuracy: 0.6125 - binary\_accuracy: 0.61 - 0s 84us/sample - loss: 0.6472 - accuracy: 0.6140 - binary\_accuracy: 0.6140 - val\_loss: 0.6485 - val\_accuracy: 0.6140 - val\_binary\_accuracy: 0.6140

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.6752 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6449 - accuracy: 0.6224 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6468 - accuracy: 0.6100 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6450 - accuracy: 0.6187 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6474 - accuracy: 0.6138 - binary\_accuracy: 0.61 - 0s 92us/sample - loss: 0.6467 - accuracy: 0.6146 - binary\_accuracy: 0.6146 - val\_loss: 0.6481 - val\_accuracy: 0.6148 - val\_binary\_accuracy: 0.6148

Epoch 8/10

3220/3220 [=====] - ETA: 0s - loss: 0.5999 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6559 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6480 - accuracy: 0.6114 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6449 - accuracy: 0.6187 - binary\_accuracy: 0.61 - 0s 84us/sample - loss: 0.6463 - accuracy: 0.6152 - binary\_accuracy: 0.6152 - val\_loss: 0.6477 - val\_accuracy: 0.6148 - val\_binary\_accuracy: 0.6148

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.5928 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6475 - accuracy: 0.6274 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6421 - accuracy: 0.6284 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6429 - accuracy: 0.6258 - binary\_accuracy: 0.62 - 0s 92us/sample - loss: 0.6458 - accuracy: 0.6149 - binary\_accuracy: 0.6149 - val\_loss: 0.6473 - val\_accuracy: 0.6155 - val\_binary\_accuracy: 0.6155  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6615 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6335 - accuracy: 0.6525 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6417 - accuracy: 0.6289 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6438 - accuracy: 0.6206 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6455 - accuracy: 0.6160 - binary\_accuracy: 0.61 - 0s 99us/sample - loss: 0.6454 - accuracy: 0.6149 - binary\_accuracy: 0.6149 - val\_loss: 0.6468 - val\_accuracy: 0.6162 - val\_binary\_accuracy: 0.6162  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 46s - loss: 0.7933 - accuracy: 0.5000 - binary\_accuracy: 0.500 - ETA: 1s - loss: 0.6857 - accuracy: 0.6150 - binary\_accuracy: 0.615 - ETA: 0s - loss: 0.6858 - accuracy: 0.6165 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6926 - accuracy: 0.6071 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6952 - accuracy: 0.6001 - binary\_accuracy: 0.60 - 1s 268us/sample - loss: 0.6949 - accuracy: 0.6009 - binary\_accuracy: 0.6009 - val\_loss: 0.6929 - val\_accuracy: 0.6025 - val\_binary\_accuracy: 0.6025  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7374 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.7084 - accuracy: 0.5775 - binary\_accuracy: 0.57 - ETA: 0s - loss: 0.6935 - accuracy: 0.5987 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6908 - accuracy: 0.6042 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6947 - accuracy: 0.5988 - binary\_accuracy: 0.59 - 0s 108us/sample - loss: 0.6938 - accuracy: 0.6003 - binary\_accuracy: 0.6003 - val\_loss: 0.6919 - val\_accuracy: 0.6032 - val\_binary\_accuracy: 0.6032  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6361 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.7066 - accuracy: 0.5872 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6988 - accuracy: 0.5943 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6906 - accuracy: 0.6014 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6918 - accuracy: 0.6016 - binary\_accuracy: 0.60 - 0s 92us/sample - loss: 0.6928 - accuracy: 0.6000 - binary\_accuracy: 0.6000 - val\_loss: 0.6909 - val\_accuracy: 0.6032 - val\_binary\_accuracy: 0.6032  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7541 - accuracy: 0.5312 - binary\_accuracy: 0.53 - ETA: 0s - loss: 0.6831 - accuracy: 0.6131 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6864 - accuracy: 0.6062 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6881 - accuracy: 0.6072 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6927 - accuracy: 0.6001 - binary\_accuracy: 0.60 - 0s 99us/sample - loss: 0.6918 - accuracy: 0.6003 - binary\_accuracy: 0.6003 - val\_loss: 0.6899 - val\_accuracy: 0.6017 - val\_binary\_accuracy: 0.6017  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6411 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6631 - accuracy: 0.6288 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6756 - accuracy: 0.6170 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6810 - accuracy: 0.6105 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6903 - accuracy: 0.6003 - binary\_accuracy: 0.60 - 0s 95us/sample - loss: 0.6908 - accuracy: 0.5997 - binary\_accuracy: 0.5997 - val\_loss: 0.6889 - val\_accuracy: 0.6010 - val\_binary\_accuracy: 0.6010  
Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7421 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6894 - accuracy: 0.5981 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6877 - accuracy: 0.6023 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6827 - accuracy: 0.6079 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6896 - accuracy: 0.5993 - binary\_accuracy: 0.59 - 0s 93us/sample - loss: 0.6898 - accuracy: 0.6000 - binary\_accuracy: 0.6000 - val\_loss: 0.6880 - val\_accuracy: 0.6010 - val\_binary\_accuracy: 0.6010  
Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.6851 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6980 - accuracy: 0.5925 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6895 - accuracy: 0.6012 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6876 - accuracy: 0.6039 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6898 - accuracy: 0.5981 - binary\_accuracy: 0.59 - 0s 88us/sample - loss: 0.6888 - accuracy: 0.5997 - binary\_accuracy: 0.5997 - val\_loss: 0.6871 - val\_accuracy: 0.6003 - val\_binary\_accuracy: 0.6003  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7333 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6812 - accuracy: 0.6080 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6862 - accuracy: 0.6050 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6896 - accuracy: 0.5977 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6876 - accuracy: 0.6000 - binary\_accuracy: 0.60 - 0s 89us/sample - loss: 0.6879 - accuracy: 0.5994 - binary\_accuracy: 0.5994 - val\_loss: 0.6862 - val\_accuracy: 0.6003 - val\_binary\_accuracy: 0.6003  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6300 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6676 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6791 - accuracy: 0.6111 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6843 - accuracy: 0.6028 - binary\_accuracy: 0.60 - 0s 81us/sample - loss: 0.6870 - accuracy: 0.5997 - binary\_accuracy: 0.5997 - val\_loss: 0.6853 - val\_accuracy: 0.6003 - val\_binary\_accuracy: 0.6003  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6740 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6673 - accuracy: 0.6115 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6759 - accuracy: 0.6062 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6797 - accuracy: 0.6029 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6851 - accuracy: 0.5992 - binary\_accuracy: 0.59 - 0s 99us/sample - loss: 0.6861 - accuracy: 0.5988 - binary\_accuracy: 0.5988 - val\_loss: 0.6845 - val\_accuracy: 0.6003 - val\_binary\_accuracy: 0.6003  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 51s - loss: 1.3191 - accuracy: 0.4062 - binary\_accuracy: 0.406 - ETA: 1s - loss: 1.3215 - accuracy: 0.4152 - binary\_accuracy: 0.415 - ETA: 0s - loss: 1.3374 - accuracy: 0.4042 - binary\_accuracy: 0.40 - ETA: 0s - loss: 1.3485 - accuracy: 0.3962 - binary\_accuracy: 0.39 - 1s 285us/sample - loss: 1.3541 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.3440 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 1.5227 - accuracy: 0.3125 - binary\_accuracy: 0.31 - ETA: 0s - loss: 1.2931 - accuracy: 0.4139 - binary\_accuracy: 0.41 - ETA: 0s - loss: 1.3312 - accuracy: 0.3940 - binary\_accuracy: 0.39 - 0s 55us/sample - loss: 1.3299 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.3201 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 1.4284 - accuracy: 0.3438 - binary\_accuracy: 0.34 - ETA: 0s - loss: 1.2620 - accuracy: 0.4231 - binary\_accuracy: 0.42 - ETA: 0s - loss: 1.3065 - accuracy: 0.3944 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.3089 - accuracy: 0.3927 - binary\_accuracy: 0.39 - 0s 82us/sample - loss: 1.3062 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.2967 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 1.2505 - accuracy: 0.4062 - binary\_accuracy: 0.40 - ETA: 0s - loss: 1.2633 - accuracy: 0.4090 - binary\_accuracy: 0.40 - ETA: 0s - loss: 1.2964 - accuracy: 0.3923 - binary\_accuracy: 0.39 - 0s 65us/sample - loss: 1.2832 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.2739 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.9407 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 1.3229 - accuracy: 0.3663 - binary\_accuracy: 0.36 - ETA: 0s - loss: 1.2737 - accuracy: 0.3883 - binary\_accuracy: 0.38 - 0s 58us/sample - loss: 1.2606 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.2517 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939

Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 1.4851 - accuracy: 0.2812 - binary\_accuracy: 0.28 - ETA: 0s - loss: 1.2588 - accuracy: 0.3906 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2517 - accuracy: 0.3877 - binary\_accuracy: 0.38 - 0s 54us/sample - loss: 1.2387 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.2300 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 1.0662 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 1.2263 - accuracy: 0.3956 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2244 - accuracy: 0.3946 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2266 - accuracy: 0.3909 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2257 - accuracy: 0.3904 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2160 - accuracy: 0.3960 - binary\_accuracy: 0.39 - 0s 120us/sample - loss: 1.2174 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.2089 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 1.3568 - accuracy: 0.3438 - binary\_accuracy: 0.34 - ETA: 0s - loss: 1.1989 - accuracy: 0.3940 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.2141 - accuracy: 0.3862 - binary\_accuracy: 0.38 - ETA: 0s - loss: 1.1992 - accuracy: 0.3931 - binary\_accuracy: 0.39 - 0s 86us/sample - loss: 1.1966 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.1884 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 1.2619 - accuracy: 0.3125 - binary\_accuracy: 0.31 - ETA: 0s - loss: 1.1704 - accuracy: 0.4004 - binary\_accuracy: 0.40 - ETA: 0s - loss: 1.1905 - accuracy: 0.3878 - binary\_accuracy: 0.38 - 0s 56us/sample - loss: 1.1764 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.1684 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.9287 - accuracy: 0.5312 - binary\_accuracy: 0.53 - ETA: 0s - loss: 1.1727 - accuracy: 0.3917 - binary\_accuracy: 0.39 - ETA: 0s - loss: 1.1539 - accuracy: 0.3982 - binary\_accuracy: 0.39 - 0s 54us/sample - loss: 1.1567 - accuracy: 0.3941 - binary\_accuracy: 0.3941 - val\_loss: 1.1491 - val\_accuracy: 0.3939 - val\_binary\_accuracy: 0.3939

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.0001

|units: 36

|Score: 0.5377745032310486

|Best step: 0

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 1:03 - loss: 0.7872 - accuracy: 0.3750 - binary\_accuracy: 0.37 - ETA: 1s - loss: 0.7291 - accuracy: 0.4256 - binary\_accuracy: 0.4256 - ETA: 0s - loss: 0.7011 - accuracy: 0.5022 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.6844 - accuracy: 0.5471 - binary\_accuracy: 0.54 - 1s 328us/sample - loss: 0.6797 - accuracy: 0.5593 - binary\_accuracy: 0.5593 - val\_loss: 0.6360 - val\_accuracy: 0.6510 - val\_binary\_accuracy: 0.6510

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.6154 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6256 - accuracy: 0.6643 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6201 - accuracy: 0.6668 - binary\_accuracy: 0.66 - 0s 75us/sample - loss: 0.6181 - accuracy: 0.6674 - binary\_accuracy: 0.6674 - val\_loss: 0.6000 - val\_accuracy: 0.6843 - val\_binary\_accuracy: 0.6843

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.6506 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.5961 - accuracy: 0.7018 - binary\_accuracy: 0.70 - ETA: 0s - loss: 0.5907 - accuracy: 0.7024 - binary\_accuracy: 0.70 - ETA: 0s - loss: 0.5860 - accuracy: 0.7070 - binary\_accuracy: 0.70 - 0s 81us/sample - loss: 0.5851 - accuracy: 0.7090 - binary\_accuracy: 0.7090 - val\_loss: 0.5689 - val\_accuracy: 0.7234 - val\_binary\_accuracy: 0.7234

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.4946 - accuracy: 0.8125 - binary\_accuracy: 0.81 - ETA: 0s - loss: 0.5595 - accuracy: 0.7393 - binary\_accuracy: 0.73 - ETA: 0s - loss: 0.5562 - accuracy: 0.7447 - binary\_accuracy: 0.74 - 0s 62us/sample - loss: 0.5550 - accuracy: 0.7472 - binary\_accuracy: 0.7472 - val\_loss: 0.5401 - val\_accuracy: 0.7610 - val\_binary\_accuracy: 0.7610

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.5856 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA: 0s - loss: 0.5284 - accuracy: 0.7897 - binary\_accuracy: 0.78 - ETA: 0s - loss: 0.5273 - accuracy: 0.7824 - binary\_accuracy: 0.78 - 0s 65us/sample - loss: 0.5272 - accuracy: 0.7835 - binary\_accuracy: 0.7835 - val\_loss: 0.5131 - val\_accuracy: 0.7907 - val\_binary\_accuracy: 0.7907

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.5006 - accuracy: 0.7812 - binary\_accuracy: 0.78 - ETA: 0s - loss: 0.5141 - accuracy: 0.8003 - binary\_accuracy: 0.80 - ETA: 0s - loss: 0.5032 - accuracy: 0.8069 - binary\_accuracy: 0.80 - 0s 62us/sample - loss: 0.5012 - accuracy: 0.8106 - binary\_accuracy: 0.8106 - val\_loss: 0.4881 - val\_accuracy: 0.8219 - val\_binary\_accuracy: 0.8219

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.4646 - accuracy: 0.8125 - binary\_accuracy: 0.81 - ETA: 0s - loss: 0.4827 - accuracy: 0.8326 - binary\_accuracy: 0.83 - ETA: 0s - loss: 0.4749 - accuracy: 0.8355 - binary\_accuracy: 0.83 - 0s 61us/sample - loss: 0.4771 - accuracy: 0.8326 - binary\_accuracy: 0.8326 - val\_loss: 0.4649 - val\_accuracy: 0.8371 - val\_binary\_accuracy: 0.8371

Epoch 8/10

3220/3220 [=====] - ETA: 0s - loss: 0.4548 - accuracy: 0.8750 - binary\_accuracy: 0.87 - ETA: 0s - loss: 0.4687 - accuracy: 0.8396 - binary\_accuracy: 0.83 - ETA: 0s - loss: 0.4581 - accuracy: 0.8438 - binary\_accuracy: 0.84 - 0s 63us/sample - loss: 0.4550 - accuracy: 0.8457 - binary\_accuracy: 0.8457 - val\_loss: 0.4436 - val\_accuracy: 0.8487 - val\_binary\_accuracy: 0.8487

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.4510 - accuracy: 0.8125 - binary\_accuracy: 0.81 - ETA: 0s - loss: 0.4436 - accuracy: 0.8551 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4339 - accuracy: 0.8542 - binary\_accuracy: 0.85 - 0s 64us/sample - loss: 0.4348 - accuracy: 0.8522 - binary\_accuracy: 0.8522 - val\_loss: 0.4241 - val\_accuracy: 0.8610 - val\_binary\_accuracy: 0.8610

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.3608 - accuracy: 0.9062 - binary\_accuracy: 0.90 - ETA:  
0s - loss: 0.4230 - accuracy: 0.8559 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4185 - accuracy: 0.8645 - binary\_accu  
racy: 0.86 - 0s 64us/sample - loss: 0.4166 - accuracy: 0.8612 - binary\_accuracy: 0.8612 - val\_loss: 0.4065 - val\_accu  
racy: 0.8660 - val\_binary\_accuracy: 0.8660  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 38s - loss: 0.7170 - accuracy: 0.5000 - binary\_accuracy: 0.500 - ET  
A: 0s - loss: 0.6976 - accuracy: 0.5076 - binary\_accuracy: 0.507 - ETA: 0s - loss: 0.6758 - accuracy: 0.5602 - binary\_  
accuracy: 0.56 - 1s 203us/sample - loss: 0.6687 - accuracy: 0.5724 - binary\_accuracy: 0.5724 - val\_loss: 0.6349 - val\_  
accuracy: 0.6365 - val\_binary\_accuracy: 0.6365  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6122 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA:  
0s - loss: 0.6289 - accuracy: 0.6316 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6196 - accuracy: 0.6433 - binary\_accu  
racy: 0.64 - 0s 60us/sample - loss: 0.6155 - accuracy: 0.6503 - binary\_accuracy: 0.6503 - val\_loss: 0.5970 - val\_accu  
racy: 0.6778 - val\_binary\_accuracy: 0.6778  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6435 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA:  
0s - loss: 0.5871 - accuracy: 0.6788 - binary\_accuracy: 0.67 - ETA: 0s - loss: 0.5809 - accuracy: 0.6860 - binary\_accu  
racy: 0.68 - 0s 55us/sample - loss: 0.5798 - accuracy: 0.6904 - binary\_accuracy: 0.6904 - val\_loss: 0.5635 - val\_accu  
racy: 0.7306 - val\_binary\_accuracy: 0.7306  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6207 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA:  
0s - loss: 0.5596 - accuracy: 0.7287 - binary\_accuracy: 0.72 - ETA: 0s - loss: 0.5503 - accuracy: 0.7444 - binary\_accu  
racy: 0.74 - 0s 55us/sample - loss: 0.5476 - accuracy: 0.7466 - binary\_accuracy: 0.7466 - val\_loss: 0.5328 - val\_accu  
racy: 0.7690 - val\_binary\_accuracy: 0.7690  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5372 - accuracy: 0.7500 - binary\_accuracy: 0.75 - ETA:  
0s - loss: 0.5344 - accuracy: 0.7642 - binary\_accuracy: 0.76 - ETA: 0s - loss: 0.5216 - accuracy: 0.7736 - binary\_accu  
racy: 0.77 - 0s 58us/sample - loss: 0.5181 - accuracy: 0.7786 - binary\_accuracy: 0.7786 - val\_loss: 0.5045 - val\_accu  
racy: 0.8009 - val\_binary\_accuracy: 0.8009  
Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4624 - accuracy: 0.8750 - binary\_accuracy: 0.87 - ETA:  
0s - loss: 0.4975 - accuracy: 0.8139 - binary\_accuracy: 0.81 - ETA: 0s - loss: 0.4931 - accuracy: 0.8147 - binary\_accu  
racy: 0.81 - 0s 53us/sample - loss: 0.4910 - accuracy: 0.8171 - binary\_accuracy: 0.8171 - val\_loss: 0.4785 - val\_accu  
racy: 0.8291 - val\_binary\_accuracy: 0.8291  
Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4806 - accuracy: 0.7500 - binary\_accuracy: 0.75 - ETA:  
0s - loss: 0.4760 - accuracy: 0.8321 - binary\_accuracy: 0.83 - ETA: 0s - loss: 0.4706 - accuracy: 0.8367 - binary\_accu  
racy: 0.83 - 0s 56us/sample - loss: 0.4661 - accuracy: 0.8373 - binary\_accuracy: 0.8373 - val\_loss: 0.4547 - val\_accu  
racy: 0.8465 - val\_binary\_accuracy: 0.8465  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5181 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.4535 - accuracy: 0.8408 - binary\_accuracy: 0.84 - ETA: 0s - loss: 0.4484 - accuracy: 0.8445 - binary\_accu  
racy: 0.84 - 0s 54us/sample - loss: 0.4436 - accuracy: 0.8512 - binary\_accuracy: 0.8512 - val\_loss: 0.4330 - val\_accu  
racy: 0.8631 - val\_binary\_accuracy: 0.8631  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4183 - accuracy: 0.9062 - binary\_accuracy: 0.90 - ETA:  
0s - loss: 0.4280 - accuracy: 0.8516 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4228 - accuracy: 0.8608 - binary\_accu



acy: 0.86 - 0s 61us/sample - loss: 0.4232 - accuracy: 0.8602 - binary\_accuracy: 0.8602 - val\_loss: 0.4135 - val\_accu  
acy: 0.8812 - val\_binary\_accuracy: 0.8812  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4215 - accuracy: 0.8750 - binary\_accuracy: 0.87 - ETA:  
0s - loss: 0.4171 - accuracy: 0.8601 - binary\_accuracy: 0.86 - ETA: 0s - loss: 0.4086 - accuracy: 0.8706 - binary\_accu  
racy: 0.87 - 0s 63us/sample - loss: 0.4048 - accuracy: 0.8714 - binary\_accuracy: 0.8714 - val\_loss: 0.3959 - val\_accu  
acy: 0.8827 - val\_binary\_accuracy: 0.8827  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 37s - loss: 0.6464 - accuracy: 0.5625 - binary\_accuracy: 0.562 - ET  
A: 0s - loss: 0.6898 - accuracy: 0.5281 - binary\_accuracy: 0.528 - ETA: 0s - loss: 0.6607 - accuracy: 0.6171 - binary\_  
accuracy: 0.61 - 1s 201us/sample - loss: 0.6492 - accuracy: 0.6373 - binary\_accuracy: 0.6373 - val\_loss: 0.6057 - val\_  
accuracy: 0.7060 - val\_binary\_accuracy: 0.7060  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6236 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.6036 - accuracy: 0.7047 - binary\_accuracy: 0.70 - ETA: 0s - loss: 0.5962 - accuracy: 0.7093 - binary\_accu  
racy: 0.70 - 0s 58us/sample - loss: 0.5921 - accuracy: 0.7127 - binary\_accuracy: 0.7127 - val\_loss: 0.5755 - val\_accu  
acy: 0.7270 - val\_binary\_accuracy: 0.7270  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5431 - accuracy: 0.7812 - binary\_accuracy: 0.78 - ETA:  
0s - loss: 0.5663 - accuracy: 0.7301 - binary\_accuracy: 0.73 - ETA: 0s - loss: 0.5694 - accuracy: 0.7254 - binary\_accu  
racy: 0.72 - 0s 56us/sample - loss: 0.5643 - accuracy: 0.7376 - binary\_accuracy: 0.7376 - val\_loss: 0.5491 - val\_accu  
acy: 0.7748 - val\_binary\_accuracy: 0.7748  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5424 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.5390 - accuracy: 0.7686 - binary\_accuracy: 0.76 - ETA: 0s - loss: 0.5385 - accuracy: 0.7673 - binary\_accu  
racy: 0.76 - 0s 57us/sample - loss: 0.5384 - accuracy: 0.7699 - binary\_accuracy: 0.7699 - val\_loss: 0.5243 - val\_accu  
acy: 0.8038 - val\_binary\_accuracy: 0.8038  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5244 - accuracy: 0.8125 - binary\_accuracy: 0.81 - ETA:  
0s - loss: 0.5150 - accuracy: 0.7904 - binary\_accuracy: 0.79 - ETA: 0s - loss: 0.5163 - accuracy: 0.7948 - binary\_accu  
racy: 0.79 - 0s 61us/sample - loss: 0.5142 - accuracy: 0.7950 - binary\_accuracy: 0.7950 - val\_loss: 0.5007 - val\_accu  
acy: 0.8154 - val\_binary\_accuracy: 0.8154  
Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4907 - accuracy: 0.7812 - binary\_accuracy: 0.78 - ETA:  
0s - loss: 0.4933 - accuracy: 0.8087 - binary\_accuracy: 0.80 - ETA: 0s - loss: 0.4899 - accuracy: 0.8163 - binary\_accu  
racy: 0.81 - ETA: 0s - loss: 0.4906 - accuracy: 0.8132 - binary\_accuracy: 0.81 - 0s 77us/sample - loss: 0.4914 - accur  
acy: 0.8130 - binary\_accuracy: 0.8130 - val\_loss: 0.4789 - val\_accuracy: 0.8320 - val\_binary\_accuracy: 0.8320  
Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4571 - accuracy: 0.8438 - binary\_accuracy: 0.84 - ETA:  
0s - loss: 0.4771 - accuracy: 0.8380 - binary\_accuracy: 0.83 - ETA: 0s - loss: 0.4747 - accuracy: 0.8297 - binary\_accu  
racy: 0.82 - ETA: 0s - loss: 0.4735 - accuracy: 0.8273 - binary\_accuracy: 0.82 - ETA: 0s - loss: 0.4710 - accuracy: 0.  
8333 - binary\_accuracy: 0.83 - 0s 129us/sample - loss: 0.4702 - accuracy: 0.8342 - binary\_accuracy: 0.8342 - val\_loss:  
0.4584 - val\_accuracy: 0.8414 - val\_binary\_accuracy: 0.8414  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5210 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.4521 - accuracy: 0.8472 - binary\_accuracy: 0.84 - ETA: 0s - loss: 0.4509 - accuracy: 0.8416 - binary\_accu  
racy: 0.84 - ETA: 0s - loss: 0.4468 - accuracy: 0.8469 - binary\_accuracy: 0.84 - ETA: 0s - loss: 0.4506 - accuracy: 0.

8406 - binary\_accuracy: 0.84 - 0s 84us/sample - loss: 0.4507 - accuracy: 0.8404 - binary\_accuracy: 0.8404 - val\_loss: 0.4399 - val\_accuracy: 0.8530 - val\_binary\_accuracy: 0.8530  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.4704 - accuracy: 0.8750 - binary\_accuracy: 0.87 - ETA: 0s - loss: 0.4382 - accuracy: 0.8586 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4362 - accuracy: 0.8562 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4374 - accuracy: 0.8535 - binary\_accuracy: 0.85 - 0s 79us/sample - loss: 0.4329 - accuracy: 0.8556 - binary\_accuracy: 0.8556 - val\_loss: 0.4227 - val\_accuracy: 0.8573 - val\_binary\_accuracy: 0.8573  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.3628 - accuracy: 0.9062 - binary\_accuracy: 0.90 - ETA: 0s - loss: 0.4152 - accuracy: 0.8587 - binary\_accuracy: 0.85 - ETA: 0s - loss: 0.4176 - accuracy: 0.8623 - binary\_accuracy: 0.86 - ETA: 0s - loss: 0.4204 - accuracy: 0.8590 - binary\_accuracy: 0.85 - 0s 82us/sample - loss: 0.4166 - accuracy: 0.8621 - binary\_accuracy: 0.8621 - val\_loss: 0.4072 - val\_accuracy: 0.8602 - val\_binary\_accuracy: 0.8602

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.01

|units: 30

|Score: 0.8696596622467041

|Best step: 0

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 48s - loss: 0.7067 - accuracy: 0.4688 - binary\_accuracy: 0.468 - ETA: 0s - loss: 0.6944 - accuracy: 0.5393 - binary\_accuracy: 0.539 - ETA: 0s - loss: 0.6926 - accuracy: 0.5537 - binary\_accuracy: 0.55 - 1s 252us/sample - loss: 0.6897 - accuracy: 0.5658 - binary\_accuracy: 0.5658 - val\_loss: 0.6785 - val\_accuracy: 0.6046 - val\_binary\_accuracy: 0.6046

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.6742 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6824 - accuracy: 0.5958 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6767 - accuracy: 0.6006 - binary\_accuracy: 0.60 - 0s 63us/sample - loss: 0.6726 - accuracy: 0.6012 - binary\_accuracy: 0.6012 - val\_loss: 0.6663 - val\_accuracy: 0.6025 - val\_binary\_accuracy: 0.6025

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.6249 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA: 0s - loss: 0.6632 - accuracy: 0.6128 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6627 - accuracy: 0.6114 - binary\_accuracy: 0.61 - 0s 66us/sample - loss: 0.6628 - accuracy: 0.6075 - binary\_accuracy: 0.6075 - val\_loss: 0.6583 - val\_accuracy: 0.6083 - val\_binary\_accuracy: 0.6083

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.5659 - accuracy: 0.7500 - binary\_accuracy: 0.75 - ETA: 0s - loss: 0.6494 - accuracy: 0.6286 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6535 - accuracy: 0.6174 - binary\_accuracy: 0.61 - 0s 63us/sample - loss: 0.6555 - accuracy: 0.6137 - binary\_accuracy: 0.6137 - val\_loss: 0.6519 - val\_accuracy: 0.6169 - val\_binary\_accuracy: 0.6169

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.7032 - accuracy: 0.4375 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.6465 - accuracy: 0.6293 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6519 - accuracy: 0.6154 - binary\_accuracy: 0.61 - 0s 64us/sample - loss: 0.6492 - accuracy: 0.6193 - binary\_accuracy: 0.6193 - val\_loss: 0.6459 - val\_accuracy: 0.6213 - val\_binary\_accuracy: 0.6213

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.7038 - accuracy: 0.4688 - binary\_accuracy: 0.46 - ETA: 0s - loss: 0.6564 - accuracy: 0.5992 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6462 - accuracy: 0.6208 - binary\_accuracy: 0.62 - 0s 64us/sample - loss: 0.6432 - accuracy: 0.6233 - binary\_accuracy: 0.6233 - val\_loss: 0.6402 - val\_accuracy: 0.6256 - val\_binary\_accuracy: 0.6256

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.6414 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6389 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6381 - accuracy: 0.6270 - binary\_accuracy: 0.62 - 0s 66us/sample - loss: 0.6374 - accuracy: 0.6292 - binary\_accuracy: 0.6292 - val\_loss: 0.6347 - val\_accuracy: 0.6264 - val\_binary\_accuracy: 0.6264

Epoch 8/10

3220/3220 [=====] - ETA: 0s - loss: 0.6602 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6302 - accuracy: 0.6406 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6336 - accuracy: 0.6320 - binary\_accuracy: 0.63 - 0s 65us/sample - loss: 0.6318 - accuracy: 0.6339 - binary\_accuracy: 0.6339 - val\_loss: 0.6293 - val\_accuracy: 0.6307 - val\_binary\_accuracy: 0.6307

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.6353 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6264 - accuracy: 0.6406 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6232 - accuracy: 0.6462 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6245 - accuracy: 0.6412 - binary\_accuracy: 0.64 - 0s 101us/sample - loss: 0.6264 - accuracy: 0.6376 - binary\_accuracy: 0.6376 - val\_loss: 0.6241 - val\_accuracy: 0.6423 - val\_binary\_accuracy: 0.6423

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.6406 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA:  
0s - loss: 0.6236 - accuracy: 0.6335 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6242 - accuracy: 0.6348 - binary\_accu  
racy: 0.63 - 0s 67us/sample - loss: 0.6210 - accuracy: 0.6425 - binary\_accuracy: 0.6425 - val\_loss: 0.6189 - val\_accu  
racy: 0.6466 - val\_binary\_accuracy: 0.6466  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 37s - loss: 0.7521 - accuracy: 0.2812 - binary\_accuracy: 0.281 - ET  
A: 0s - loss: 0.7248 - accuracy: 0.4535 - binary\_accuracy: 0.453 - ETA: 0s - loss: 0.7165 - accuracy: 0.4769 - binary\_  
accuracy: 0.47 - 1s 208us/sample - loss: 0.7109 - accuracy: 0.4932 - binary\_accuracy: 0.4932 - val\_loss: 0.6998 - val\_  
accuracy: 0.5243 - val\_binary\_accuracy: 0.5243  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6522 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA:  
0s - loss: 0.6955 - accuracy: 0.5320 - binary\_accuracy: 0.53 - ETA: 0s - loss: 0.6905 - accuracy: 0.5522 - binary\_accu  
racy: 0.55 - 0s 59us/sample - loss: 0.6875 - accuracy: 0.5584 - binary\_accuracy: 0.5584 - val\_loss: 0.6837 - val\_accu  
racy: 0.5670 - val\_binary\_accuracy: 0.5670  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6458 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA:  
0s - loss: 0.6723 - accuracy: 0.5883 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6742 - accuracy: 0.5864 - binary\_accu  
racy: 0.58 - 0s 57us/sample - loss: 0.6750 - accuracy: 0.5829 - binary\_accuracy: 0.5829 - val\_loss: 0.6736 - val\_accu  
racy: 0.5851 - val\_binary\_accuracy: 0.5851  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6247 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.6728 - accuracy: 0.5841 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6708 - accuracy: 0.5837 - binary\_accu  
racy: 0.58 - 0s 56us/sample - loss: 0.6663 - accuracy: 0.5929 - binary\_accuracy: 0.5929 - val\_loss: 0.6657 - val\_accu  
racy: 0.5894 - val\_binary\_accuracy: 0.5894  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6193 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA:  
0s - loss: 0.6611 - accuracy: 0.6057 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6620 - accuracy: 0.5926 - binary\_accu  
racy: 0.59 - 0s 55us/sample - loss: 0.6590 - accuracy: 0.5994 - binary\_accuracy: 0.5994 - val\_loss: 0.6588 - val\_accu  
racy: 0.5988 - val\_binary\_accuracy: 0.5988  
Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6757 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA:  
0s - loss: 0.6494 - accuracy: 0.6190 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6523 - accuracy: 0.6052 - binary\_accu  
racy: 0.60 - 0s 56us/sample - loss: 0.6524 - accuracy: 0.6062 - binary\_accuracy: 0.6062 - val\_loss: 0.6524 - val\_accu  
racy: 0.5996 - val\_binary\_accuracy: 0.5996  
Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7108 - accuracy: 0.4375 - binary\_accuracy: 0.43 - ETA:  
0s - loss: 0.6449 - accuracy: 0.6193 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6461 - accuracy: 0.6129 - binary\_accu  
racy: 0.61 - 0s 54us/sample - loss: 0.6460 - accuracy: 0.6130 - binary\_accuracy: 0.6130 - val\_loss: 0.6462 - val\_accu  
racy: 0.6075 - val\_binary\_accuracy: 0.6075  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6611 - accuracy: 0.5312 - binary\_accuracy: 0.53 - ETA:  
0s - loss: 0.6395 - accuracy: 0.6265 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6400 - accuracy: 0.6198 - binary\_accu  
racy: 0.61 - 0s 55us/sample - loss: 0.6399 - accuracy: 0.6171 - binary\_accuracy: 0.6171 - val\_loss: 0.6402 - val\_accu  
racy: 0.6177 - val\_binary\_accuracy: 0.6177  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6671 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA:  
0s - loss: 0.6340 - accuracy: 0.6288 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6351 - accuracy: 0.6227 - binary\_accu

accuracy: 0.62 - 0s 56us/sample - loss: 0.6340 - accuracy: 0.6248 - binary\_accuracy: 0.6248 - val\_loss: 0.6343 - val\_accuracy: 0.6249 - val\_binary\_accuracy: 0.6249

Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6336 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6324 - accuracy: 0.6186 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6299 - accuracy: 0.6250 - binary\_accuracy: 0.62 - 0s 55us/sample - loss: 0.6282 - accuracy: 0.6295 - binary\_accuracy: 0.6295 - val\_loss: 0.6286 - val\_accuracy: 0.6314 - val\_binary\_accuracy: 0.6314

Train on 3220 samples, validate on 1381 samples

Epoch 1/10  
3220/3220 [=====] - ETA: 37s - loss: 0.6566 - accuracy: 0.5938 - binary\_accuracy: 0.593 - ETA: 0s - loss: 0.7723 - accuracy: 0.4492 - binary\_accuracy: 0.449 - ETA: 0s - loss: 0.7541 - accuracy: 0.4497 - binary\_accuracy: 0.44 - 1s 205us/sample - loss: 0.7448 - accuracy: 0.4531 - binary\_accuracy: 0.4531 - val\_loss: 0.7095 - val\_accuracy: 0.4707 - val\_binary\_accuracy: 0.4707

Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6670 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6994 - accuracy: 0.4789 - binary\_accuracy: 0.47 - ETA: 0s - loss: 0.6946 - accuracy: 0.4938 - binary\_accuracy: 0.49 - 0s 65us/sample - loss: 0.6916 - accuracy: 0.5065 - binary\_accuracy: 0.5065 - val\_loss: 0.6775 - val\_accuracy: 0.5489 - val\_binary\_accuracy: 0.5489

Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6214 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6755 - accuracy: 0.5675 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6730 - accuracy: 0.5820 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6725 - accuracy: 0.5840 - binary\_accuracy: 0.58 - 0s 79us/sample - loss: 0.6693 - accuracy: 0.5932 - binary\_accuracy: 0.5932 - val\_loss: 0.6620 - val\_accuracy: 0.6010 - val\_binary\_accuracy: 0.6010

Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6715 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6590 - accuracy: 0.6230 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6559 - accuracy: 0.6276 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6587 - accuracy: 0.6143 - binary\_accuracy: 0.61 - 0s 77us/sample - loss: 0.6575 - accuracy: 0.6177 - binary\_accuracy: 0.6177 - val\_loss: 0.6531 - val\_accuracy: 0.6119 - val\_binary\_accuracy: 0.6119

Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5823 - accuracy: 0.7500 - binary\_accuracy: 0.75 - ETA: 0s - loss: 0.6435 - accuracy: 0.6528 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6516 - accuracy: 0.6376 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6508 - accuracy: 0.6348 - binary\_accuracy: 0.63 - 0s 81us/sample - loss: 0.6497 - accuracy: 0.6345 - binary\_accuracy: 0.6345 - val\_loss: 0.6466 - val\_accuracy: 0.6235 - val\_binary\_accuracy: 0.6235

Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6313 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6495 - accuracy: 0.6406 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6479 - accuracy: 0.6339 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6434 - accuracy: 0.6401 - binary\_accuracy: 0.64 - 0s 77us/sample - loss: 0.6433 - accuracy: 0.6404 - binary\_accuracy: 0.6404 - val\_loss: 0.6406 - val\_accuracy: 0.6293 - val\_binary\_accuracy: 0.6293

Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5881 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6286 - accuracy: 0.6657 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6377 - accuracy: 0.6416 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6376 - accuracy: 0.6379 - binary\_accuracy: 0.63 - 0s 71us/sample - loss: 0.6374 - accuracy: 0.6398 - binary\_accuracy: 0.6398 - val\_loss: 0.6351 - val\_accuracy: 0.6358 - val\_binary\_accuracy: 0.6358

Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6500 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6242 - accuracy: 0.6736 - binary\_accuracy: 0.67 - ETA: 0s - loss: 0.6311 - accuracy: 0.6465 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6310 - accuracy: 0.6467 - binary\_accuracy: 0.64 - 0s 72us/sample - loss: 0.6317 - accuracy: 0.6457 - binary\_accuracy: 0.6457 - val\_loss: 0.6297 - val\_accuracy: 0.6430 - val\_binary\_accuracy: 0.6430

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.6511 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6325 - accuracy: 0.6464 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6294 - accuracy: 0.6467 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6255 - accuracy: 0.6525 - binary\_accuracy: 0.65 - 0s 79us/sample - loss: 0.6263 - accuracy: 0.6481 - binary\_accuracy: 0.6481 - val\_loss: 0.6245 - val\_accuracy: 0.6488 - val\_binary\_accuracy: 0.6488

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.5866 - accuracy: 0.7812 - binary\_accuracy: 0.78 - ETA: 0s - loss: 0.6230 - accuracy: 0.6447 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6232 - accuracy: 0.6538 - binary\_accuracy: 0.65 - 0s 64us/sample - loss: 0.6209 - accuracy: 0.6562 - binary\_accuracy: 0.6562 - val\_loss: 0.6194 - val\_accuracy: 0.6568 - val\_binary\_accuracy: 0.6568

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.001

|units: 82

|Score: 0.6449432969093323

|Best step: 0

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 39s - loss: 0.6979 - accuracy: 0.5625 - binary\_accuracy: 0.562 - ETA: 0s - loss: 0.6740 - accuracy: 0.5785 - binary\_accuracy: 0.578 - ETA: 0s - loss: 0.6711 - accuracy: 0.5962 - binary\_accuracy: 0.59 - 1s 219us/sample - loss: 0.6691 - accuracy: 0.6065 - binary\_accuracy: 0.6065 - val\_loss: 0.6622 - val\_accuracy: 0.6097 - val\_binary\_accuracy: 0.6097

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.6398 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6590 - accuracy: 0.6380 - binary\_accuracy: 0.63 - ETA: 0s - loss: 0.6557 - accuracy: 0.6461 - binary\_accuracy: 0.64 - 0s 67us/sample - loss: 0.6543 - accuracy: 0.6478 - binary\_accuracy: 0.6478 - val\_loss: 0.6513 - val\_accuracy: 0.6488 - val\_binary\_accuracy: 0.6488

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.6749 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6504 - accuracy: 0.6484 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6467 - accuracy: 0.6483 - binary\_accuracy: 0.64 - 0s 67us/sample - loss: 0.6458 - accuracy: 0.6500 - binary\_accuracy: 0.6500 - val\_loss: 0.6441 - val\_accuracy: 0.6575 - val\_binary\_accuracy: 0.6575

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.7146 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.6411 - accuracy: 0.6445 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6422 - accuracy: 0.6477 - binary\_accuracy: 0.64 - 0s 65us/sample - loss: 0.6395 - accuracy: 0.6506 - binary\_accuracy: 0.6506 - val\_loss: 0.6381 - val\_accuracy: 0.6604 - val\_binary\_accuracy: 0.6604

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.6351 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6392 - accuracy: 0.6417 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6342 - accuracy: 0.6511 - binary\_accuracy: 0.65 - 0s 64us/sample - loss: 0.6340 - accuracy: 0.6522 - binary\_accuracy: 0.6522 - val\_loss: 0.6327 - val\_accuracy: 0.6676 - val\_binary\_accuracy: 0.6676

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.6146 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6319 - accuracy: 0.6532 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6282 - accuracy: 0.6562 - binary\_accuracy: 0.65 - 0s 64us/sample - loss: 0.6288 - accuracy: 0.6547 - binary\_accuracy: 0.6547 - val\_loss: 0.6276 - val\_accuracy: 0.6734 - val\_binary\_accuracy: 0.6734

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.6420 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6247 - accuracy: 0.6600 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6255 - accuracy: 0.6570 - binary\_accuracy: 0.65 - 0s 65us/sample - loss: 0.6237 - accuracy: 0.6596 - binary\_accuracy: 0.6596 - val\_loss: 0.6226 - val\_accuracy: 0.6785 - val\_binary\_accuracy: 0.6785

Epoch 8/10

3220/3220 [=====] - ETA: 0s - loss: 0.6181 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6026 - accuracy: 0.6916 - binary\_accuracy: 0.69 - ETA: 0s - loss: 0.6124 - accuracy: 0.6713 - binary\_accuracy: 0.67 - 0s 67us/sample - loss: 0.6187 - accuracy: 0.6640 - binary\_accuracy: 0.6640 - val\_loss: 0.6178 - val\_accuracy: 0.6807 - val\_binary\_accuracy: 0.6807

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.6642 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6152 - accuracy: 0.6726 - binary\_accuracy: 0.67 - ETA: 0s - loss: 0.6155 - accuracy: 0.6675 - binary\_accuracy: 0.66 - 0s 66us/sample - loss: 0.6139 - accuracy: 0.6711 - binary\_accuracy: 0.6711 - val\_loss: 0.6130 - val\_accuracy: 0.6857 - val\_binary\_accuracy: 0.6857

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.5701 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.6220 - accuracy: 0.6639 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6110 - accuracy: 0.6788 - binary\_accu  
racy: 0.67 - 0s 70us/sample - loss: 0.6091 - accuracy: 0.6811 - binary\_accuracy: 0.6811 - val\_loss: 0.6082 - val\_accu  
racy: 0.6894 - val\_binary\_accuracy: 0.6894  
Train on 3220 samples, validate on 1381 samples  
Epoch 1/10  
3220/3220 [=====] - ETA: 36s - loss: 0.6163 - accuracy: 0.7500 - binary\_accuracy: 0.750 - ET  
A: 0s - loss: 0.6565 - accuracy: 0.6601 - binary\_accuracy: 0.660 - ETA: 0s - loss: 0.6560 - accuracy: 0.6547 - binary\_  
accuracy: 0.65 - 1s 200us/sample - loss: 0.6557 - accuracy: 0.6531 - binary\_accuracy: 0.6531 - val\_loss: 0.6464 - val\_  
accuracy: 0.6626 - val\_binary\_accuracy: 0.6626  
Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6451 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA:  
0s - loss: 0.6496 - accuracy: 0.6598 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6458 - accuracy: 0.6668 - binary\_accu  
racy: 0.66 - 0s 59us/sample - loss: 0.6483 - accuracy: 0.6590 - binary\_accuracy: 0.6590 - val\_loss: 0.6397 - val\_accu  
racy: 0.6749 - val\_binary\_accuracy: 0.6749  
Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5788 - accuracy: 0.8125 - binary\_accuracy: 0.81 - ETA:  
0s - loss: 0.6435 - accuracy: 0.6689 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6411 - accuracy: 0.6685 - binary\_accu  
racy: 0.66 - 0s 55us/sample - loss: 0.6420 - accuracy: 0.6630 - binary\_accuracy: 0.6630 - val\_loss: 0.6338 - val\_accu  
racy: 0.6734 - val\_binary\_accuracy: 0.6734  
Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6855 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA:  
0s - loss: 0.6418 - accuracy: 0.6584 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6380 - accuracy: 0.6684 - binary\_accu  
racy: 0.66 - 0s 55us/sample - loss: 0.6361 - accuracy: 0.6668 - binary\_accuracy: 0.6668 - val\_loss: 0.6282 - val\_accu  
racy: 0.6698 - val\_binary\_accuracy: 0.6698  
Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6305 - accuracy: 0.7188 - binary\_accuracy: 0.71 - ETA:  
0s - loss: 0.6434 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6342 - accuracy: 0.6635 - binary\_accu  
racy: 0.66 - 0s 54us/sample - loss: 0.6304 - accuracy: 0.6711 - binary\_accuracy: 0.6711 - val\_loss: 0.6228 - val\_accu  
racy: 0.6720 - val\_binary\_accuracy: 0.6720  
Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6347 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA:  
0s - loss: 0.6298 - accuracy: 0.6690 - binary\_accuracy: 0.66 - ETA: 0s - loss: 0.6262 - accuracy: 0.6732 - binary\_accu  
racy: 0.67 - 0s 55us/sample - loss: 0.6249 - accuracy: 0.6727 - binary\_accuracy: 0.6727 - val\_loss: 0.6175 - val\_accu  
racy: 0.6799 - val\_binary\_accuracy: 0.6799  
Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5858 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA:  
0s - loss: 0.6257 - accuracy: 0.6737 - binary\_accuracy: 0.67 - ETA: 0s - loss: 0.6219 - accuracy: 0.6795 - binary\_accu  
racy: 0.67 - 0s 57us/sample - loss: 0.6196 - accuracy: 0.6817 - binary\_accuracy: 0.6817 - val\_loss: 0.6124 - val\_accu  
racy: 0.6857 - val\_binary\_accuracy: 0.6857  
Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5952 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA:  
0s - loss: 0.6117 - accuracy: 0.6969 - binary\_accuracy: 0.69 - ETA: 0s - loss: 0.6137 - accuracy: 0.6835 - binary\_accu  
racy: 0.68 - 0s 62us/sample - loss: 0.6144 - accuracy: 0.6842 - binary\_accuracy: 0.6842 - val\_loss: 0.6073 - val\_accu  
racy: 0.6937 - val\_binary\_accuracy: 0.6937  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6385 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA:  
0s - loss: 0.6169 - accuracy: 0.6845 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6066 - accuracy: 0.7001 - binary\_accu



acy: 0.70 - 0s 65us/sample - loss: 0.6092 - accuracy: 0.6932 - binary\_accuracy: 0.6932 - val\_loss: 0.6024 - val\_accuracy: 0.7031 - val\_binary\_accuracy: 0.7031

Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5964 - accuracy: 0.8438 - binary\_accuracy: 0.84 - ETA: 0s - loss: 0.5979 - accuracy: 0.7228 - binary\_accuracy: 0.72 - ETA: 0s - loss: 0.5991 - accuracy: 0.7079 - binary\_accuracy: 0.70 - 0s 69us/sample - loss: 0.6042 - accuracy: 0.6984 - binary\_accuracy: 0.6984 - val\_loss: 0.5975 - val\_accuracy: 0.7082 - val\_binary\_accuracy: 0.7082

Train on 3220 samples, validate on 1381 samples

Epoch 1/10  
3220/3220 [=====] - ETA: 37s - loss: 0.7037 - accuracy: 0.5625 - binary\_accuracy: 0.562 - ETA: 0s - loss: 0.6951 - accuracy: 0.5961 - binary\_accuracy: 0.596 - ETA: 0s - loss: 0.6928 - accuracy: 0.5869 - binary\_accuracy: 0.58 - 1s 204us/sample - loss: 0.6876 - accuracy: 0.5885 - binary\_accuracy: 0.5885 - val\_loss: 0.6758 - val\_accuracy: 0.5844 - val\_binary\_accuracy: 0.5844

Epoch 2/10  
3220/3220 [=====] - ETA: 0s - loss: 0.7330 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.6705 - accuracy: 0.5945 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6705 - accuracy: 0.5849 - binary\_accuracy: 0.58 - 0s 58us/sample - loss: 0.6728 - accuracy: 0.5786 - binary\_accuracy: 0.5786 - val\_loss: 0.6646 - val\_accuracy: 0.5858 - val\_binary\_accuracy: 0.5858

Epoch 3/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6575 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6624 - accuracy: 0.5861 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6624 - accuracy: 0.5831 - binary\_accuracy: 0.58 - 0s 60us/sample - loss: 0.6634 - accuracy: 0.5773 - binary\_accuracy: 0.5773 - val\_loss: 0.6566 - val\_accuracy: 0.5844 - val\_binary\_accuracy: 0.5844

Epoch 4/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6892 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.6539 - accuracy: 0.5872 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6569 - accuracy: 0.5806 - binary\_accuracy: 0.58 - 0s 59us/sample - loss: 0.6560 - accuracy: 0.5770 - binary\_accuracy: 0.5770 - val\_loss: 0.6499 - val\_accuracy: 0.5865 - val\_binary\_accuracy: 0.5865

Epoch 5/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6373 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6493 - accuracy: 0.5828 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6545 - accuracy: 0.5724 - binary\_accuracy: 0.57 - 0s 58us/sample - loss: 0.6493 - accuracy: 0.5823 - binary\_accuracy: 0.5823 - val\_loss: 0.6437 - val\_accuracy: 0.5873 - val\_binary\_accuracy: 0.5873

Epoch 6/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6012 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6461 - accuracy: 0.5861 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6431 - accuracy: 0.5875 - binary\_accuracy: 0.58 - 0s 59us/sample - loss: 0.6431 - accuracy: 0.5866 - binary\_accuracy: 0.5866 - val\_loss: 0.6378 - val\_accuracy: 0.5923 - val\_binary\_accuracy: 0.5923

Epoch 7/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6436 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6342 - accuracy: 0.6067 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6351 - accuracy: 0.6019 - binary\_accuracy: 0.60 - 0s 57us/sample - loss: 0.6371 - accuracy: 0.5957 - binary\_accuracy: 0.5957 - val\_loss: 0.6320 - val\_accuracy: 0.6025 - val\_binary\_accuracy: 0.6025

Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.5935 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6359 - accuracy: 0.5977 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6340 - accuracy: 0.5977 - binary\_accuracy: 0.59 - 0s 57us/sample - loss: 0.6311 - accuracy: 0.6053 - binary\_accuracy: 0.6053 - val\_loss: 0.6263 - val\_accuracy: 0.6133 - val\_binary\_accuracy: 0.6133

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.6486 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6236 - accuracy: 0.6273 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6267 - accuracy: 0.6100 - binary\_accuracy: 0.61 - 0s 58us/sample - loss: 0.6254 - accuracy: 0.6121 - binary\_accuracy: 0.6121 - val\_loss: 0.6208 - val\_accuracy: 0.6271 - val\_binary\_accuracy: 0.6271

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.6472 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6207 - accuracy: 0.6124 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6211 - accuracy: 0.6139 - binary\_accuracy: 0.61 - 0s 56us/sample - loss: 0.6198 - accuracy: 0.6189 - binary\_accuracy: 0.6189 - val\_loss: 0.6154 - val\_accuracy: 0.6307 - val\_binary\_accuracy: 0.6307

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.001

|units: 88

|Score: 0.676080048084259

|Best step: 0

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 37s - loss: 0.7316 - accuracy: 0.4688 - binary\_accuracy: 0.468 - ETA: 0s - loss: 0.7599 - accuracy: 0.4199 - binary\_accuracy: 0.419 - ETA: 0s - loss: 0.7643 - accuracy: 0.3947 - binary\_accuracy: 0.39 - 1s 216us/sample - loss: 0.7630 - accuracy: 0.3938 - binary\_accuracy: 0.3938 - val\_loss: 0.7671 - val\_accuracy: 0.3845 - val\_binary\_accuracy: 0.3845

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.7992 - accuracy: 0.4375 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.7582 - accuracy: 0.3971 - binary\_accuracy: 0.39 - ETA: 0s - loss: 0.7576 - accuracy: 0.3898 - binary\_accuracy: 0.38 - 0s 60us/sample - loss: 0.7547 - accuracy: 0.3953 - binary\_accuracy: 0.3953 - val\_loss: 0.7590 - val\_accuracy: 0.3838 - val\_binary\_accuracy: 0.3838

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.8158 - accuracy: 0.2188 - binary\_accuracy: 0.21 - ETA: 0s - loss: 0.7446 - accuracy: 0.4121 - binary\_accuracy: 0.41 - ETA: 0s - loss: 0.7477 - accuracy: 0.4032 - binary\_accuracy: 0.40 - 0s 65us/sample - loss: 0.7471 - accuracy: 0.4040 - binary\_accuracy: 0.4040 - val\_loss: 0.7516 - val\_accuracy: 0.3874 - val\_binary\_accuracy: 0.3874

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.7631 - accuracy: 0.2812 - binary\_accuracy: 0.28 - ETA: 0s - loss: 0.7417 - accuracy: 0.4055 - binary\_accuracy: 0.40 - ETA: 0s - loss: 0.7418 - accuracy: 0.4030 - binary\_accuracy: 0.40 - 0s 65us/sample - loss: 0.7402 - accuracy: 0.4053 - binary\_accuracy: 0.4053 - val\_loss: 0.7449 - val\_accuracy: 0.3925 - val\_binary\_accuracy: 0.3925

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.7386 - accuracy: 0.3750 - binary\_accuracy: 0.37 - ETA: 0s - loss: 0.7334 - accuracy: 0.4040 - binary\_accuracy: 0.40 - ETA: 0s - loss: 0.7320 - accuracy: 0.4093 - binary\_accuracy: 0.40 - 0s 65us/sample - loss: 0.7339 - accuracy: 0.4087 - binary\_accuracy: 0.4087 - val\_loss: 0.7388 - val\_accuracy: 0.3968 - val\_binary\_accuracy: 0.3968

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.6628 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.7210 - accuracy: 0.4360 - binary\_accuracy: 0.43 - ETA: 0s - loss: 0.7268 - accuracy: 0.4224 - binary\_accuracy: 0.42 - 0s 65us/sample - loss: 0.7282 - accuracy: 0.4199 - binary\_accuracy: 0.4199 - val\_loss: 0.7332 - val\_accuracy: 0.4120 - val\_binary\_accuracy: 0.4120

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.7396 - accuracy: 0.3750 - binary\_accuracy: 0.37 - ETA: 0s - loss: 0.7286 - accuracy: 0.4238 - binary\_accuracy: 0.42 - ETA: 0s - loss: 0.7240 - accuracy: 0.4225 - binary\_accuracy: 0.42 - 0s 67us/sample - loss: 0.7229 - accuracy: 0.4248 - binary\_accuracy: 0.4248 - val\_loss: 0.7281 - val\_accuracy: 0.4164 - val\_binary\_accuracy: 0.4164

Epoch 8/10

3220/3220 [=====] - ETA: 0s - loss: 0.6960 - accuracy: 0.5000 - binary\_accuracy: 0.50 - ETA: 0s - loss: 0.7180 - accuracy: 0.4286 - binary\_accuracy: 0.42 - ETA: 0s - loss: 0.7204 - accuracy: 0.4328 - binary\_accuracy: 0.43 - 0s 67us/sample - loss: 0.7182 - accuracy: 0.4295 - binary\_accuracy: 0.4295 - val\_loss: 0.7235 - val\_accuracy: 0.4258 - val\_binary\_accuracy: 0.4258

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.7186 - accuracy: 0.4062 - binary\_accuracy: 0.40 - ETA: 0s - loss: 0.7154 - accuracy: 0.4505 - binary\_accuracy: 0.45 - ETA: 0s - loss: 0.7134 - accuracy: 0.4539 - binary\_accuracy: 0.45 - 0s 67us/sample - loss: 0.7138 - accuracy: 0.4503 - binary\_accuracy: 0.4503 - val\_loss: 0.7192 - val\_accuracy: 0.4424 - val\_binary\_accuracy: 0.4424

Epoch 10/10

```
3220/3220 [=====] - ETA: 0s - loss: 0.6794 - accuracy: 0.5625 - binary_accuracy: 0.56 - ETA:
0s - loss: 0.7076 - accuracy: 0.4531 - binary_accuracy: 0.45 - ETA: 0s - loss: 0.7094 - accuracy: 0.4639 - binary_accu
racy: 0.46 - 0s 67us/sample - loss: 0.7099 - accuracy: 0.4686 - binary_accuracy: 0.4686 - val_loss: 0.7153 - val_accu
racy: 0.4736 - val_binary_accuracy: 0.4736
Train on 3220 samples, validate on 1381 samples
Epoch 1/10
3220/3220 [=====] - ETA: 54s - loss: 0.4705 - accuracy: 0.8438 - binary_accuracy: 0.843 - ET
A: 0s - loss: 0.6655 - accuracy: 0.6098 - binary_accuracy: 0.609 - ETA: 0s - loss: 0.6689 - accuracy: 0.6082 - binary_
accuracy: 0.60 - 1s 273us/sample - loss: 0.6696 - accuracy: 0.6071 - binary_accuracy: 0.6071 - val_loss: 0.6675 - val_
accuracy: 0.6090 - val_binary_accuracy: 0.6090
Epoch 2/10
3220/3220 [=====] - ETA: 0s - loss: 0.5715 - accuracy: 0.6875 - binary_accuracy: 0.68 - ETA:
0s - loss: 0.6668 - accuracy: 0.6106 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6744 - accuracy: 0.6011 - binary_accu
racy: 0.60 - ETA: 0s - loss: 0.6629 - accuracy: 0.6132 - binary_accuracy: 0.61 - 0s 93us/sample - loss: 0.6668 - accur
acy: 0.6096 - binary_accuracy: 0.6096 - val_loss: 0.6649 - val_accuracy: 0.6104 - val_binary_accuracy: 0.6104
Epoch 3/10
3220/3220 [=====] - ETA: 0s - loss: 0.6912 - accuracy: 0.6250 - binary_accuracy: 0.62 - ETA:
0s - loss: 0.6565 - accuracy: 0.6159 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6646 - accuracy: 0.6107 - binary_accu
racy: 0.61 - ETA: 0s - loss: 0.6640 - accuracy: 0.6097 - binary_accuracy: 0.60 - 0s 75us/sample - loss: 0.6643 - accur
acy: 0.6102 - binary_accuracy: 0.6102 - val_loss: 0.6625 - val_accuracy: 0.6119 - val_binary_accuracy: 0.6119
Epoch 4/10
3220/3220 [=====] - ETA: 0s - loss: 0.5812 - accuracy: 0.7188 - binary_accuracy: 0.71 - ETA:
0s - loss: 0.6545 - accuracy: 0.6166 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6638 - accuracy: 0.6127 - binary_accu
racy: 0.61 - 0s 65us/sample - loss: 0.6619 - accuracy: 0.6134 - binary_accuracy: 0.6134 - val_loss: 0.6602 - val_accu
racy: 0.6162 - val_binary_accuracy: 0.6162
Epoch 5/10
3220/3220 [=====] - ETA: 0s - loss: 0.6132 - accuracy: 0.6875 - binary_accuracy: 0.68 - ETA:
0s - loss: 0.6677 - accuracy: 0.6066 - binary_accuracy: 0.60 - ETA: 0s - loss: 0.6623 - accuracy: 0.6167 - binary_accu
racy: 0.61 - ETA: 0s - loss: 0.6597 - accuracy: 0.6171 - binary_accuracy: 0.61 - 0s 81us/sample - loss: 0.6596 - accur
acy: 0.6168 - binary_accuracy: 0.6168 - val_loss: 0.6581 - val_accuracy: 0.6191 - val_binary_accuracy: 0.6191
Epoch 6/10
3220/3220 [=====] - ETA: 0s - loss: 0.6317 - accuracy: 0.6562 - binary_accuracy: 0.65 - ETA:
0s - loss: 0.6616 - accuracy: 0.6149 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6590 - accuracy: 0.6155 - binary_accu
racy: 0.61 - 0s 67us/sample - loss: 0.6575 - accuracy: 0.6174 - binary_accuracy: 0.6174 - val_loss: 0.6561 - val_accu
racy: 0.6191 - val_binary_accuracy: 0.6191
Epoch 7/10
3220/3220 [=====] - ETA: 0s - loss: 0.6974 - accuracy: 0.6250 - binary_accuracy: 0.62 - ETA:
0s - loss: 0.6681 - accuracy: 0.6078 - binary_accuracy: 0.60 - ETA: 0s - loss: 0.6717 - accuracy: 0.5984 - binary_accu
racy: 0.59 - ETA: 0s - loss: 0.6654 - accuracy: 0.6049 - binary_accuracy: 0.60 - ETA: 0s - loss: 0.6623 - accuracy: 0.
6091 - binary_accuracy: 0.60 - ETA: 0s - loss: 0.6558 - accuracy: 0.6173 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6
546 - accuracy: 0.6202 - binary_accuracy: 0.62 - ETA: 0s - loss: 0.6519 - accuracy: 0.6240 - binary_accuracy: 0.62 - 1
s 203us/sample - loss: 0.6555 - accuracy: 0.6183 - binary_accuracy: 0.6183 - val_loss: 0.6542 - val_accuracy: 0.6235 -
val_binary_accuracy: 0.6235
Epoch 8/10
3220/3220 [=====] - ETA: 0s - loss: 0.7335 - accuracy: 0.5000 - binary_accuracy: 0.50 - ETA:
0s - loss: 0.6637 - accuracy: 0.6000 - binary_accuracy: 0.60 - ETA: 0s - loss: 0.6704 - accuracy: 0.5974 - binary_accu
racy: 0.59 - ETA: 0s - loss: 0.6595 - accuracy: 0.6125 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6587 - accuracy: 0.
6118 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6542 - accuracy: 0.6174 - binary_accuracy: 0.61 - ETA: 0s - loss: 0.6
```

516 - accuracy: 0.6246 - binary\_accuracy: 0.62 - 1s 161us/sample - loss: 0.6536 - accuracy: 0.6214 - binary\_accuracy: 0.6214 - val\_loss: 0.6525 - val\_accuracy: 0.6242 - val\_binary\_accuracy: 0.6242

Epoch 9/10

3220/3220 [=====] - ETA: 0s - loss: 0.6102 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6641 - accuracy: 0.6060 - binary\_accuracy: 0.60 - ETA: 0s - loss: 0.6551 - accuracy: 0.6171 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6534 - accuracy: 0.6204 - binary\_accuracy: 0.62 - 0s 89us/sample - loss: 0.6519 - accuracy: 0.6239 - binary\_accuracy: 0.6239 - val\_loss: 0.6508 - val\_accuracy: 0.6249 - val\_binary\_accuracy: 0.6249

Epoch 10/10

3220/3220 [=====] - ETA: 0s - loss: 0.6507 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6353 - accuracy: 0.6463 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6403 - accuracy: 0.6431 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6490 - accuracy: 0.6254 - binary\_accuracy: 0.62 - 0s 83us/sample - loss: 0.6502 - accuracy: 0.6252 - binary\_accuracy: 0.6252 - val\_loss: 0.6492 - val\_accuracy: 0.6249 - val\_binary\_accuracy: 0.6249

Train on 3220 samples, validate on 1381 samples

Epoch 1/10

3220/3220 [=====] - ETA: 42s - loss: 0.6774 - accuracy: 0.5312 - binary\_accuracy: 0.531 - ETA: 0s - loss: 0.6731 - accuracy: 0.5570 - binary\_accuracy: 0.557 - ETA: 0s - loss: 0.6751 - accuracy: 0.5642 - binary\_accuracy: 0.56 - 1s 222us/sample - loss: 0.6724 - accuracy: 0.5717 - binary\_accuracy: 0.5717 - val\_loss: 0.6668 - val\_accuracy: 0.5764 - val\_binary\_accuracy: 0.5764

Epoch 2/10

3220/3220 [=====] - ETA: 0s - loss: 0.6407 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6636 - accuracy: 0.5884 - binary\_accuracy: 0.58 - ETA: 0s - loss: 0.6633 - accuracy: 0.5933 - binary\_accuracy: 0.59 - 0s 60us/sample - loss: 0.6664 - accuracy: 0.5901 - binary\_accuracy: 0.5901 - val\_loss: 0.6611 - val\_accuracy: 0.6032 - val\_binary\_accuracy: 0.6032

Epoch 3/10

3220/3220 [=====] - ETA: 0s - loss: 0.6817 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6614 - accuracy: 0.5994 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6628 - accuracy: 0.6025 - binary\_accuracy: 0.60 - 0s 57us/sample - loss: 0.6610 - accuracy: 0.6109 - binary\_accuracy: 0.6109 - val\_loss: 0.6558 - val\_accuracy: 0.6227 - val\_binary\_accuracy: 0.6227

Epoch 4/10

3220/3220 [=====] - ETA: 0s - loss: 0.6237 - accuracy: 0.6562 - binary\_accuracy: 0.65 - ETA: 0s - loss: 0.6557 - accuracy: 0.6198 - binary\_accuracy: 0.61 - ETA: 0s - loss: 0.6556 - accuracy: 0.6299 - binary\_accuracy: 0.62 - 0s 65us/sample - loss: 0.6560 - accuracy: 0.6252 - binary\_accuracy: 0.6252 - val\_loss: 0.6511 - val\_accuracy: 0.6437 - val\_binary\_accuracy: 0.6437

Epoch 5/10

3220/3220 [=====] - ETA: 0s - loss: 0.6350 - accuracy: 0.5938 - binary\_accuracy: 0.59 - ETA: 0s - loss: 0.6475 - accuracy: 0.6471 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6520 - accuracy: 0.6431 - binary\_accuracy: 0.64 - 0s 67us/sample - loss: 0.6516 - accuracy: 0.6441 - binary\_accuracy: 0.6441 - val\_loss: 0.6468 - val\_accuracy: 0.6698 - val\_binary\_accuracy: 0.6698

Epoch 6/10

3220/3220 [=====] - ETA: 0s - loss: 0.6065 - accuracy: 0.7812 - binary\_accuracy: 0.78 - ETA: 0s - loss: 0.6529 - accuracy: 0.6482 - binary\_accuracy: 0.64 - ETA: 0s - loss: 0.6491 - accuracy: 0.6567 - binary\_accuracy: 0.65 - 0s 73us/sample - loss: 0.6475 - accuracy: 0.6590 - binary\_accuracy: 0.6590 - val\_loss: 0.6429 - val\_accuracy: 0.6894 - val\_binary\_accuracy: 0.6894

Epoch 7/10

3220/3220 [=====] - ETA: 0s - loss: 0.6864 - accuracy: 0.5625 - binary\_accuracy: 0.56 - ETA: 0s - loss: 0.6449 - accuracy: 0.6757 - binary\_accuracy: 0.67 - ETA: 0s - loss: 0.6466 - accuracy: 0.6757 - binary\_accuracy: 0.67 - 0s 69us/sample - loss: 0.6438 - accuracy: 0.6801 - binary\_accuracy: 0.6801 - val\_loss: 0.6393 - val\_accuracy: 0.6980 - val\_binary\_accuracy: 0.6980

Epoch 8/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6454 - accuracy: 0.6875 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6432 - accuracy: 0.6815 - binary\_accuracy: 0.68 - ETA: 0s - loss: 0.6427 - accuracy: 0.6871 - binary\_accuracy: 0.68 - 0s 68us/sample - loss: 0.6404 - accuracy: 0.6910 - binary\_accuracy: 0.6910 - val\_loss: 0.6361 - val\_accuracy: 0.7161 - val\_binary\_accuracy: 0.7161  
Epoch 9/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6533 - accuracy: 0.6250 - binary\_accuracy: 0.62 - ETA: 0s - loss: 0.6364 - accuracy: 0.7214 - binary\_accuracy: 0.72 - ETA: 0s - loss: 0.6384 - accuracy: 0.7146 - binary\_accuracy: 0.71 - 0s 68us/sample - loss: 0.6373 - accuracy: 0.7134 - binary\_accuracy: 0.7134 - val\_loss: 0.6331 - val\_accuracy: 0.7335 - val\_binary\_accuracy: 0.7335  
Epoch 10/10  
3220/3220 [=====] - ETA: 0s - loss: 0.6659 - accuracy: 0.7500 - binary\_accuracy: 0.75 - ETA: 0s - loss: 0.6377 - accuracy: 0.7257 - binary\_accuracy: 0.72 - ETA: 0s - loss: 0.6318 - accuracy: 0.7332 - binary\_accuracy: 0.73 - ETA: 0s - loss: 0.6345 - accuracy: 0.7207 - binary\_accuracy: 0.72 - ETA: 0s - loss: 0.6349 - accuracy: 0.7300 - binary\_accuracy: 0.73 - 0s 110us/sample - loss: 0.6345 - accuracy: 0.7311 - binary\_accuracy: 0.7311 - val\_loss: 0.6304 - val\_accuracy: 0.7567 - val\_binary\_accuracy: 0.7567

## Trial complete

## Trial summary

### Hp values:

|learning\_rate: 0.0001

|units: 84

|Score: 0.6183924674987793

|Best step: 0

```
In [123]: best_linear_model = classtuner.get_best_models(num_models = 1)

best_linear_model = best_linear_model[0]
```

```
In [124]: ### Use our best model and reverify it against our test data

spam_linear_history = best_linear_model.fit(X_train_standard, y_train,
                                             validation_data = (X_test_standard, y_test), epochs = 50, verbose = 0)
```

```
In [125]: train_acc = best_linear_model.evaluate(X_train_standard, y_train, verbose=0)
test_acc = best_linear_model.evaluate(X_test_standard, y_test, verbose=0)
```

```
In [126]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Plot scores on each trial for nested CV


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot nested scores for each classifier - quickly visual the best performing model
## This is WITHOUT having changed any of the default parameters
plt.plot(spam_linear_history.history['accuracy'], label = "Training Data Accuracy")
plt.plot(spam_linear_history.history['val_accuracy'], label = "Testing Data Accuracy")

## Give some labels and title
plt.xlabel("No. of Epochs")
plt.ylabel("Score")

## Title and Legend
plt.title("Model Accuracy Comparison - Training & Test Data")
plt.legend()

## Show the graph
plt.show()
```



```
In [127]: ## Create three empty lists to store my precision, recall, and average precision scores
precision, recall, average_precision = [], [], []

## Get probabilities for our labels
probas_ = best_linear_model.predict_proba(X_test)
```



```
In [128]: ## Create my false positive rate, true positive rate, and threshold using my test data
fpr, tpr, thresholds = roc_curve(y_test, linear_probas_, pos_label = 1)

## Create precision and recall scores to plot with
precision_score, recall_score, _ = precision_recall_curve(y_test, probas_)

## Calculate the overall AUC for the model
auc = np.trapz(tpr, fpr)

## Save the average precision for our model
avg_precision = average_precision_score(y_test, probas_)
```

```
In [129]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Create a new figure to plot
plt.figure(figsize= (15, 10))

lw = 2

## Draw the line for my fpr and tpr
plt.plot(fpr, tpr, color = 'darkorange',
         label = 'ROC Curve (area = %0.2f)' % auc)

## Put in a line to demonstrate blind luck
plt.plot([0, 1], [0, 1], color = 'navy', linestyle = '--', label = 'Luck')


## Set the limits of the plot for better visualization
plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.05])

## Set labels for x and y
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

## Set a title and legend
plt.title('ROC Curve')
plt.legend(loc = 'lower right')

## Show the curve!
plt.show()
```

ROC Curve



```
In [130]: ### Reset seaborn to the default background - for better viewing
sns.set_style("white")

## Create a new figure to plot
plt.figure(figsize= (15, 10))

lw = 2

## Plot the model's Precision Recall Curve
plt.plot(precision_score, recall_score, label='Precision-Recall curve (area = {})'.format(round(avg_precision, 4)))

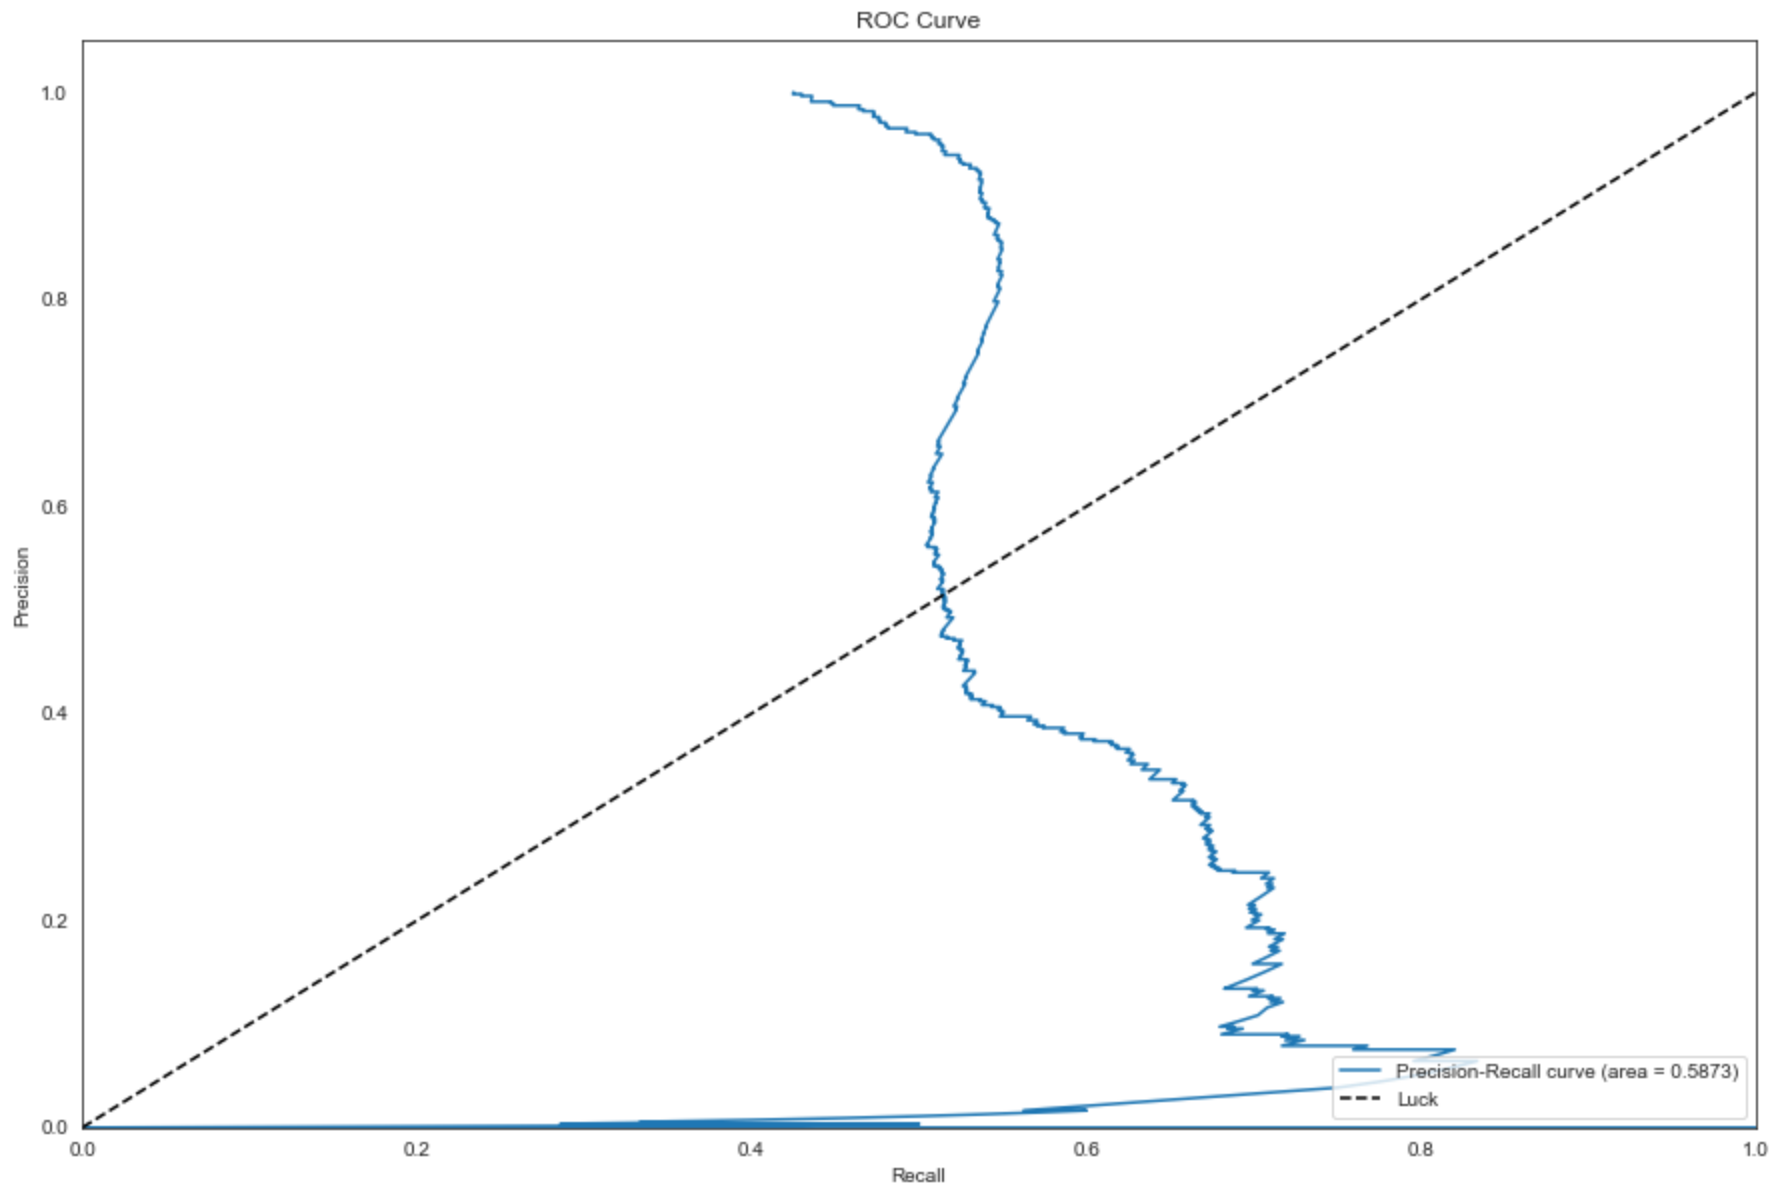
## Put in a line to demonstrate blind luck
plt.plot([0, 1], [0, 1], color = 'black', linestyle = '--', label = 'Luck')

## Set the limits so they start at zero
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])

## Set labels for x and y
plt.xlabel('Recall')
plt.ylabel('Precision')

## Set a title and legend
plt.title('ROC Curve')
plt.legend(loc = 'lower right')

## Show the curve!
plt.show()
```



**Interesting! The classifier behaved pretty similarly to the one generated above but the precision-recall curve is wildly different, and converges very strangely. I would be extremely cautious using this model.**

```
In [132]: index = 0

for t in thresholds:

    predict_thre = np.where(probas_ > t, 1, 0) ## prediction based on the preset threshold
    clf_matrix = confusion_matrix(y_test, predict_thre)
    Cost_List[index] = clf_matrix[0][0]*cost_matrix[0][0] +clf_matrix[0][1]*cost_matrix[0][1] +clf_matrix[1][0]*cost_ma
    trix[1][0] +clf_matrix[1][1]*cost_matrix[1][1]
    index+=1


## Set the figure size
plt.figure(figsize= (15, 10))

## Plot each Cost Line individually
plt.plot(thresholds, Cost_List, 'r--', label = "Average Cost for Misclassification")

## Give some labels
plt.xlabel("thresholds")
plt.ylabel("Cost")

## Title and Legend
plt.title("Cost Matrix")
plt.legend(loc = 'right')

## Show the Cost Matrix Analysis
plt.show()
```




The average cost of misclassification is pretty similar to our first Keras model.

## D. Conclusions / Model Evaluation

Based on the analysis completed, my selection for the best identifier for this classification would be using a Support Vector Machine with these hyper parameters tuned:


```
▼ ## Get our best params and their scores  
print(svclass.best_params_)  
print()  
  
## Print out how well it performed using the best params  
print(svclass.best_score_)  
  
## save our best params so we can use them in our actual SVC model!  
best_svc_params = svclass.best_params_  
  
{'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}  
  
0.9347826086956522
```






**As we are dealing with unbalanced classes, this metric is not enough to make sure the model is accurate. That's where a precision recall curve helps accurately determine model performance.**

Precision Recall Curve



**With the model's performance, as compared to random luck / guessing, and being a perfect world, it does very well. When you compare this Precision-Recall curves to the ones generated by the classifier we created with Keras, it is a much better performing classifier.**

Average Misclassification Cost Function



As it gets closer to identifying spam, the "preferred" Support Vector Machine classifier punishes mistakes more harshly, which is the behavior that I was hoping for. Again, the idea is to punish misidentifying potential spam as not spam, so we want the model to be accurate here.

```
## Create a report to show our precision(accuracy), recall, and f1 for predictions
report = classification_report(y_test, predicted)
print(report)
```

```
[[798  39]
 [ 51 493]]
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	837
1	0.93	0.91	0.92	544
accuracy			0.93	1381
macro avg	0.93	0.93	0.93	1381
weighted avg	0.93	0.93	0.93	1381

**The support vector machine does well at both classification tasks, and might be improved with additional data to learn and refine its search.**

**Thank you for reading my analysis!**